

Process Management

Process Abstraction in Unix

Lecture 2b – Unix Case study

Overview

- Process in Unix
 - Identification
 - Information
 - Creation
 - Termination
 - Parent-Child Synchronization
- Process states in Unix
- Implementation Issues

Process Abstraction in Unix

Identification

- PID: Process ID (an integer value)

Information

- Process State:
 - Running, Sleeping, Stopped, **Zombie**
- Parent PID:
 - PID of the parent process
- Cumulative CPU time:
 - Total amount of CPU time used so far
- etc

- Unix Command for process information:
 - **ps** (short for process status)

Process Creation in Unix: **fork()**

- The main way to create a new process

Header File	<code>#include <unistd.h></code>
Syntax	<code>int <i>fork</i>();</code>

- Returns:
 - **PID** of the newly created process (for parent process) OR
 - 0 (for child process)
- Header files are system dependent
 - "*man fork*" to locate the right files for your system!

Process Creation in Unix: **fork ()** (cont)

■ Behavior:

- ❑ Creates a new process (known as ***child process***)
- ❑ Child process is a **duplicate** of the current executable image
 - i.e., same code, same address space, etc.
 - Memory in child is a **COPY** of the parent (i.e., not shared)
 - ❑ Implemented using copy-on-write
- ❑ Child **differs only in:**
 - Process id (PID)
 - Parent (PPID)
 - ❑ Parent = The process which executed the fork()
 - **fork ()** return value

fork () : Example

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main()
{
    printf("I am ONE\n");
    fork();
    printf("I am seeing DOUBLE\n");

    return 0;
}
```

■ Question:

- What do you think is the output?

`fork ()`: Example Explained

- **Both** parent and child processes continue executing after `fork ()`
- A common usage is to use the parent/child process differently
 - For example:
 - The parent spawn off a child to carry out some work
 - And then the parent is ready to take another task
 - **Use the return value of `fork ()` to distinguish parent and child**

fork () : Parent and Child Example

```
... ..  
int result;  
  
result = fork();  
{ if (result != 0) {  
    printf("P:My Id is %i\n", getpid());  
    printf("P:Child Id is %i\n", result);  
} else {  
    printf("C:My Id is %i\n", getpid() );  
    printf("C:Parent Id is %i\n", getppid() );  
}  
... ..
```

Parent Process

Child Process

`fork()` : Independent Memory Space

```
... ..
int var = 1234;
int result;

result = fork();
if (result != 0) {
    printf("Parent: Var is %i\n", var);
    var++;
    printf("Parent: Var is %i\n", var);
} else {
    printf("Child: Var is %i\n", var);
    var--;
    printf("Child: Var is %i\n", var);
}
... ..
```

■ Question:

- Is there ONE or TWO `var` variable?

Executing A New Program/Image

- **fork ()** itself is not useful:
 - ❑ You still need to provide the full code for the child process
 - ❑ What if we want to execute ***another existing program*** instead?
- Make use of the **exec ()** system calls family
 - ❑ Many variants:
 - **execv, execl, execl, execlv, execlp**, etc.
 - ❑ Will touch on:
 - **execl**
 - ❑ Others are similar ("man XXX" to find out more)

Sidetrack: Command Line Argument in C

- You can pass arguments to a program in C

- e.g. **a.exe 1 2 3 hello**

```
int main( int argc, char* argv[] )  
{  
    //use argc and argv  
}
```

- **argc:**

- Number of command line arguments
 - Including the program name itself

- **argv:**

- A char strings array
 - Each element in **argv[]** is a C character string

C Command Line Argument: Example

```
int main( int argc, char* argv[] )
{
    int i;

    for (i = 0; i < argc; i++){
        printf("Arg %i: %s\n",i, argv[i] );
    }
    return 0;
}
```

■ Example Run:

a.out 123 hello world

■ Output:

Arg 0: a.out

Arg 1: 123

Arg 2: hello

Arg 3: world

exec1 () System Call

- To **replace** current executing process image with a new one
 - ❑ Code replacement
 - ❑ PID and other information still intact

Header File	<code>#include <unistd.h></code>
Syntax	<pre>int exec1(const char *path, const char *arg0, ..., const char *argN, NULL);</pre>

- ❑ **path**: Location of the executable
- ❑ **arg0, ..., argN**: Command Line Argument(s)
- ❑ **NULL**: To indicate end of argument list

`exec1()` : Simple Example

```
int main()
{
    exec1( "/bin/ls", "ls", "-al", NULL);
}
```

■ Note:

- ❑ `Path = "/bin/ls"`

- The "`dir`" command in unix, to list the files in directory

- ❑ `arg0 = "ls"`

- The program name

- ❑ `arg1 = "-al"`

- The above is exactly the same as executing:

`ls -al`

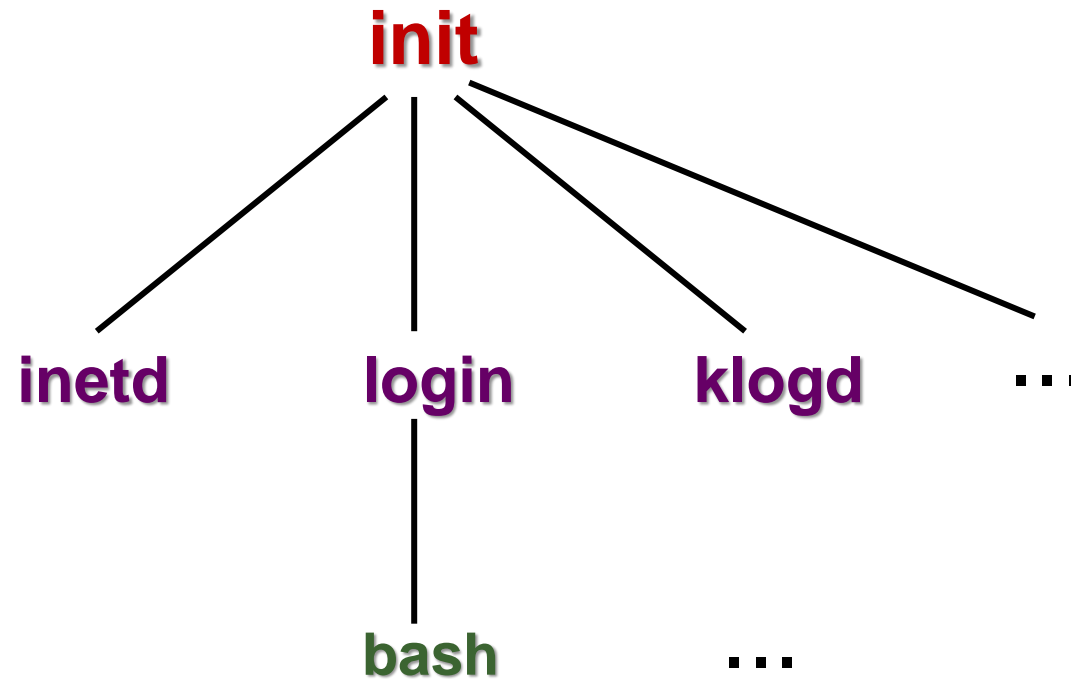
Hmm... **fork()** + **exec()** ?

- By combining the two mechanisms, we can:
 - Spawn off a child process
 - Let the child process perform a task through **exec()**
 - Meanwhile, the parent process is still around
 - To accept another request
- This combination of mechanisms is the main way in Unix:
 - To get a new process for running a new program

The Master Process

- Question:
 - ❑ If every process has parent, then which process is the "commonest ancestor"?
- Special initial process:
 - ❑ `init` process
 - ❑ Created in kernel at boot up time
 - ❑ Traditionally has a PID = 1
 - ❑ Watches for other processes and respawns where needed
- **`fork()`** creates process tree:
 - ❑ `init` is the root process

Process Tree Example (simplified)



Note: just a simple example, actual process tree varies according to Unix setup

Process Termination in Unix

■ To end execution of process:

Header File	<code>#include <stdlib.h></code>
Syntax	<code>void exit(int status);</code>

- ❑ Status is returned to the parent process (more later)
- ❑ Unix Convention:
 - 0 = Normal Termination (successful execution)
 - ! 0 = To indicate problematic execution
- ❑ The function **does not return!**

Process On Exit

■ Process finished execution

- **Most** system resources used by process are released on exit

- E.g. File descriptors

- Each opened file in C has a file descriptor attach to it

- Similar to File object in Java, File Stream Object in C++

- Some basic process resources **not releasable**:

- PID & status needed

- For parent-children synchronization

- Process accounting info, e.g., cpu time

➔ Process table entry **may be** still needed

Implicit `exit()`

- Most programs have no explicit `exit()` call
- Example:

```
int main()  
{  
    printf("Just to say goodbye!\n");  
}
```

- Return from `main()` implicitly calls `exit()`
 - Open files also get flushed automatically!

Parent/Child Synchronization in Unix

- Parent process can wait for child process to terminates

Header
File

```
#include <sys/types.h>
#include <sys/wait.h>
```

Syntax

```
int wait( int *status );
```

- ❑ Returns the PID of the terminated child process
- ❑ status (passed by address):
 - Stores the exit status of the terminated child process
 - Use **NULL** if you do not need/want this info

Parent/Child Synchronization in Unix

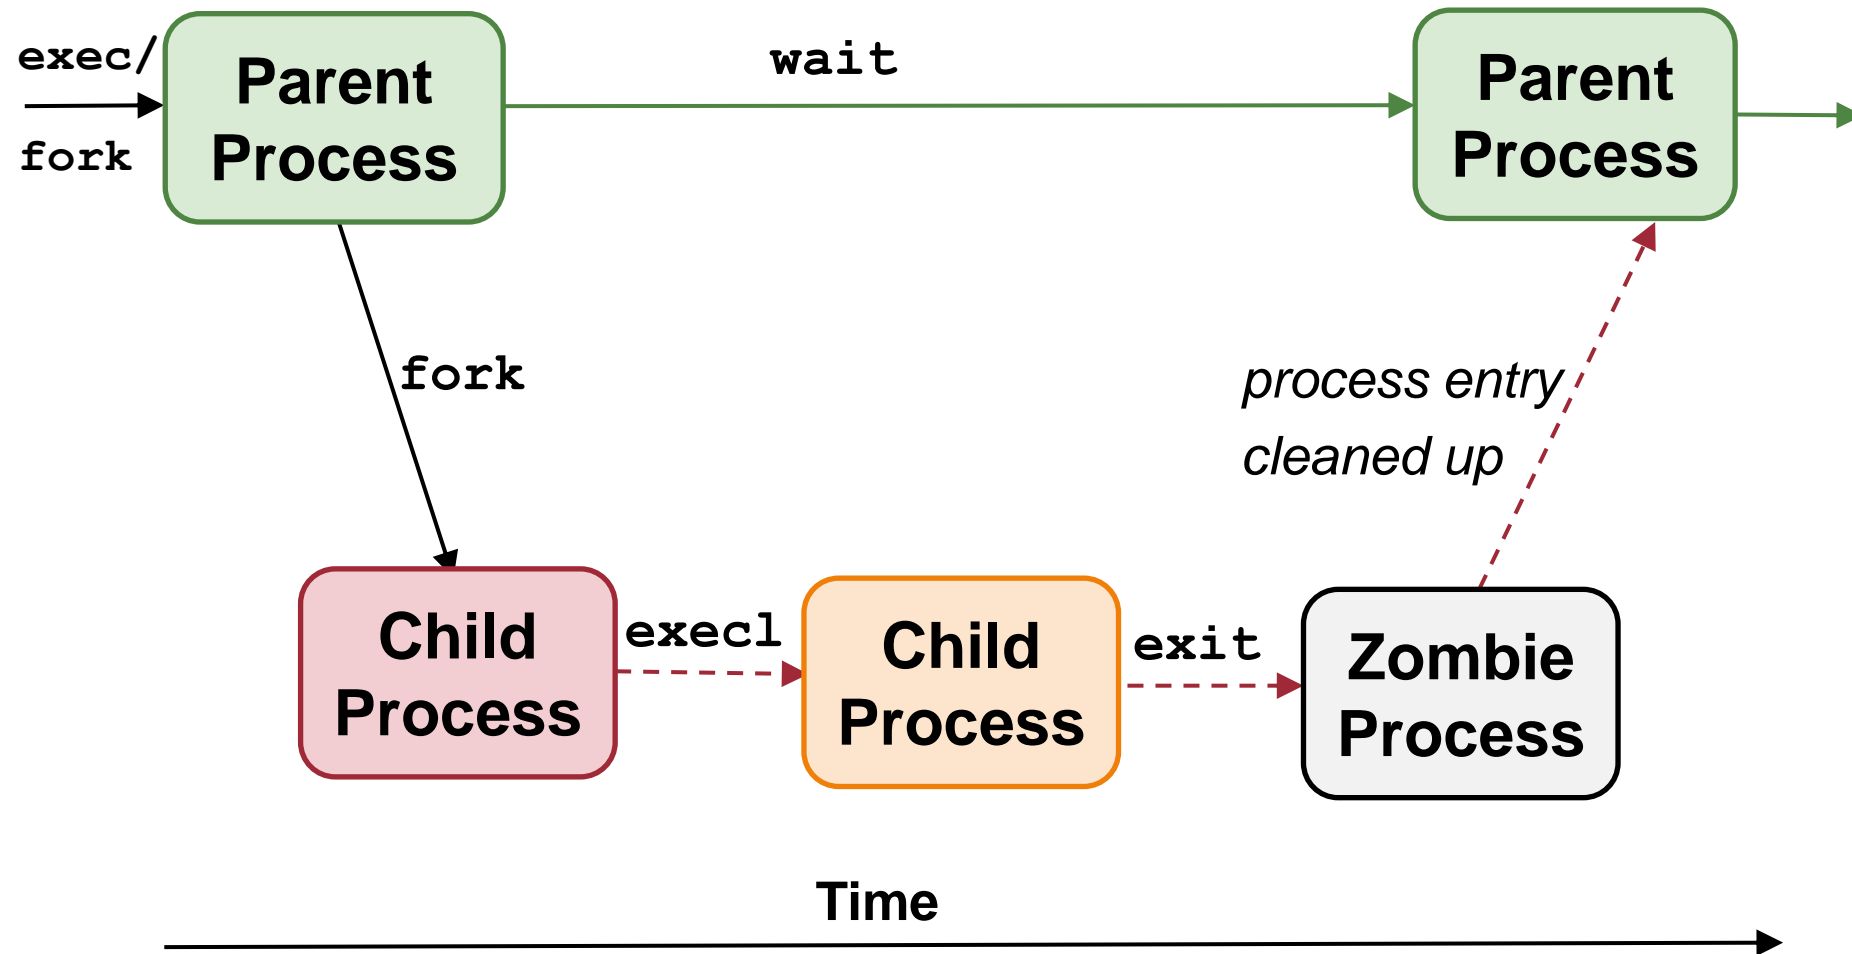
■ Behavior:

- ❑ The call is blocking:
 - Parent process blocks until at least one child terminates
- ❑ The call cleans up ***remainder*** of child system resources
 - Those not removed on `exit()`
 - Kill zombie process 😊

■ Other variants of `wait()` :

- ❑ `waitpid()`
 - Wait for a specific child process
- ❑ `waitid()`
 - Wait for any child process to **change status**
- ❑ etc.

Process Interaction in Unix



Note: example uses one ordering of execution, others are possible!

`wait()` "creates" zombies!!

- On process exit: (see previous slide)
 - ❑ most of the resources are released
 - ❑ becomes **zombie**
- ❑ **Cannot delete** all process info
 - What if parent asks for the info in a `wait()` call?
 - Can be **cleaned up** only when `wait()` happens
- ❑ **Cannot kill** zombie
 - The process is already **dead**!

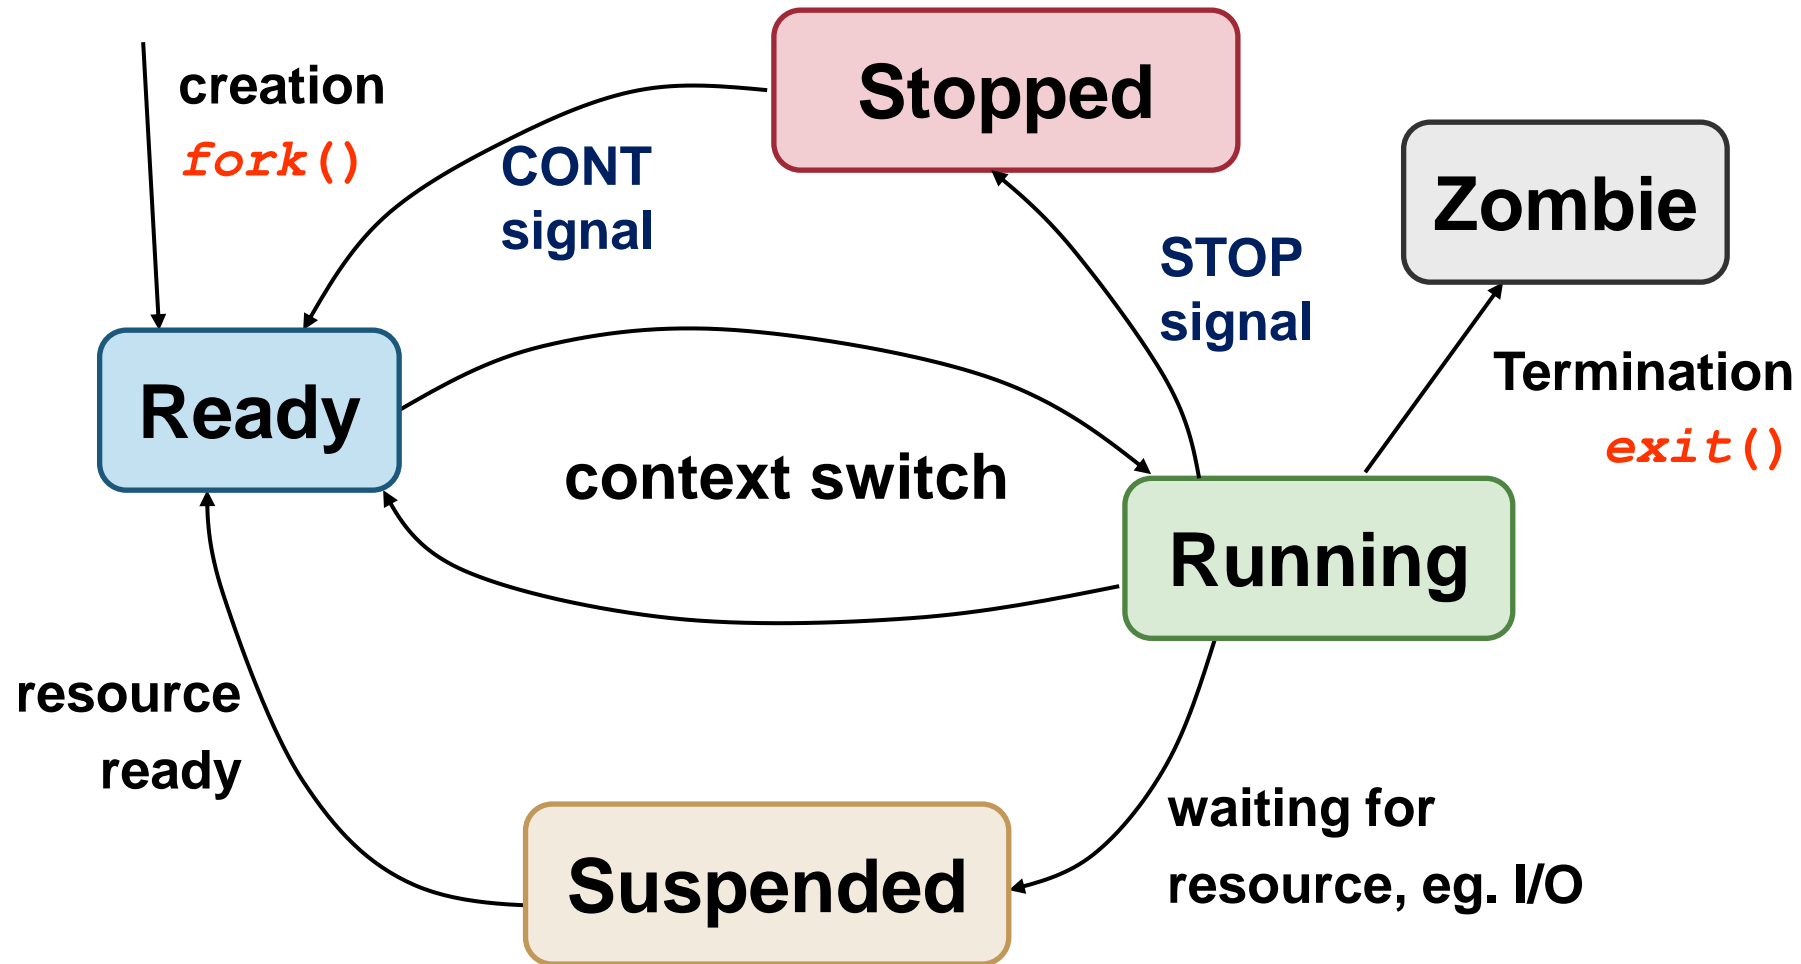
Zombie Process and Orphan Process

1. Orphan: parent process terminates before child process:
 - ❑ `init` process becomes "pseudo" parent of child processes
 - ❑ Child termination sends signal to `init`, which utilizes `wait()` to cleanup
2. Zombie: Child process terminates before parent but parent did not call `wait`:
 - ❑ Child process become a zombie process
 - ❑ Can fill up process table
 - May need a reboot to clear the table on older Unix implementations

Summary of Unix Process System calls

- **fork()** :
 - Process creation
- **exec()** family:
 - Change executing image/program
 - **execl**, **execv**, **execve**, **execle**, **execvp**
- **exit()** :
 - Process termination
- **wait()** family:
 - Get exit status, synchronize with child
 - **wait**, **waitpid**, **waitid**, **etc**
- **getpid()** family:
 - Get process information
 - **getpid**, **getppid**, **etc**

Process State Diagram in Unix



IMPLEMENTATION ISSUES

Implementing **fork()**

- Behavior of **fork()** :
 - ❑ Makes an almost exact copy of parent process
- Simplified implementation:
 1. Create address space of child process
 2. Allocate **p' = new PID**
 3. Create kernel process data structures
 - E.g. Entry in Process Table
 4. Copy kernel environment of parent process
 - E.g., Priority (for process scheduling)
 5. Initialize child process context:
 - **PID= p' , PPID=parent id, zero CPU time**

Implementing **fork()** (cont)

6. Copy memory regions from parent
 - Program, Data, Stack
 - Very expensive operation that can be optimized (more later)
7. Acquires shared resources:
 - Open files, current working directory, etc.
8. Initialize hardware context for child process:
 - Copy registers, etc., from parent process
9. Child process is now ready to run
 - add to scheduler queue

Memory Copy Operation

- Memory copy is very expensive:
 - ❑ Potentially need to copy the whole memory space
- Observations:
 - ❑ The child process will not access the whole memory range right away
 - ❑ Additionally:
 - If child just read from a location:
 - ❑ Remain unchanged
 - ❑ Can use a shared version
 - Only when write is perform on a location:
 - ❑ Then two independent copies are needed

Memory Copy Optimization

- **Copy on Write** is a possible optimization for memory copy operation:
 - ❑ Only duplicate a “memory location” when it is written to
 - ❑ Otherwise parent and child share the same “memory location”
- Note that, actually:
 - ❑ Memory is organized into memory pages
 - A consecutive range of memory locations
 - ❑ Memory is managed on a page level
 - Instead of individual location
 - ❑ **Will be covered in details in Memory Management part of lecture**

Modern Take on **fork()**

- **fork()** system call is part of the Unix design
 - inherited by most (all?) variants
- However, it is not versatile:
 - A thorough duplication of the parent process
- There are scenarios where a partial duplication may be preferred:
 - e.g. parent and child shares some of the memory regions, or some other resources
- Linux provides **clone()** which supersedes **fork()**

Summary

- Covered most of the process operations available in Unix:
 - ❑ Creation through `fork()`
 - ❑ Change execution through `exec()`
 - ❑ Termination through `exit()`
 - ❑ Synchronization (Parent \leftrightarrow Child) through `wait()`
- Process States
 - ❑ Process state diagram
- Implementation issues with `fork()`