

Memory Management

Disjoint Memory Schemes

Lecture 8

Overview

■ Paging Scheme

- ❑ Basic Idea
- ❑ Logical Address Translation
- ❑ Hardware Support
- ❑ Protection & Page Sharing

■ Segmentation Scheme

- ❑ Motivation
- ❑ Basic Idea
- ❑ Logical Address Translation
- ❑ Hardware Support

■ Segmentation with Paging Scheme

Disjoined Memory Space with Paging

- In lecture 7, memory management was discussed with two assumptions:
 1. Process occupies a contiguous physical memory region
 2. Physical memory is large enough to hold the whole process

- Let us remove assumption(1) in this lecture:
 - ❑ Process memory space can now be in **disjoint physical memory locations**
 - ❑ Via a main mechanism known as **paging**

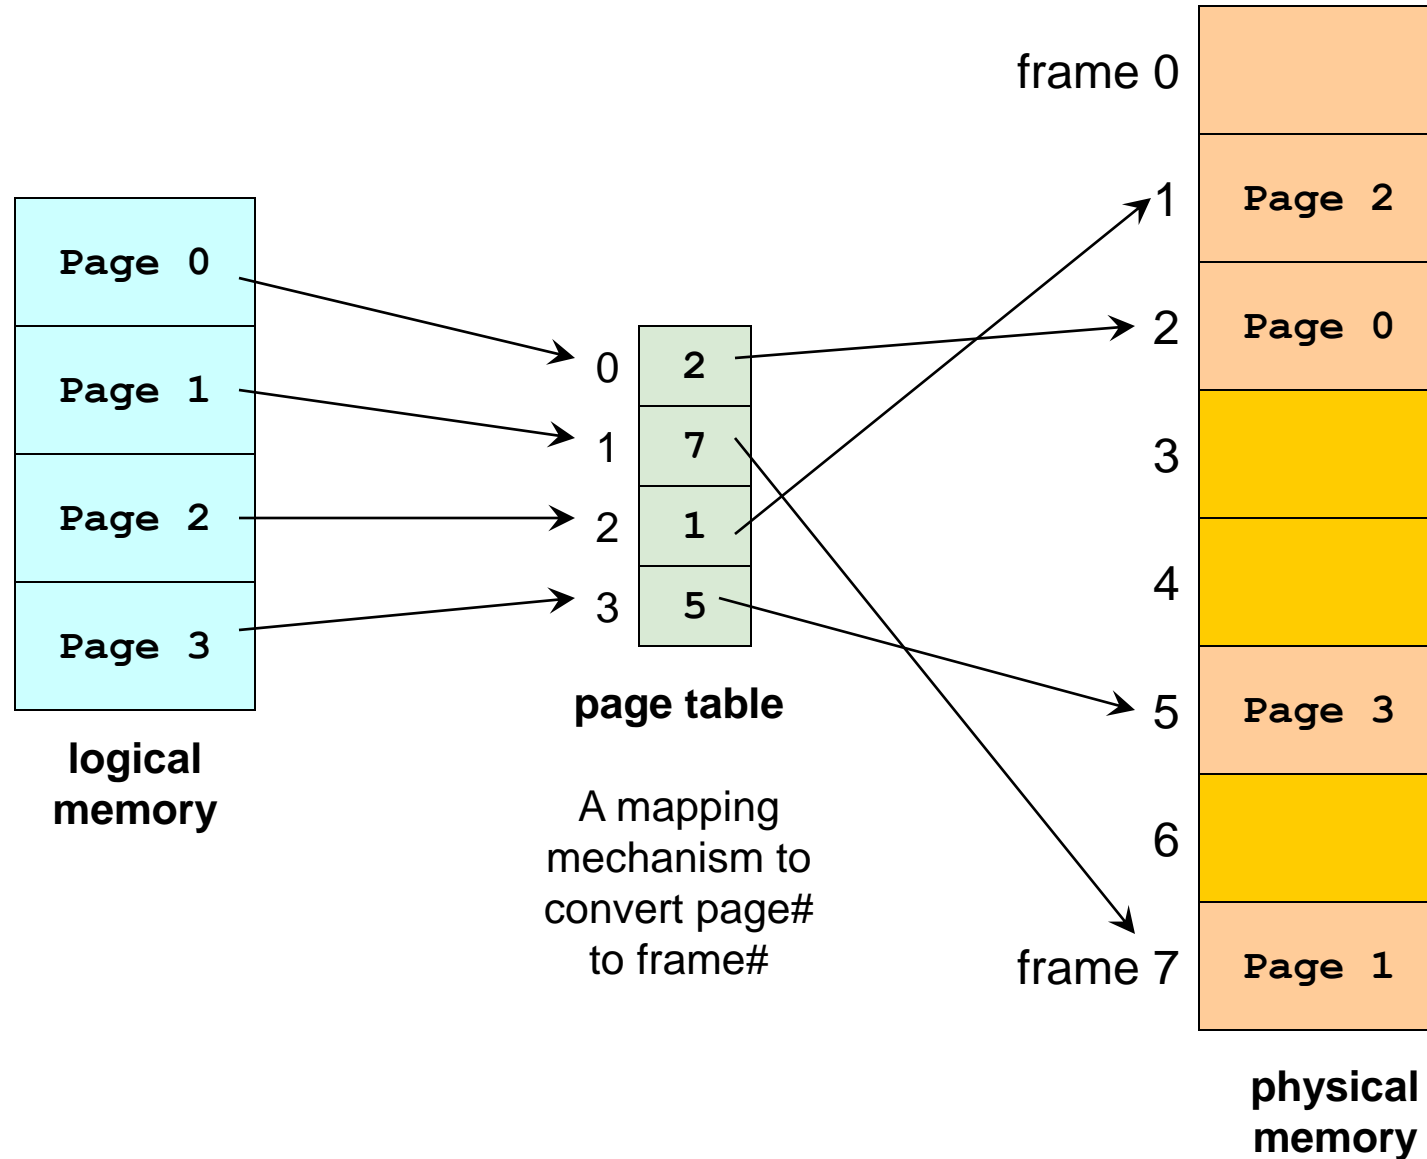
Paging Scheme: Basic Idea

- **Physical memory** is split into regions of fixed size (decided by hardware)
 - known as **physical frames**
- **Logical memory** of a process is split into regions of the **same size**
 - known as **logical pages**
- At execution time, pages of a process are loaded into **any available** memory frame
 - ➔ Logical memory space remain contiguous
 - ➔ Occupied physical memory region can be disjointed!

Page Table: Lookup Mechanism

- In **contiguous memory allocation**, it is very simple to keep track of the usage of a process:
 - Minimally, **starting address** and **size of process**
- Under **paging scheme**:
 - Logical page \leftrightarrow physical frame mapping is no longer straightforward
 - Need a **lookup table** to provide the translation
 - This structure is known as a **Page Table**

Paging: Illustration

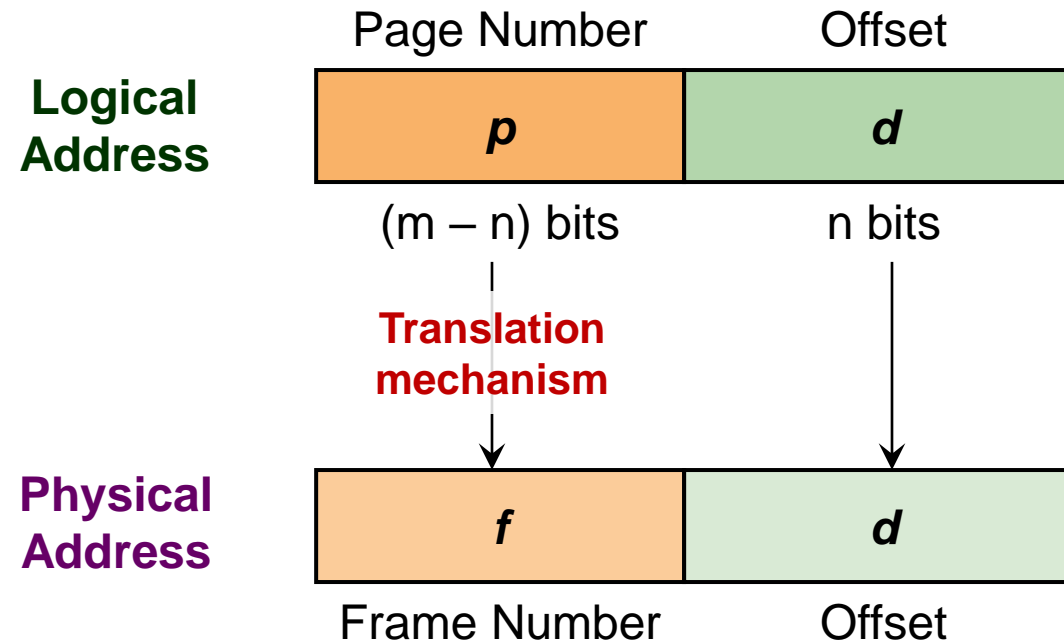


Logical Address Translation

- Program code uses **logical memory address**
 - However, to actually access the value, **physical memory address** is needed
- To locate a value in physical memory in the paging scheme, we need to know:
 - **F**, the physical frame number
 - **Offset**, displacement from the beginning of the physical frame
- **Physical Address**
= F x size of physical frame + Offset

Logical Address Translation: Essential Tricks

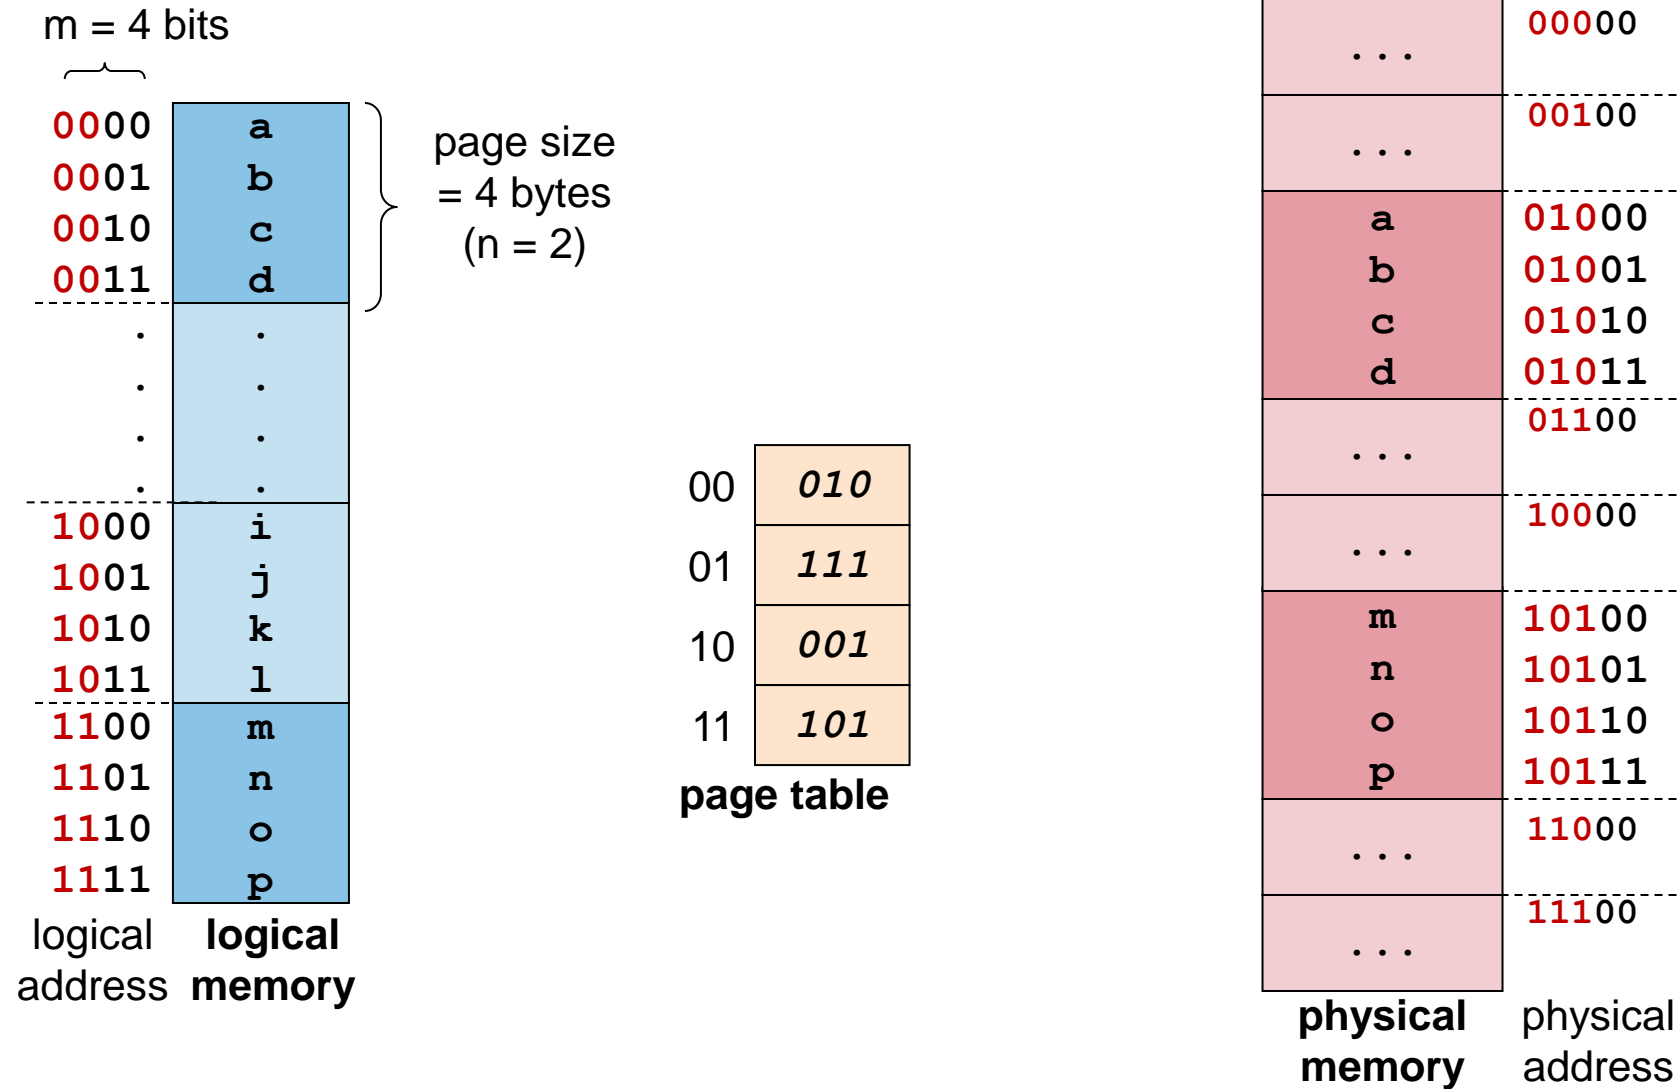
- Two important design decisions simplify address translation calculation
 1. Keep frame size (page size) as a power-of-2
 2. Physical frame size == Logical page size



Logical Address Translation: Formula

- Given:
 - Page/Frame size of 2^n
 - m bits of logical address
- Logical Address LA :
 - p = Most significant $m-n$ bits of LA
 - d = Remaining n bits of LA
- Use p to find frame number f
 - from mapping mechanism like page table
- Physical Address PA :
 - $PA = f * 2^n + d$

Example: 4 Logical Pages, 8 Physical Frames



Paging: Observations

- Paging removes **external fragmentation**
 - ❑ No left-over physical memory region
 - ❑ All free frame can be used with no wastage
- Paging can still have **internal fragmentation**
 - ❑ Logical memory space may not be multiple of page size
- Clear separation of logical and physical address space
 - ❑ Allow great flexibility
 - ❑ Simple address translation

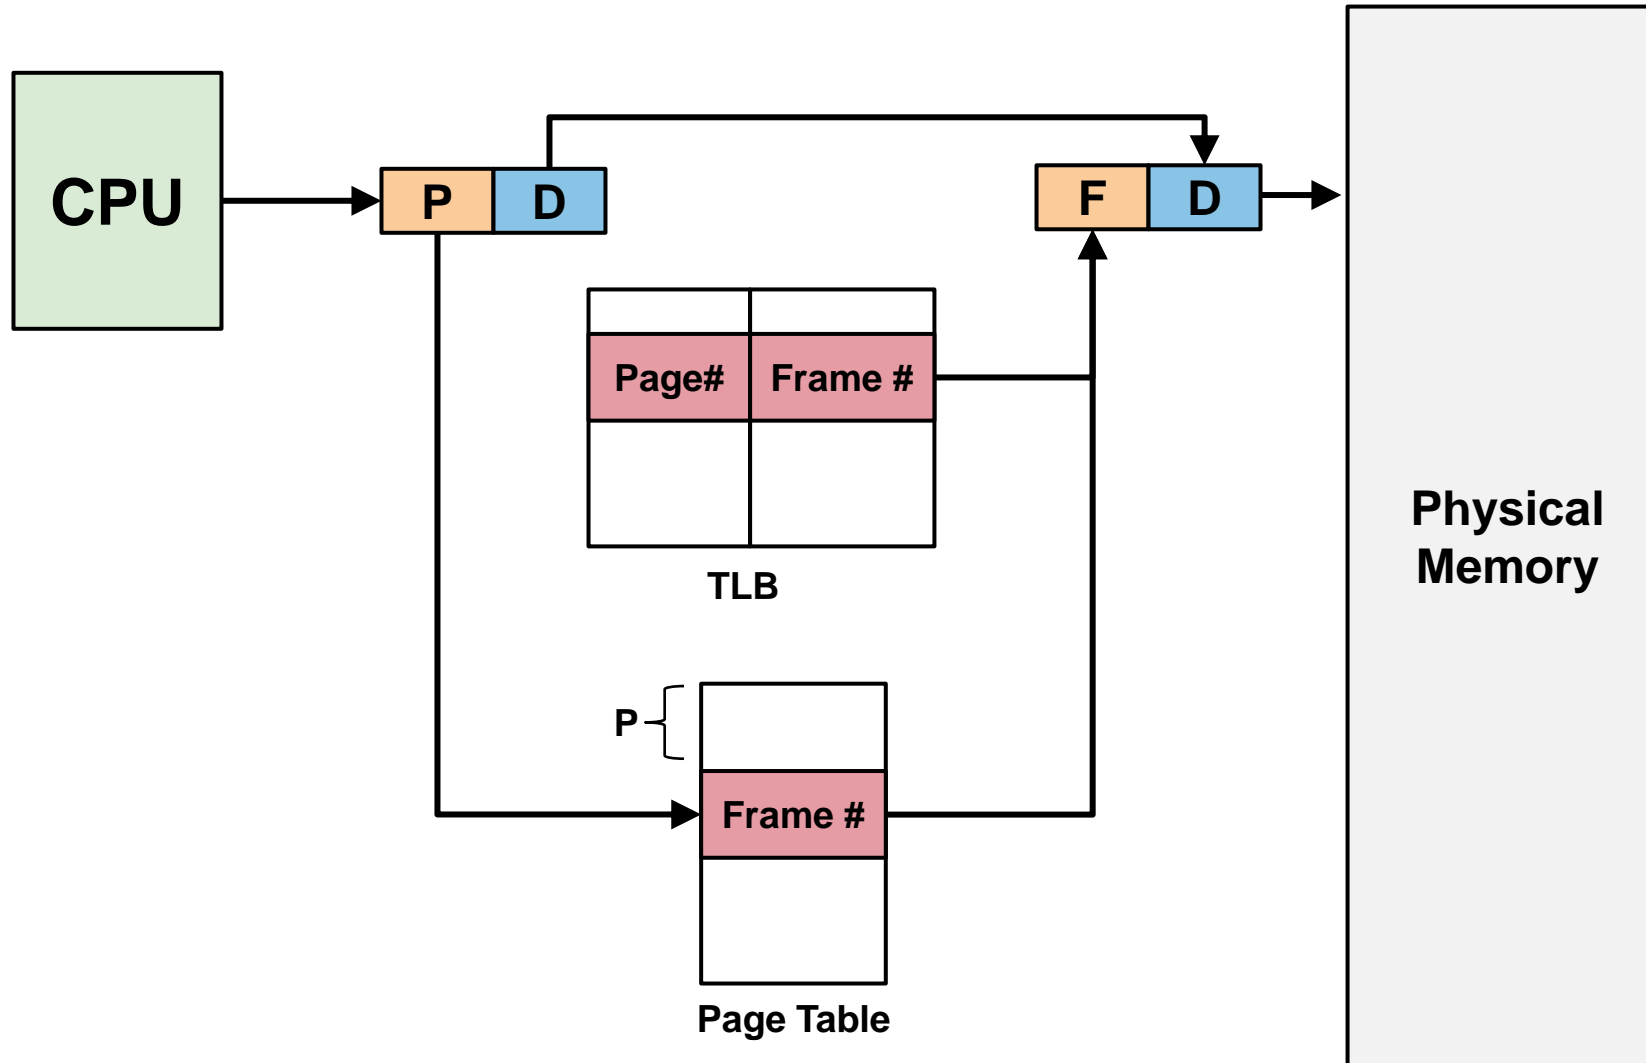
Implementing Paging Scheme

- Common pure-software solution:
 - OS stores page table alongside process information (e.g., PCB)
 - **Improved understanding:** Memory context of a process == Page table
- Issues:
 - Require two memory accesses for every memory reference
 - 1st access is to read the indexed page table entry to get frame number
 - 2nd access is to access the actual memory item

Paging Scheme: Hardware Support

- Modern processors provide specialized **on-chip** component to support paging
 - Known as **Translation Look-Aside Buffer (TLB)**
 - TLB acts as a **cache** for a **few** page table entries
- Logical address translation with TLB:
 - Use **page number** to search TLB associatively
 - Entry found (**TLB-Hit**):
 - Frame number is retrieved to **generate physical address**
 - Entry not found (**TLB-Miss**):
 - Access the full page table
 - Retrieved frame number is used to **generate physical address** and **update TLB**

Translation Look-Aside Buffer: Illustration



TLB: Impact on Memory Access Time

- Suppose:

- TLB access takes 1ns
- Main memory access takes 50ns
- What is the average memory access time if TLB contains 40% of the whole page table?

- Memory access time

= TLB hit + TLB miss

= 40% x (1ns + 50ns) + 60%(1ns + 50ns + 50ns)

= 81ns

- Note:

- Overhead of filling in TLB entry and impact of cache ignored.

TLB and Process Switching

- **Improved Understanding:** TLB is part of the hardware context of a process
- When a context switch occurs:
 - TLB entries are flushed
 - So that new process will not get incorrect translation
- Hence, when a process resumes running:
 - Will encounter many TLB misses to fill the TLB
 - It is possible to place some entries initially, e.g., the code pages, to reduce TLB misses

Paging Scheme: Protection

- The basic paging scheme can be easily extended to provide memory protection between processes using:
 - ❑ Access-Right Bits
 - ❑ Valid Bit
- Access Right Bits:
 - ❑ Each Page Table Entry includes several bits to indicate:
 - Whether the page is writable, readable, executable
 - E.g., page containing code should be executable, page containing data should be readable and writable, etc.
 - ❑ Memory access is checked against the access right bits

Paging Scheme: Protection (cont)

■ Observation:

- ❑ The logical memory range is usually the same for **all** processes
- ❑ However, not all processes utilize the whole range
 - ➔ Some pages are out-of-range for a particular process

■ Valid bit:

- ❑ Attached to each page table entry to indicate:
 - Whether the page is valid for the process to access
- ❑ OS will set the valid bits when a process is running
- ❑ Memory access is checked against this bit:
 - Out-of-range access will be caught by OS

Paging Scheme: Page sharing

- Page table can allow several processes to share the same physical memory frame
 - Use the same physical frame number in the page table entries
- Possible usage:
 - Shared code page:
 - Some code are being used by many processes, e.g., C standard library, system calls, etc.
 - Implement Copy-On-Write:
 - As discussed earlier in "Process Abstraction" lecture, parent and child process can share a page until one tries to change a value in it

Paging Scheme: **Page sharing**

Page 0
Page 1
Page 2
Page 3

Process P's
logical
memory

Page 0
Page 1
Page 2
Page 3

Process Q's
logical
memory

0	2
1	7
2	1
3	5

Process P's
page table

0	
1	
2	
3	

Process Q's
page table

frame 0	Page 1
1	Page 2
2	Page 0
3	Page 2
4	Page 0
5	Page 3
6	Page 3
frame 7	Page 1

**physical
memory**

Paging Scheme: **Copy-On-Write**

Page 0
Page 1
Page 2
Page 3

Process P's
logical
memory

0	2
1	7
2	1
3	5

Process P's
page table

Page 0
Page 1
Page 2
Page 3

Process Q's
logical
memory

0	
1	
2	
3	

Process Q's
page table

frame 0	
1	Page 2
2	Page 0
3	
4	
5	Page 3
6	
frame 7	Page 1

**physical
memory**

Why memory error is often called ***segmentation fault***?

SEGMENTATION SCHEME

Segmentation Scheme: **Motivation**

- **Memory space** of a process is treated as a single entity so far
 - ❑ However, there are several **memory regions** with different usage in a process
- For example, in a typical C program:
 - ❑ User Code Region
 - ❑ Global Variables Region
 - Static Data (persistent as long as program executes)
 - ❑ Heap Region:
 - Dynamic data (persistent as long as user doesn't free)
 - ❑ Stack Region
 - ❑ Library Code Region
 - ❑ etc.

Segmentation Scheme: Motivation (cont)

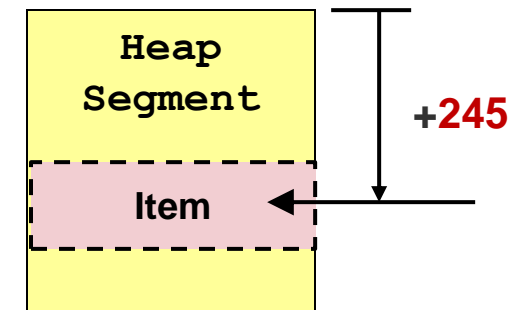
- Some regions may grow/shrink at execution time:
 - E.g., the stack region, the heap region, the library code region
- It is hard to place different regions in a contiguous memory space and still allow them grow/shrink freely
 - Also hard to check whether a memory access in a region is within its bound

Segmentation Scheme: Basic Idea

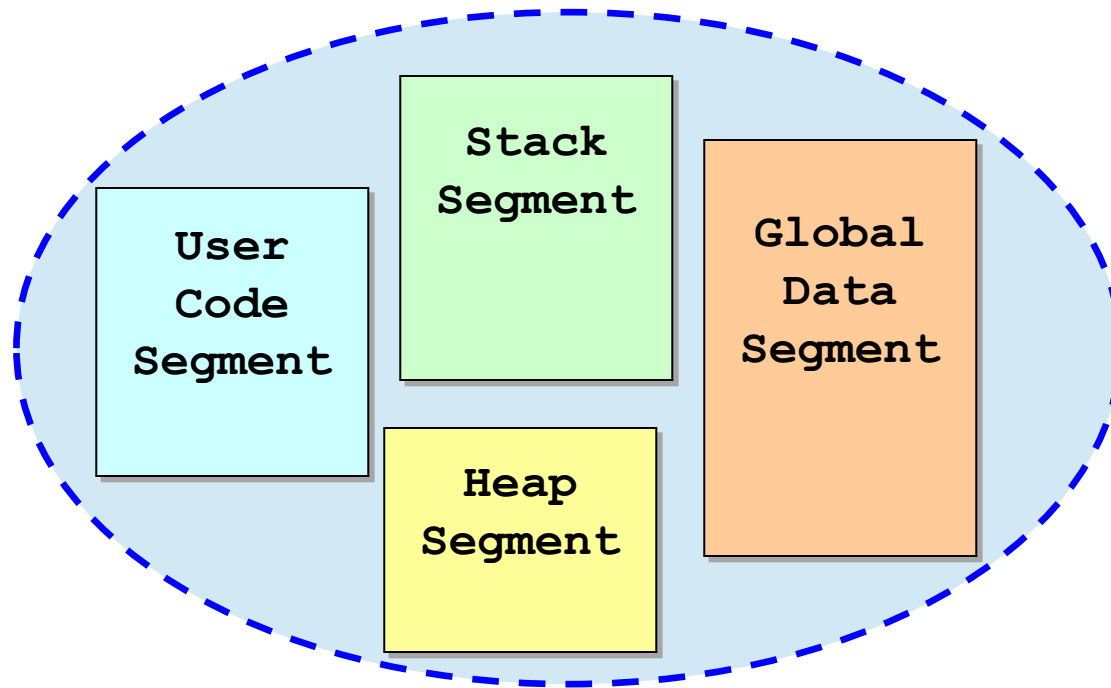
- Separate regions into multiple **memory segments**
 - Logical memory space of a process is now a collection of **segments**
- Each memory segment:
 - Has a **name**
 - For ease of reference
 - Has a **limit**
 - Indicate the range of the segment
- All **memory reference** is now specified as:
 - **Segment name + Offset**

Memory Access
Example:

"*Heap*" + 245

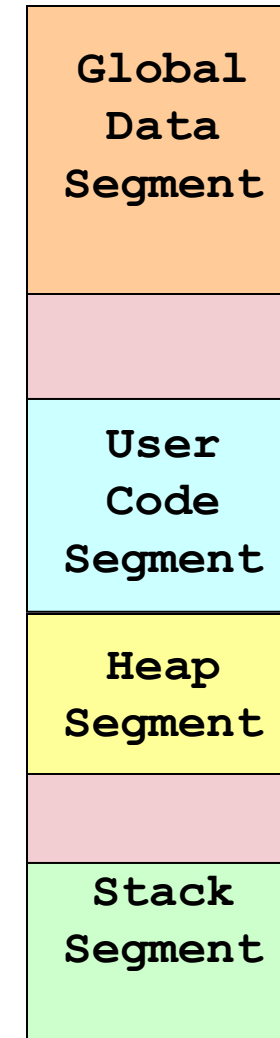


Segmentation: Illustration



Logical Address Space

(Segments can be of different sizes)



Physical Memory

(Segments are disjoint in RAM)

Segmentation: Logical Address Translation

- Each **segment** is mapped to a **contiguous physical memory region**
 - with a **base address** and a **limit**
- The **segmentation name** is usually represented as a **single number**
 - Known as **segment id**
- Logical address **<SegID, Offset>**:
 - **SegID** is used to look up **<Base, Limit>** of the segment in a **segment table**
 - Physical Address $PA = Base + Offset$
 - **Offset < Limit** for valid access

Logical Address Translation: Illustration

Assume:

User Code Segment = 0

Global Data Segment = 1

Heap Segment = 2

Stack Segment = 3

Memory Access

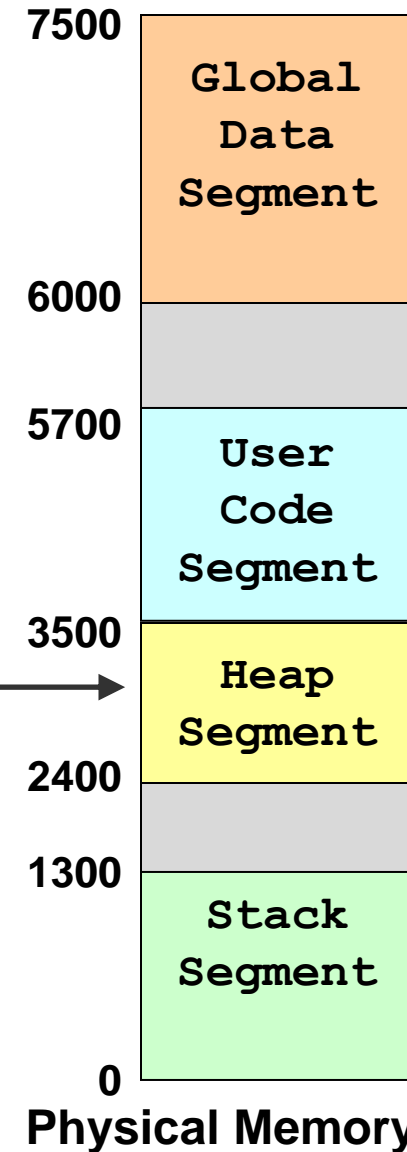
< Segment Id, Offset >

	Base	Limit
0	3500	2200
1	6000	1500
2	2400	1100
3	0	1300

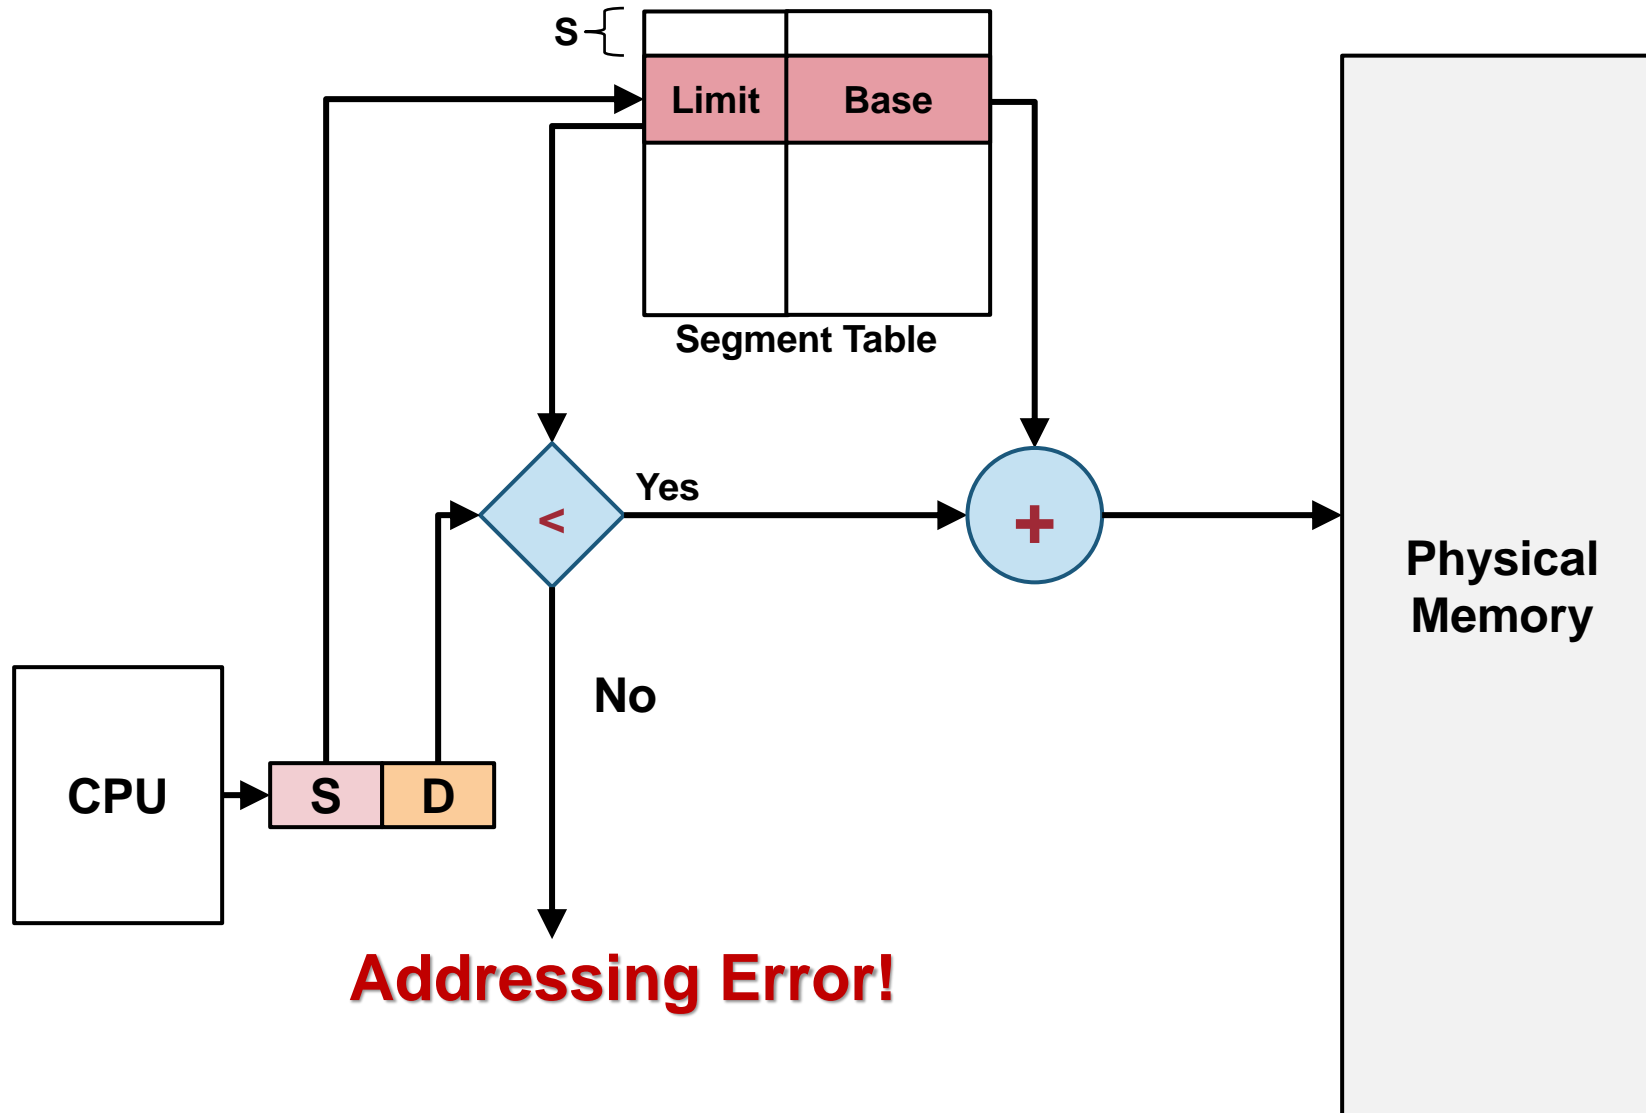
segment table

< 2, 500 >

2900



Segmentation: Hardware Support



Segmentation: Summary

■ Pros:

- ❑ Each segment is an **independent contiguous** memory space
 - ➔ Can grow / shrink independently
 - ➔ Can be protected / shared independently

■ Cons:

- ❑ Segmentation requires **variable-size contiguous** memory regions
 - ➔ can cause **external fragmentation**

■ Important observation:

- ❑ Segmentation is **not** the same as paging
- ❑ Each of them trying to solve a different problem

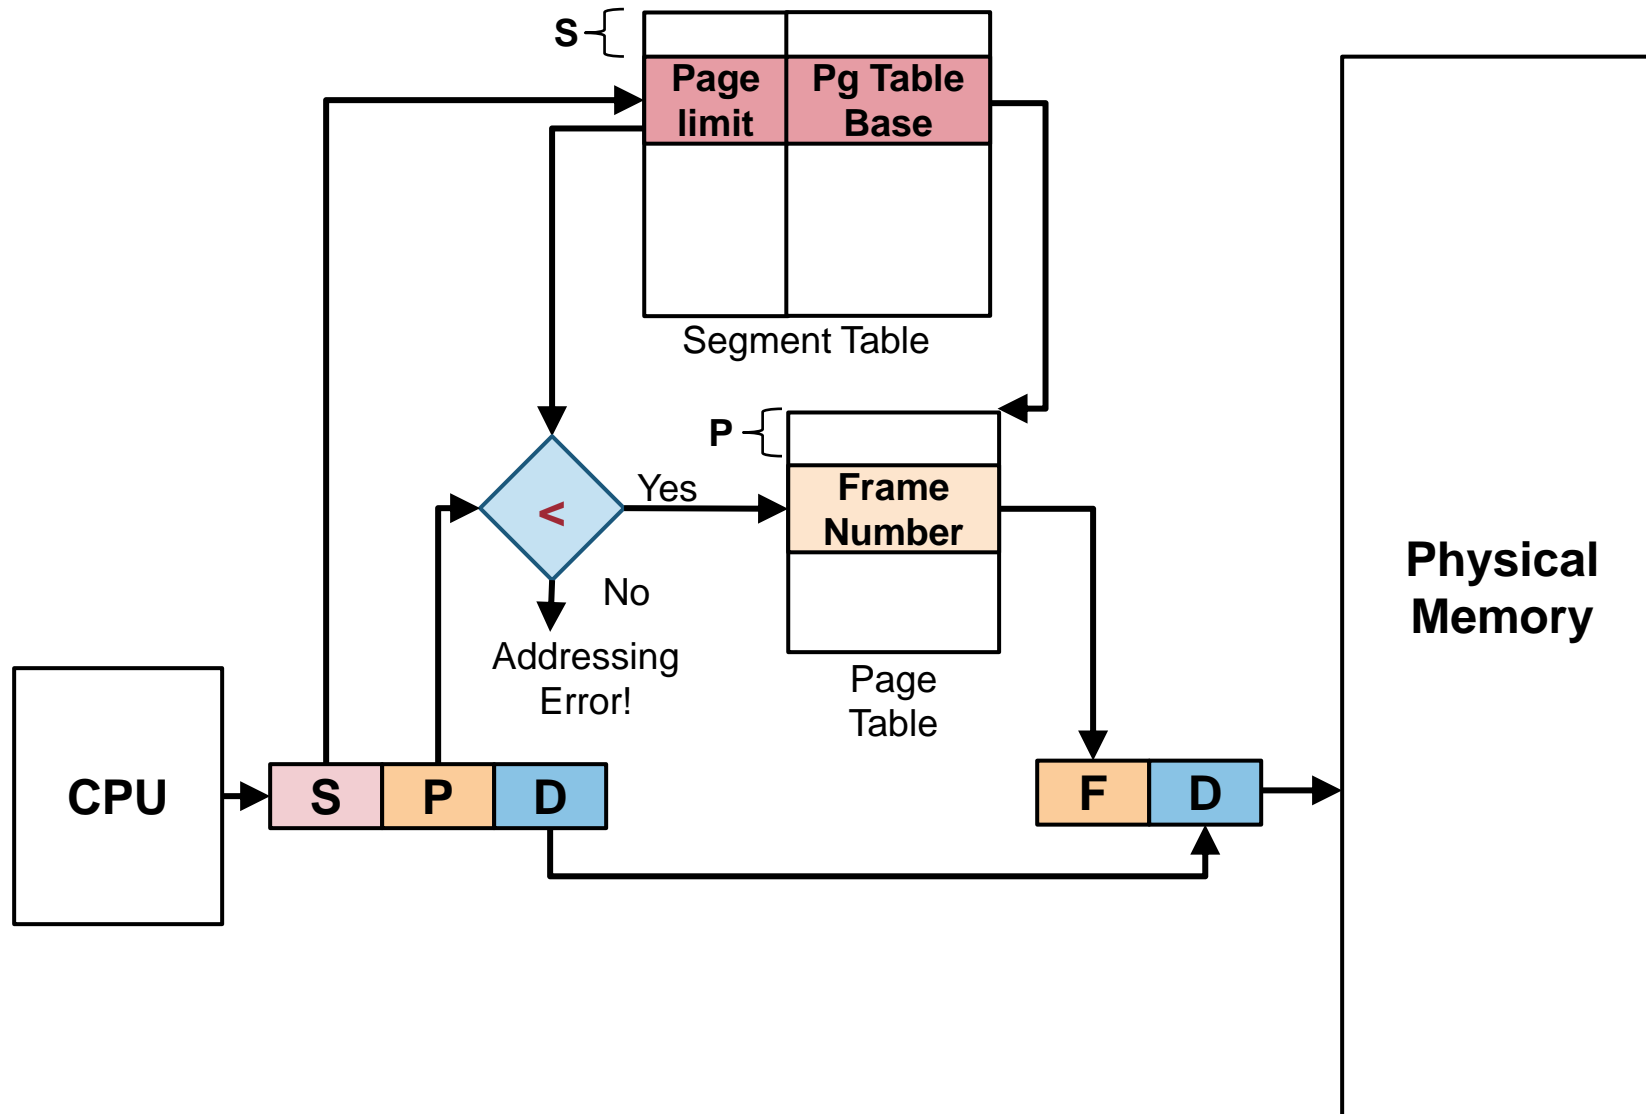
P is good; S is also good; Let's do S+P then!

SEGMENTATION WITH PAGING

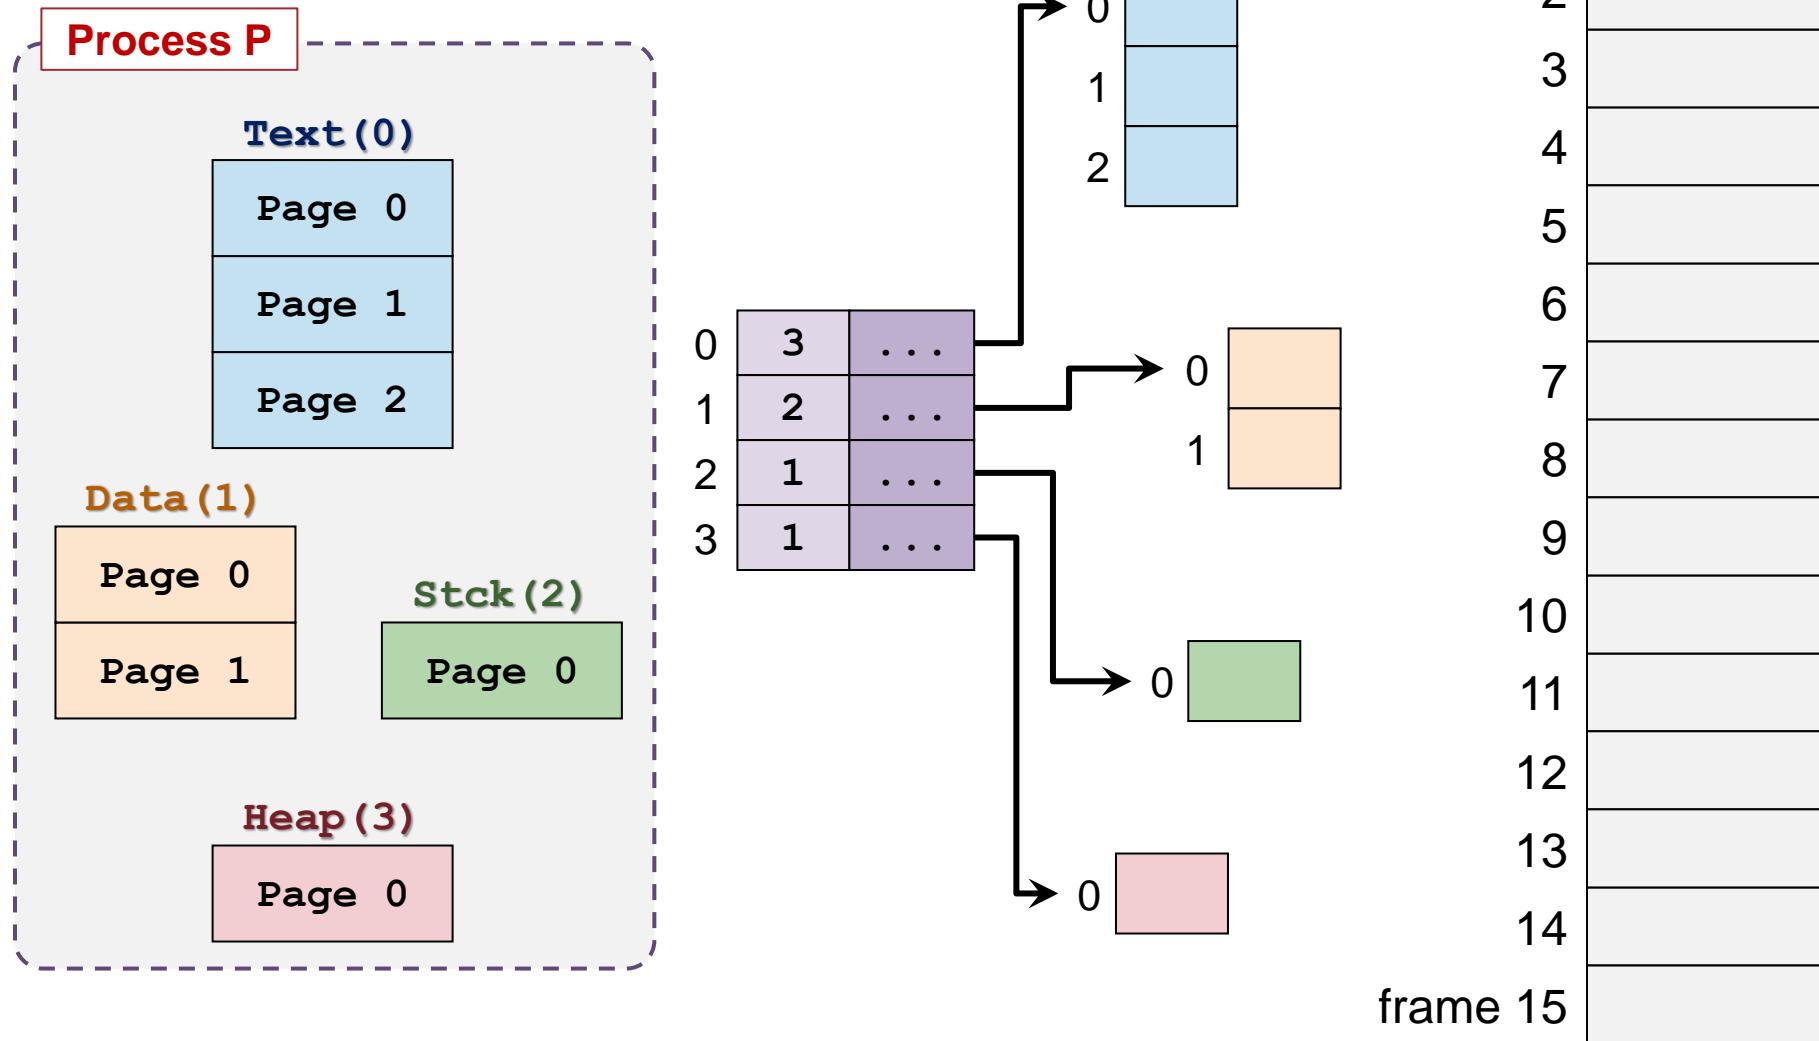
Segmentation with Paging: Basic Idea

- The intuitive next step is to combine segmentation and paging
- Basic Idea:
 - ❑ Each segment is now composed of several pages instead of a contiguous memory region
 - Essentially, each segment has a page table
 - ❑ Segment can grow by allocating new pages then add to its page table, and similarly for shrinking

Segmentation with Paging: Illustration



Segmentation with Paging



Summary

- Discussed two popular memory management schemes
 - Both allow the logical address space to be in disjoint physical regions
- Paging split the logical address into fixed size pages
 - Store in fixed size physical memory frames
- Segmentation split the logical address into variable size segments according to their usage
 - Stored in variable-sized physical partitions
- Combining segmentation and paging