

---

Process Management

# Process Abstraction

---

Lecture 2

# Overview

- Program execution:
  - Hardware Context
  - Memory Context
    - Code & Data
    - Function call
    - Dynamically allocated memory
- Introduction to Process Management
  - OS Context
    - Process State
  - Process Control Block and Process Table
- OS interaction with Process

# Recap: C Sample Program and Assembly Code

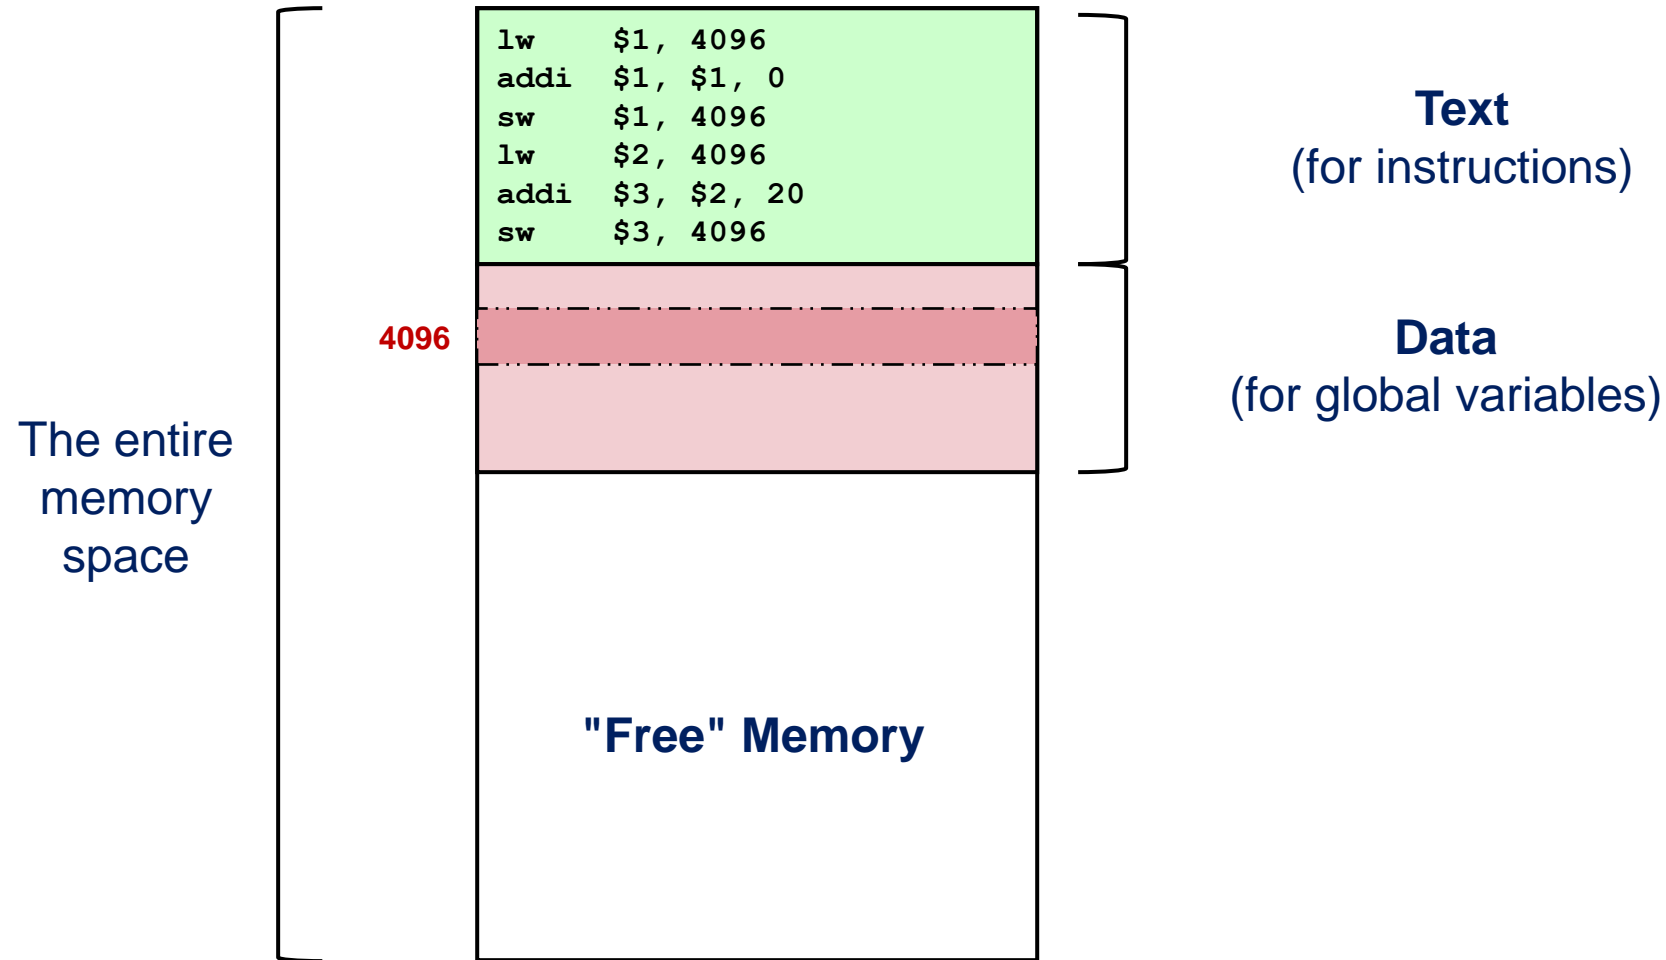
```
int i = 0;  
  
i = i + 20;
```

C Code Fragment

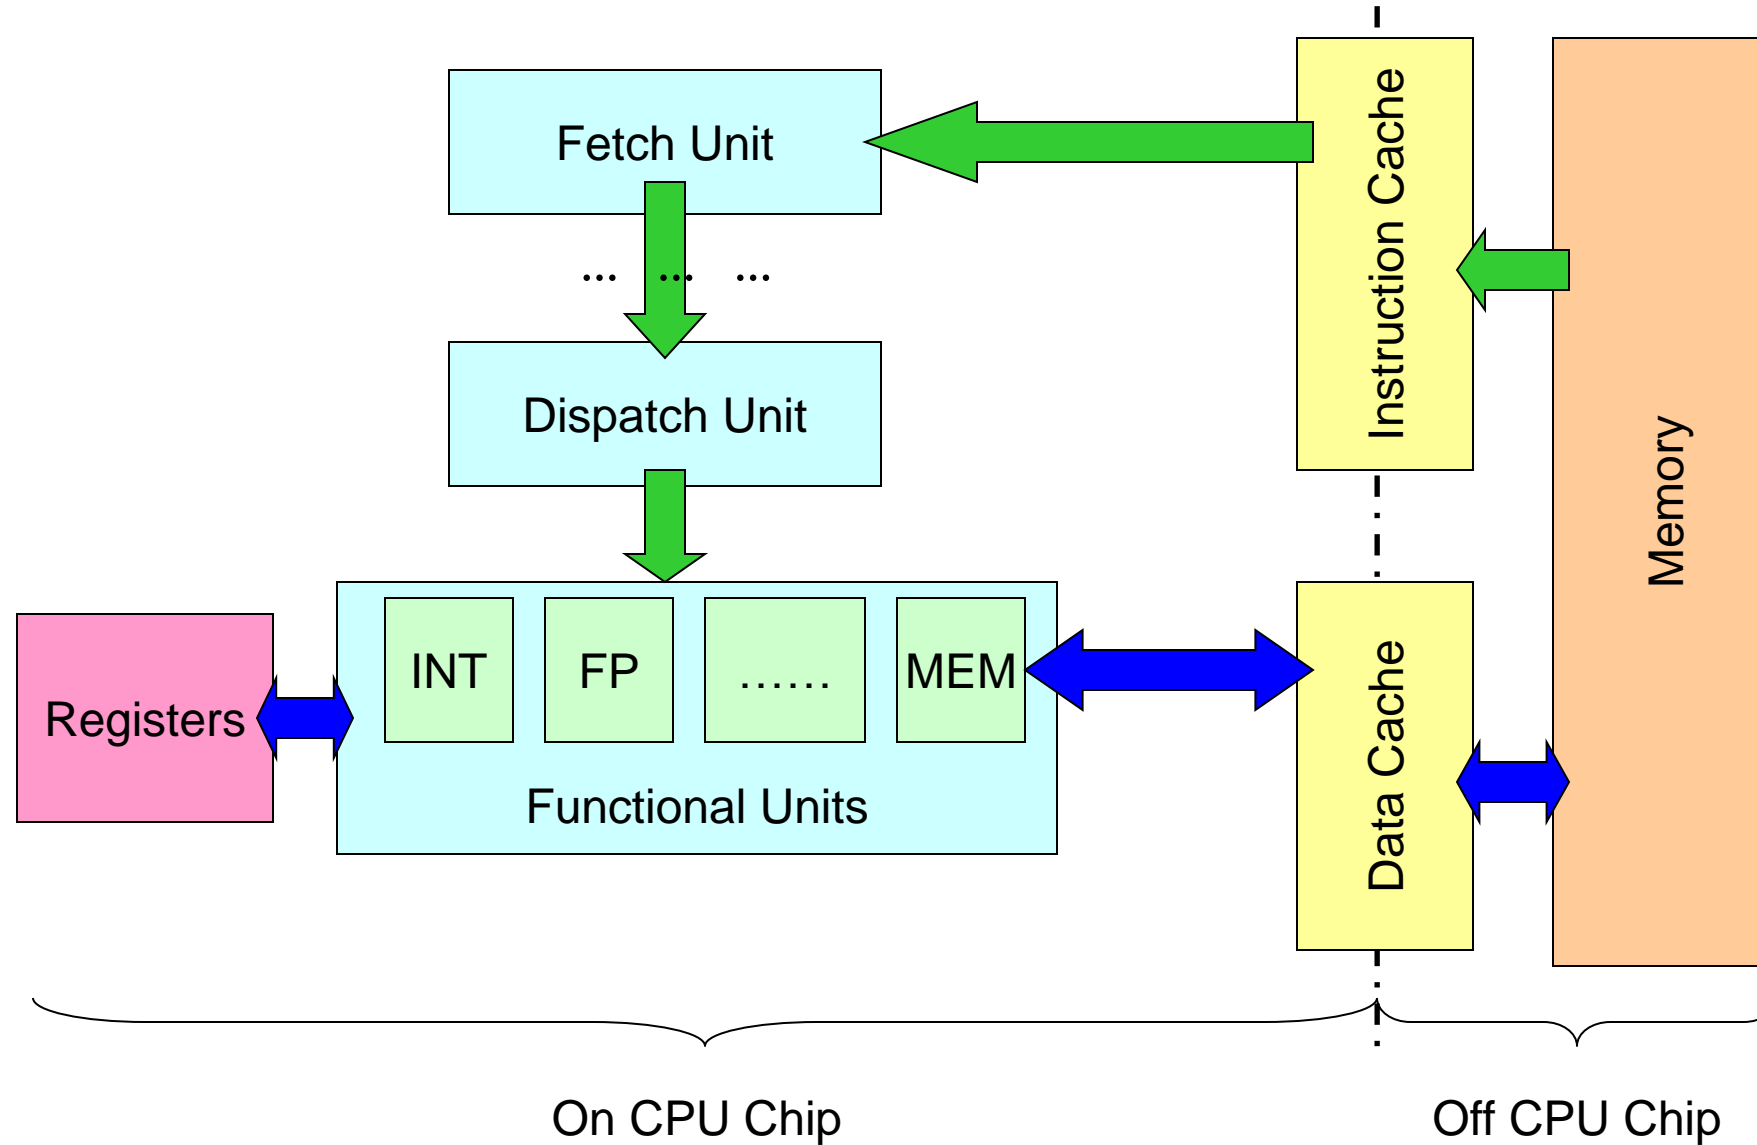
```
lw      $1, 4096      //Assume address of i = 4096  
addi    $1, $0, 0     //register $1 = 0  
sw      $1, 4096      //i = 0  
  
lw      $2, 4096      //read i  
addi    $3, $2, 20     //$3 = $2 + 20  
sw      $3, 4096      //i = i + 20
```

Corresponding MIPS-like Assembly Code

# Recap: Program Execution (Memory)



# Recap: Generic Computer Organization



# Recap: Component Description

## ■ Memory:

- ❑ Storage for instruction and data

## ■ Cache:

- ❑ Duplicate part of the memory for faster access
- ❑ Usually split into instruction cache and data cache

## ■ Fetch unit:

- ❑ Loads instruction from memory
- ❑ Location indicated by a special register: **P**rogram **C**ounter (PC)

# Recap: Component Description (cont)

## ■ Functional units:

- ❑ Carry out the instruction execution
- ❑ Dedicated to different instruction type

## ■ Registers:

- ❑ Internal storage for the fastest access speed
- ❑ General Purpose Register (GPR):
  - Accessible by user program (i.e. visible to compiler)
- ❑ Special Register:
  - Program Counter (PC)
  - Stack Pointer (SP)
  - Frame Pointer (FP)
  - etc.

# Recap: Basic Instruction Execution

- Instruction X is fetched
  - Memory location indicated by **P**rogram **C**ounter
- Instruction X dispatched to the corresponding Functional Unit
  - Read operands if applicable
    - Usually from memory or GPR
  - Result computed
  - Write value if applicable
    - Usually to memory or GPR
- Instruction X is completed
  - PC updated for the next instruction



## Recap: What you should know 😊

- An executable (binary) consists of two major components:
  - Instructions and Data
- When a program is **under execution**, there are **more information**:
  - Memory context:
    - **Text** and **Data**, ...
  - Hardware context:
    - **General purpose registers, Program Counter**, ...
- Actually, there are **other types of memory usage** during program execution
  - Coming up next

Memory Context

# Function Call

---

What if  $\mathbf{f}()$  calls  $\mathbf{u}()$  calls  $\mathbf{n}()$  ?

# Function Call : Challenges

```
int i = 0;
```

```
i = i + 20;
```

C Code Fragment

VS

```
int g(int i, int j)
{
    int a;

    a = i + j;
    return a;
}
```

C Code with Function

## ■ Consider:

- ❑ How do we allocate memory space for variables **i**, **j** and **a**?
  - Can we just make use of the "**data**" memory space?
- ❑ What are the key issues?

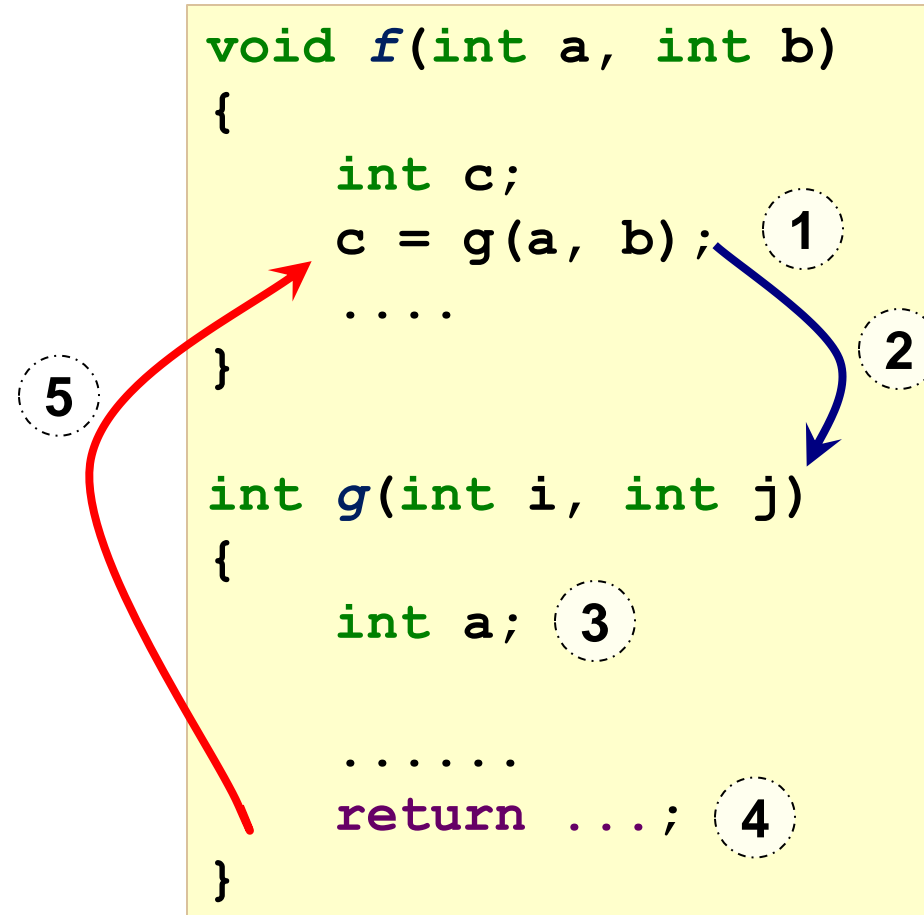
# Function Call : Control Flow and Data

- **f ()** calls **g ()**

- **f ()** is the **caller**
- **g ()** is the **callee**

- Important Steps:

1. Setup the parameters
2. Transfer control to callee
3. Setup local variable
4. Store result if applicable
5. Return to caller



# Function Call : Control Flow and Data

## ■ Control Flow Issues:

- ❑ Need to jump to the function body
- ❑ Need to resume when the function call is done
- ➔ Minimally, need to store the PC of the caller

## ■ Data Storage Issues:

- ❑ Need to pass parameters to the function
- ❑ Need to capture the return result
- ❑ May have local variables declaration

➔ Need a **new region of memory** that dynamically used by function invocations

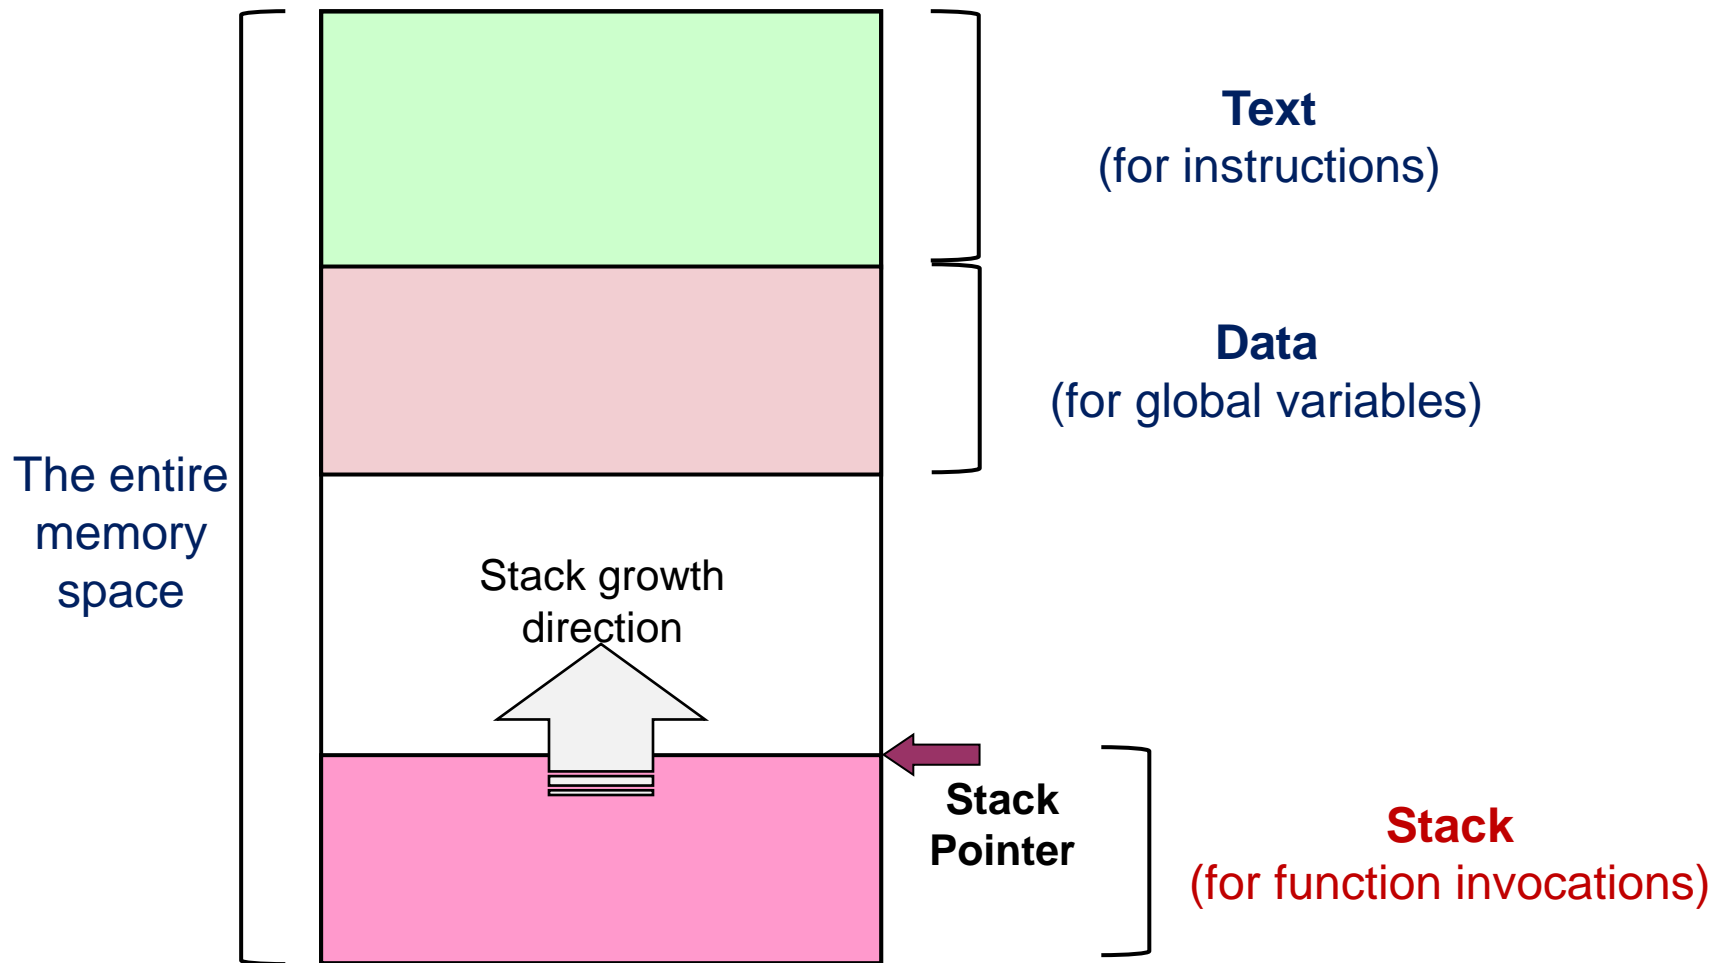
# Introducing Stack Memory

- **Stack Memory Region:**
  - The new memory region to store information for function invocation
- Information of a function invocation is described by a **stack frame**
- Stack frame contains:
  - Return address of the caller
  - Arguments (Parameters) for the function
  - Storage for local variables
  - Other information.... (more later)

# Stack Pointer

- The top of stack region (first unused location) is logically indicated by a **Stack Pointer**:
  - Most CPU has a specialized register for this purpose
  - Stack frame is added on top when a function is invoked
    - Stack “grows”
  - Stack frame is removed from top when a function call ends
    - Stack “shrinks”

# Illustration: Stack Memory





# Illustration: Stack Memory Usage (1 / 5)

```
void f()
```

```
{
```

```
...
```

```
g();
```

```
...
```

```
}
```

```
void g()
```

```
{
```

```
h();
```

```
...
```

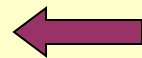
```
}
```

```
void h()
```

```
{
```

```
...
```

```
}
```



At this  
point

Stack Frame  
for f()

...


...

# Illustration: Stack Memory Usage (2 / 5)

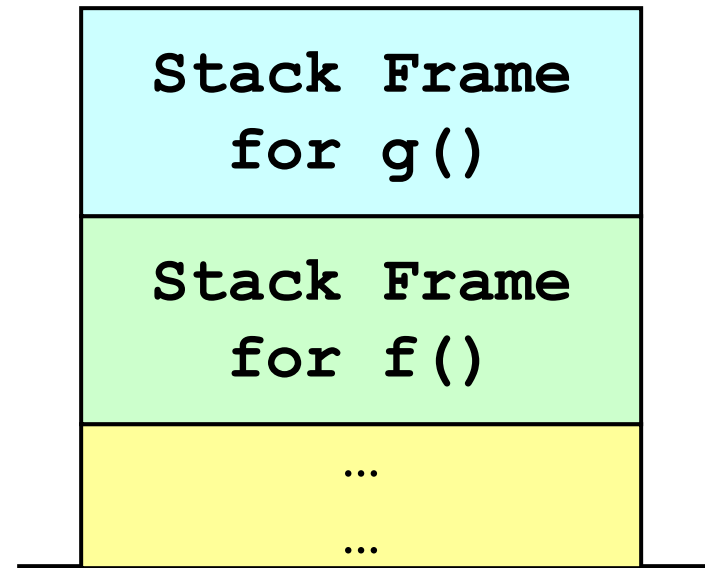
```
void f()  
{  
    ...  
    g();  
    ...  
}
```

```
void g()  
{  
    h();  
    ...  
}
```

```
void h()  
{  
    ...  
}
```



At this  
point



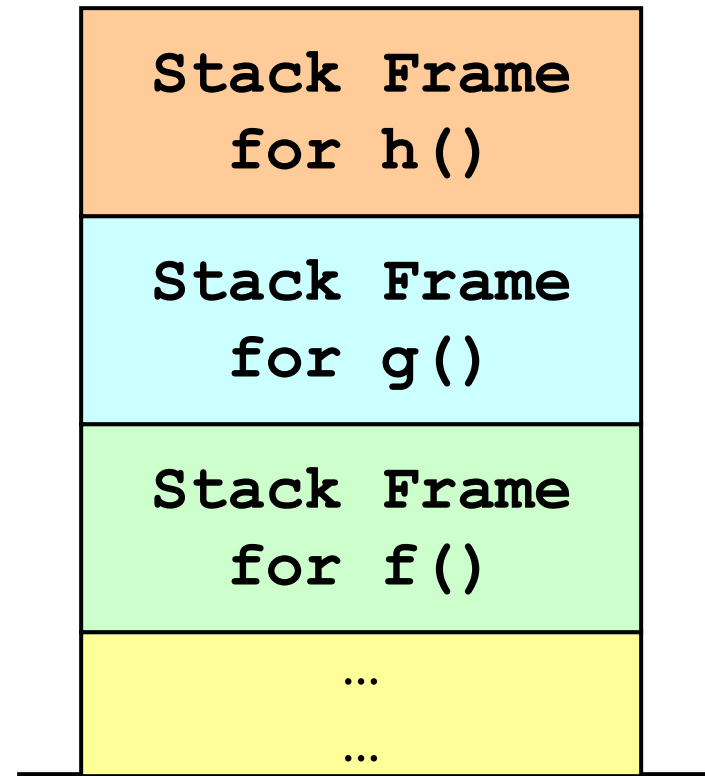
# Illustration: Stack Memory Usage (3 / 5)

```
void f()  
{  
    ...  
    g();  
    ...  
}
```

```
void g()  
{  
    h();  
    ...  
}
```

```
void h()  
{  
    ...  
}
```

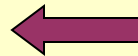
At this  
point



# Illustration: Stack Memory Usage (4 / 5)

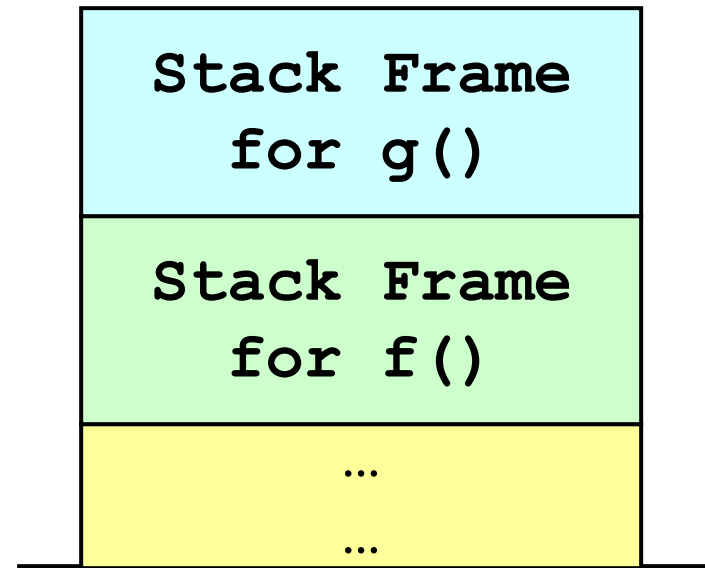
```
void f()  
{  
    ...  
    g();  
    ...  
}
```

```
void g()  
{  
    h();  
    ...  
}
```



At this  
point

```
void h()  
{  
    ...  
}
```



# Illustration: Stack Memory Usage (5 / 5)

```
void f()
```


```
{
```

```
...
```

```
g();
```

```
...
```

```
}
```



At this  
point

```
void g()
```

```
{
```

```
h();
```

```
...
```

```
}
```

```
void h()
```

```
{
```

```
...
```

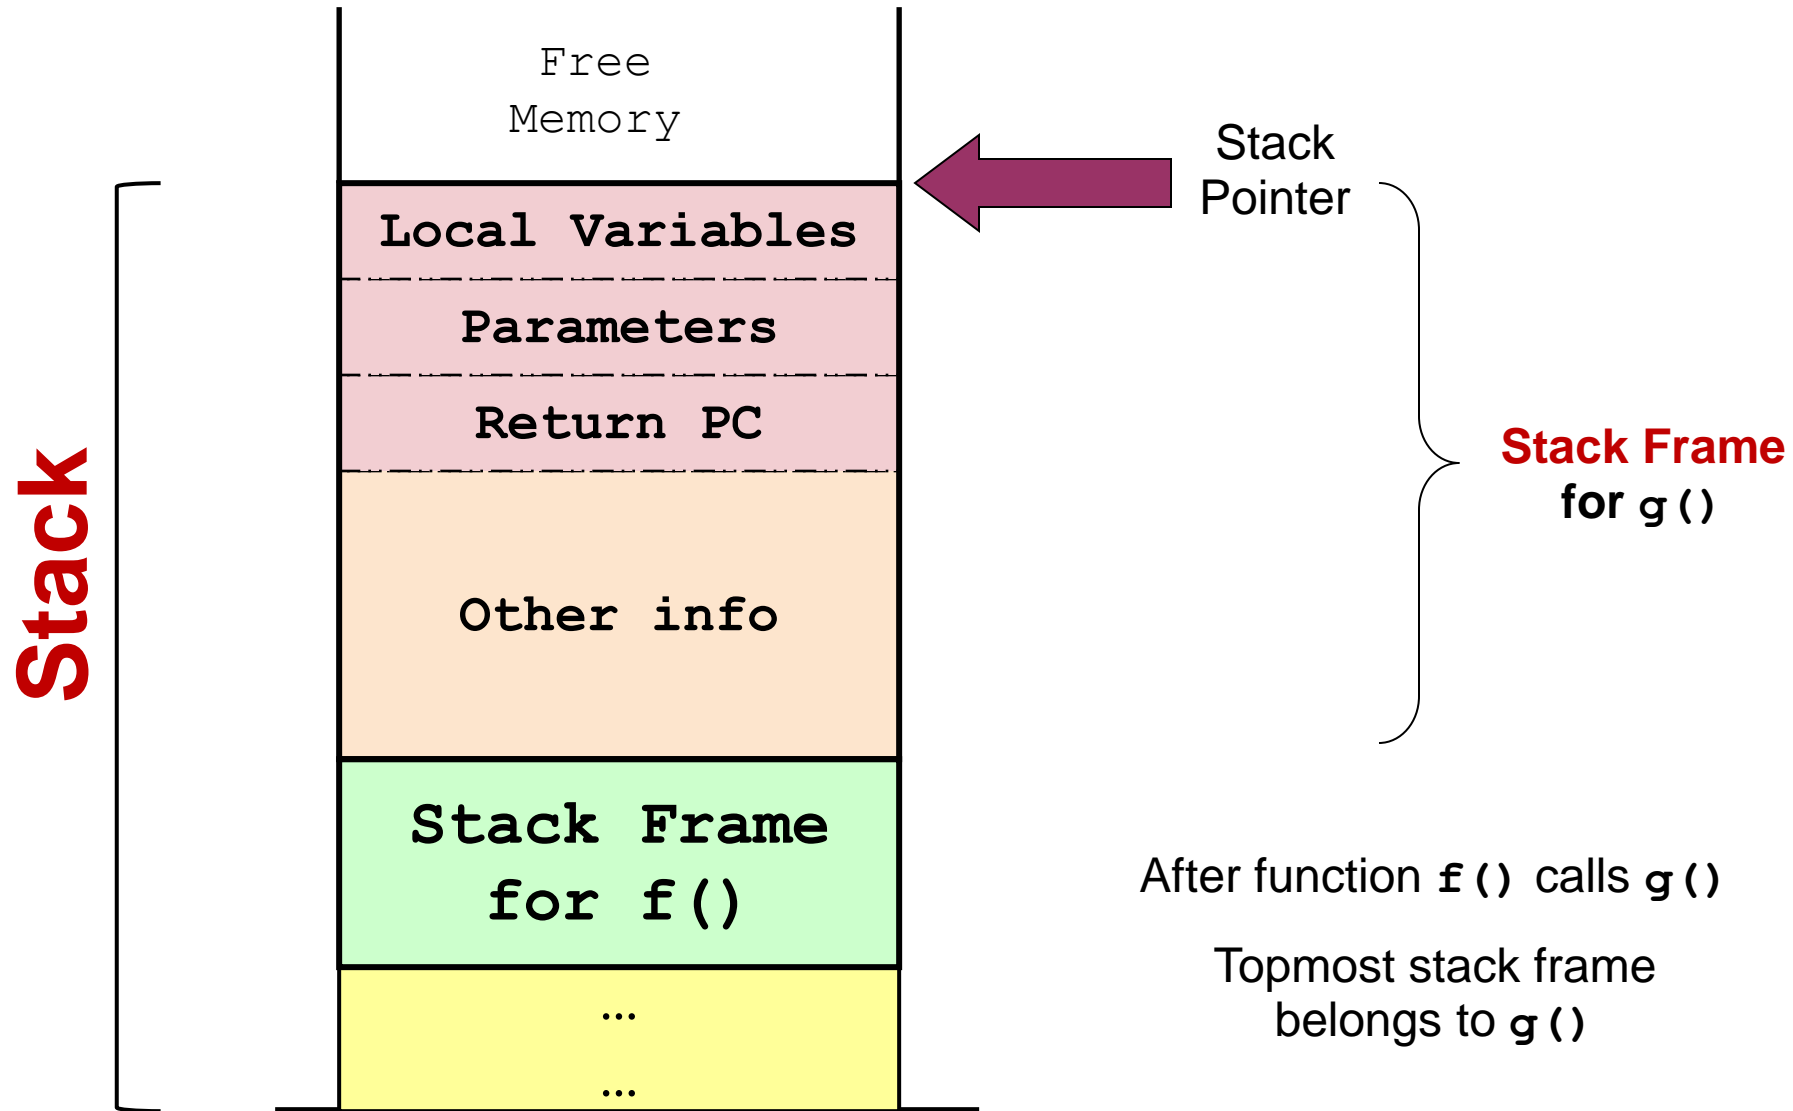
```
}
```

Stack Frame  
for f()

...

...

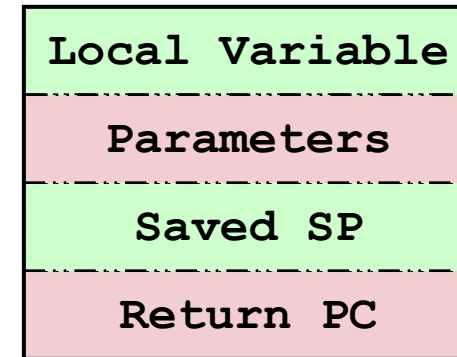
# Illustration: Stack Frame v1.0



# Function Call Convention

- Different ways to setup stack frame:
  - Known as **function call convention**
  - Main differences:
    - What information is stored in stack frame vs. in registers?
    - Which portion of stack frame is prepared by caller / callee?
    - Which portion of stack frame is cleared by caller / callee?
    - Who between caller / callee to adjust the stack pointer?
- No universal way
  - Hardware and programming language dependent
- An example scheme is described next

# Stack Frame Setup



## ■ Prepare to make a function call:

- ❑ **Caller:** Pass parameters using registers and/or stack
- ❑ **Caller:** Save Return PC on stack

## ❑ **Transfer Control from Caller to Callee**

- ❑ **Callee:** Save the old Stack Pointer (SP)
- ❑ **Callee:** Allocate space for local variables of callee on stack
- ❑ **Callee:** Adjust SP to point to new stack top



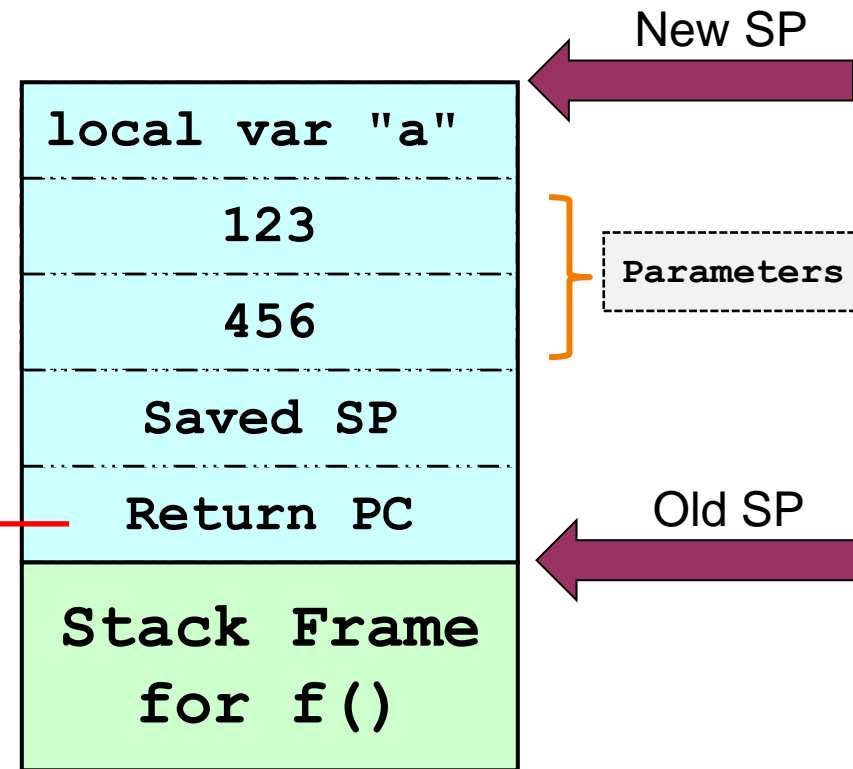
# Illustration: Calling function *g* ( )

```
void f(int a, int b)
{
    int c;

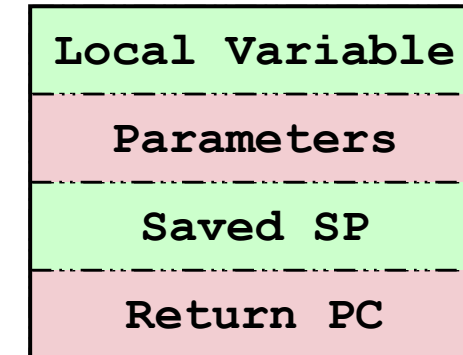
    a = 123;
    b = 456;
    c = g(a, b);
    ....
}

int g(int i, int j)
{
    int a;

    a = i + j;
    return a * 2;
}
```



# Stack Frame Teardown



## ■ On returning from function call:

- ❑ **Callee:** Place return result in register (if applicable)

- ❑ **Callee:** Restore saved Stack Pointer

- ❑ Transfer control back to caller using saved PC

- ❑ **Caller:** Utilize return result (if applicable)

- ❑ **Caller:** Continues execution in caller

# Illustration: Function *g* ( ) finishes

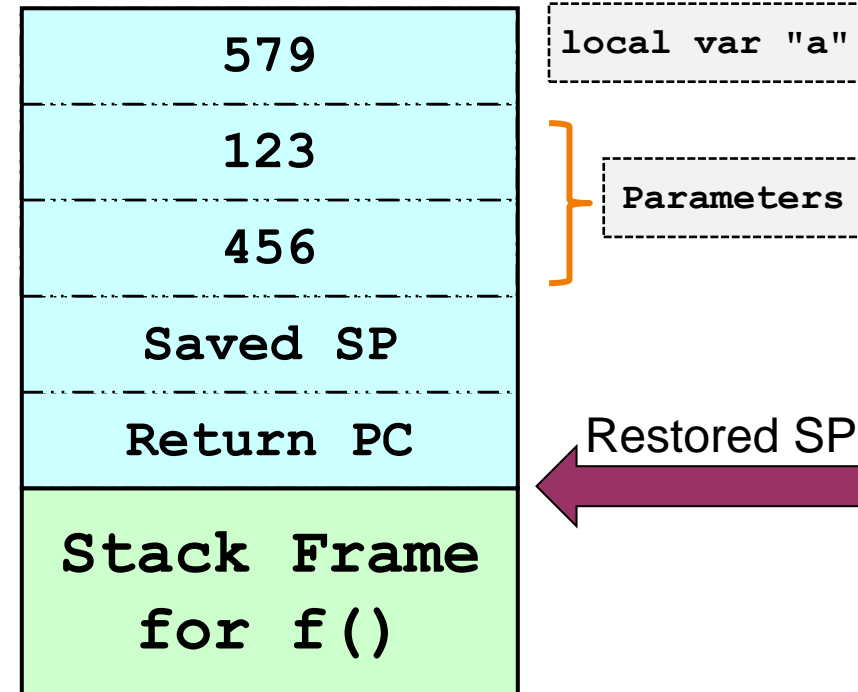
```
void f(int a, int b)
{
    int c;

    a = 123;
    b = 456;
    c = g(a, b);
    ....
}
```

```
int g(int i, int j)
{
    int a;

    a = i + j;
    return a * 2;
}
```

Execution  
resumes here



reg \$2

1158

# Other Information in Stack Frame

- We have described the basic idea of:
  - Stack frame
  - Calling Convention: Setup and Teardown
- A few common additional information in the stack frame:
  - Frame Pointer
  - Saved Registers

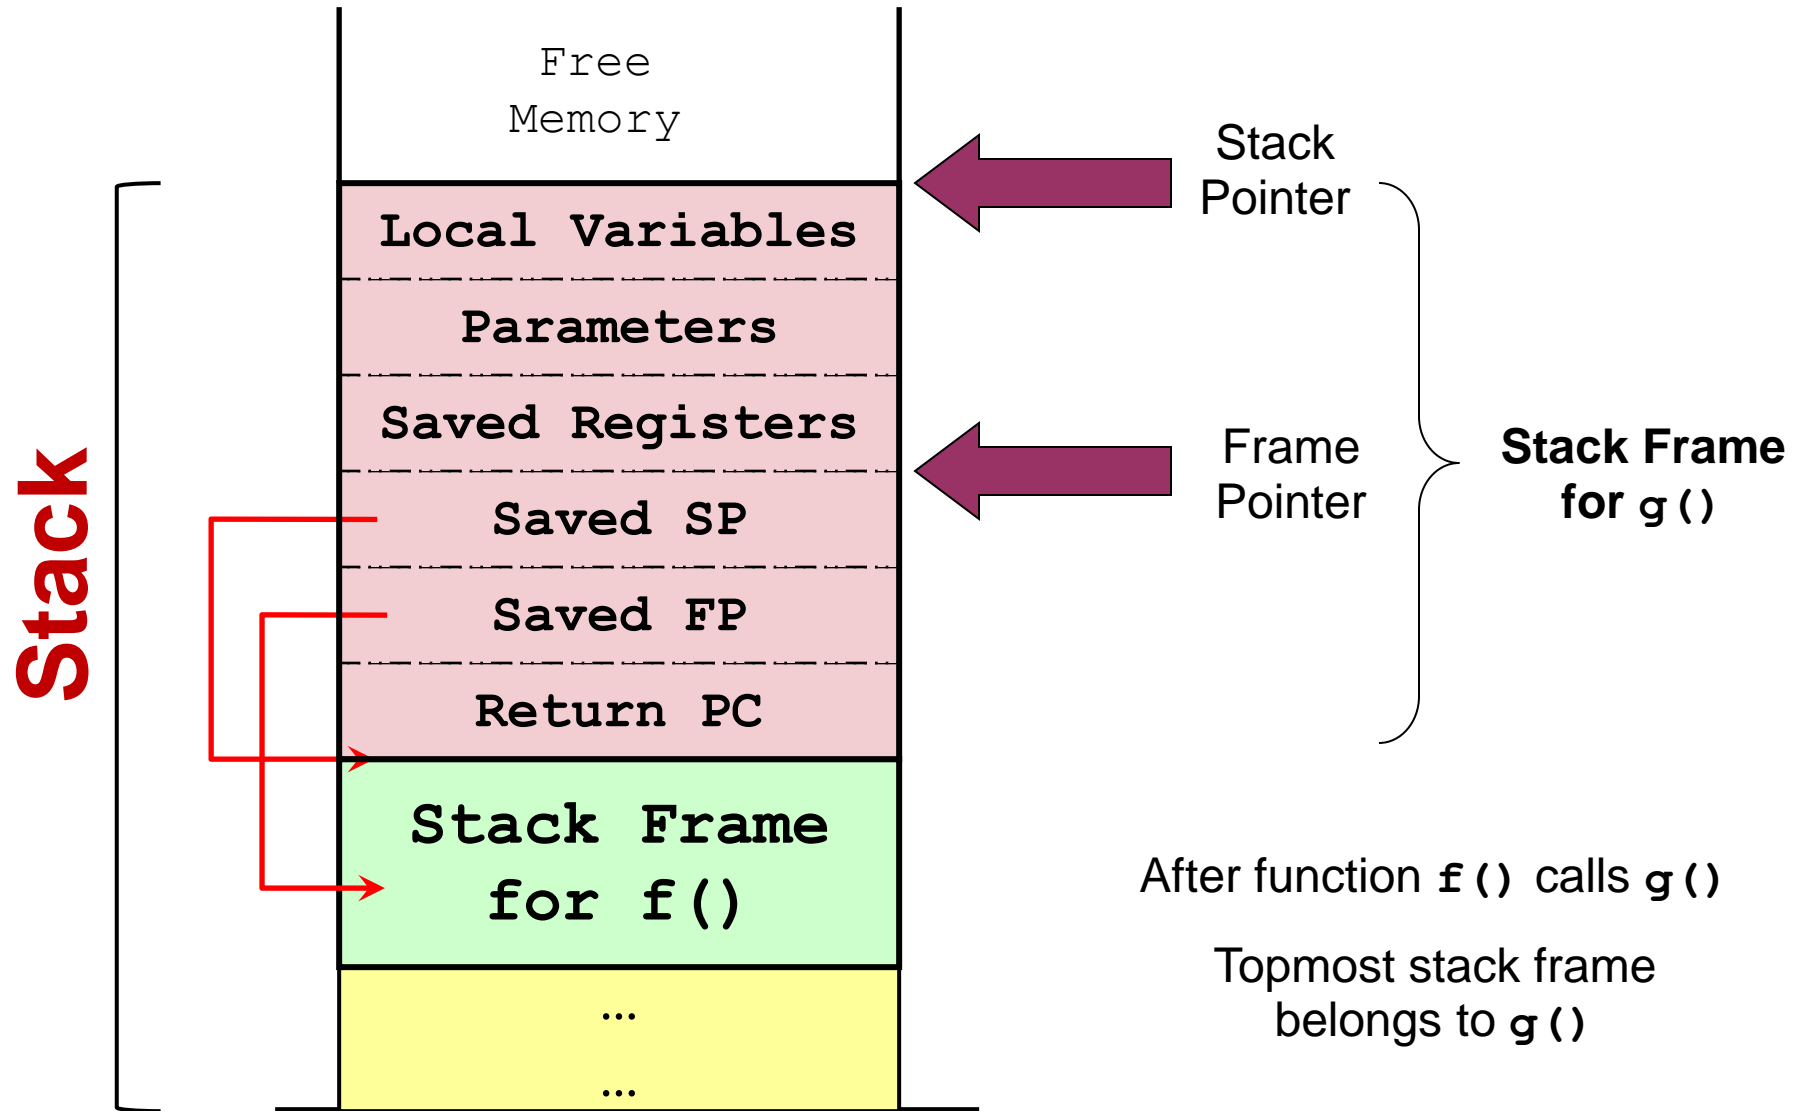
# Frame Pointer

- To facilitate the access of various stack frame items:
  - Stack Pointer is hard to use as it can change
  - ➔ Some processors provide a dedicated register ***Frame Pointer***
- The frame pointer points to a fixed location in a stack frame
  - Other items are accessed as a displacement from the frame pointer
- The usage of FP is platform dependent

# Saved Registers

- The number of general purpose register (GPR) on most processors are very limited:
  - E.g. MIPS has 32 GPRs, x86 has 16 GPRs
- When GPRs are exhausted:
  - Use memory to temporary hold the GPR value
  - That GPR can then be reused for other purpose
  - The GPR value can be restored afterwards
  - known as **register spilling**
- Similarly, a function can spill the registers it intend to use before the function starts (callee-saved)
  - Restore those registers at the end of function

# Illustration: Stack Frame v2.0



# Stack Frame Setup / Teardown [Updated]

## ■ On executing function call:

- ❑ **Caller:** Pass arguments with registers and/or stack
- ❑ **Caller:** Save Return PC on stack

### ❑ **Transfer control from caller to callee**

- ❑ **Callee:** Save registers used by callee. Save old FP, SP
- ❑ **Callee:** Allocate space for local variables of callee on stack
- ❑ **Callee:** Adjust SP to point to new stack top

## ■ On returning from function call:

- ❑ **Callee:** Restore saved registers, FP, SP
- ❑ Transfer control from callee to caller using saved PC
- ❑ **Caller:** Continues execution in caller

## ■ Remember, just an example!



# Function Call Summary

- In this part, we learned:
  - Another portion of memory space is used as a **Stack Memory**
  - Stack Memory stores the executing function using **Stack Frame**
    - Typical information stored on a stack frame
    - Typical scheme of setting up and tearing down a stack frame
  - The usage of Stack Pointer and Frame Pointer

Memory Context

# **Dynamically Allocated Memory**

---

Hmm... I need more memory

# Dynamically Allocated Memory

- Most programming languages allow dynamically allocated memory:
  - i.e. acquire memory space during **execution time**
- Examples:
  - In C, the ***malloc()*** function call
  - In C++, the ***new*** keyword
  - In Java, the ***new*** keyword
- Question:
  - Can we use the existing "Data" or "Stack" memory regions?

# Dynamically Allocated Memory

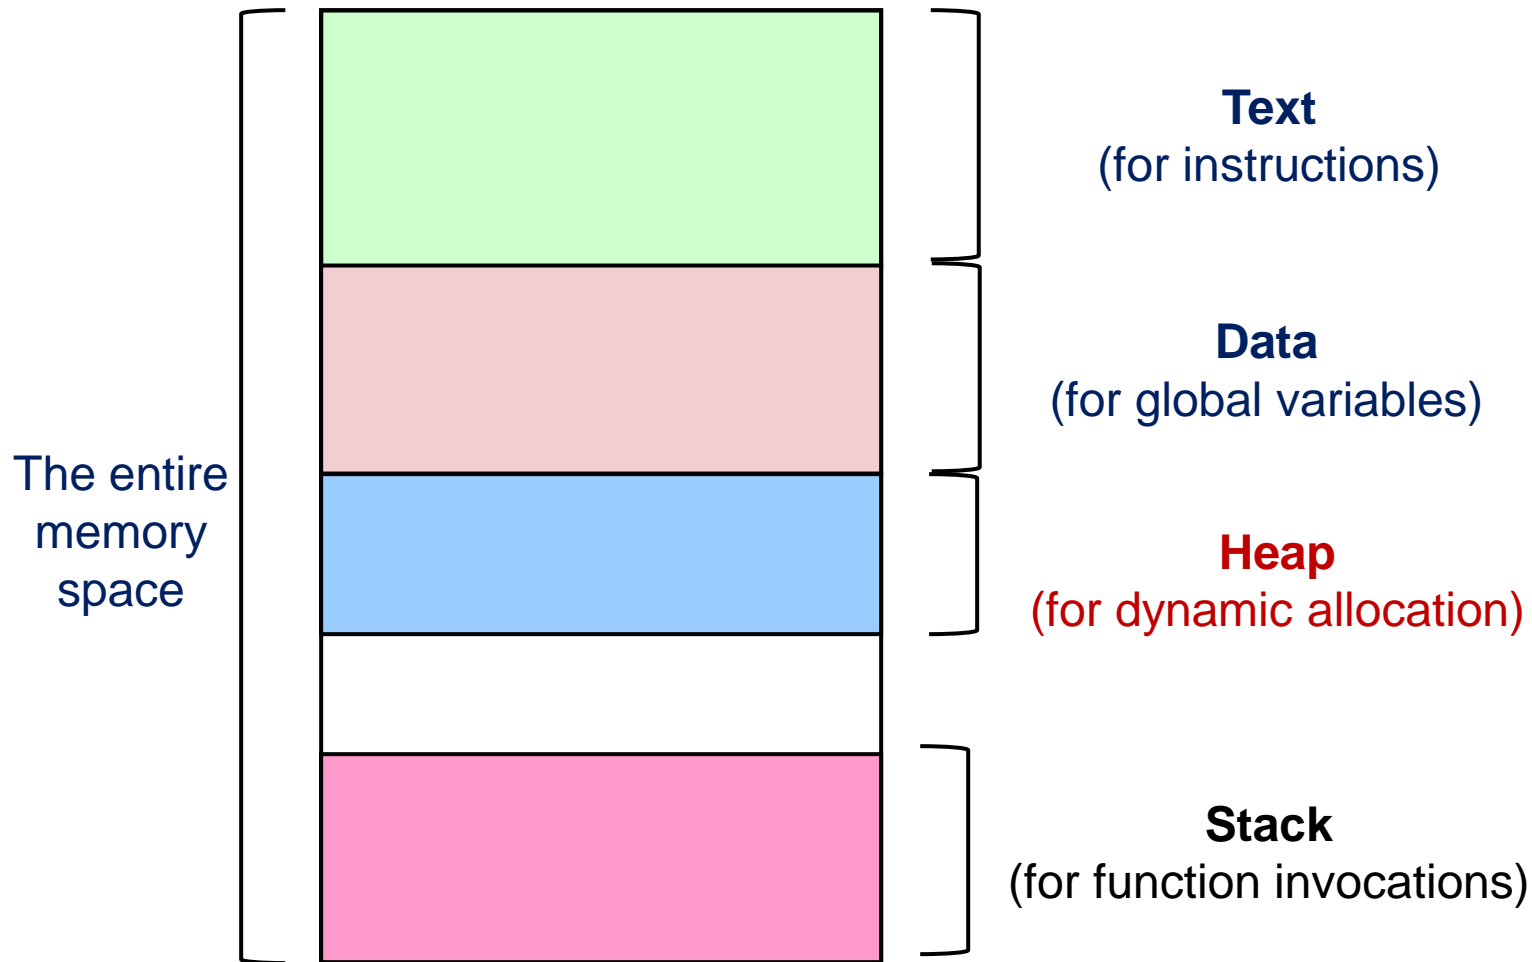
## ■ Observations:

- ❑ These memory blocks have different behaviors:
  1. Allocated only at runtime, i.e. size is not known during compilation time → Cannot place in **Data** region
  2. No definite deallocation timing, e.g. can be explicitly freed by programmer in C/C++, can be implicitly freed by garbage collector in Java → Cannot place in **Stack** region

## ■ Solution:

- ❑ Setup a separate **heap memory region**

# Illustration for Heap Memory



# Managing Heap Memory

- Heap memory is a lot trickier to manage due to its nature:
  - Variable size
  - Variable allocation / deallocation timing
- You can easily construct a scenario where heap memory are allocated /deallocated in such a way to create "holes" in the memory
  - Free memory block squeezed in between of occupied memory block
- We will learn more in the memory management (much) later in the course

# Checkpoint: Contexts updated

- Information describing a **program execution**:

- Memory context:

- Text, Data, **Stack** and **Heap**

- Hardware context:

- General purpose registers, Program Counter, **Stack pointer**, **Stack frame pointer**, ....

# Overview

- Program execution:
  - Hardware Context
  - Memory Context
    - Code & Data
    - Function call
    - Dynamically allocated memory
- Introduction to Process Management
  - OS Context
    - Process State
  - Process Control Block and Process Table
- OS interaction with Process



# Recap: Efficient Hardware Utilization

- OS should provide efficient use of the hardware resource:
  - By managing the programs executing on the hardware
- Observation:
  - If there is only **one program executing at any point in time**, how can we utilize hardware resources effectively?
    - Batch processing?
- Solution:
  - Allow **multiple programs to share the hardware**
    - e.g. Multiprogramming, Time-sharing

# Introduction to Process Management

- As the OS, to be able to switch from running program **A** to program **B** requires:
  1. Information regarding the execution of program **A** needs to be stored
  2. Program **A**'s information is replaced with the information required to run program **B**
- Hence, we need:
  - An **abstraction** to describe a running program
  - aka **process**

# Key Topics

## Process Abstraction

- Information describing an executing program

## Process Scheduling

- Deciding which process get to execute

## Inter-Process Communication & Synchronization

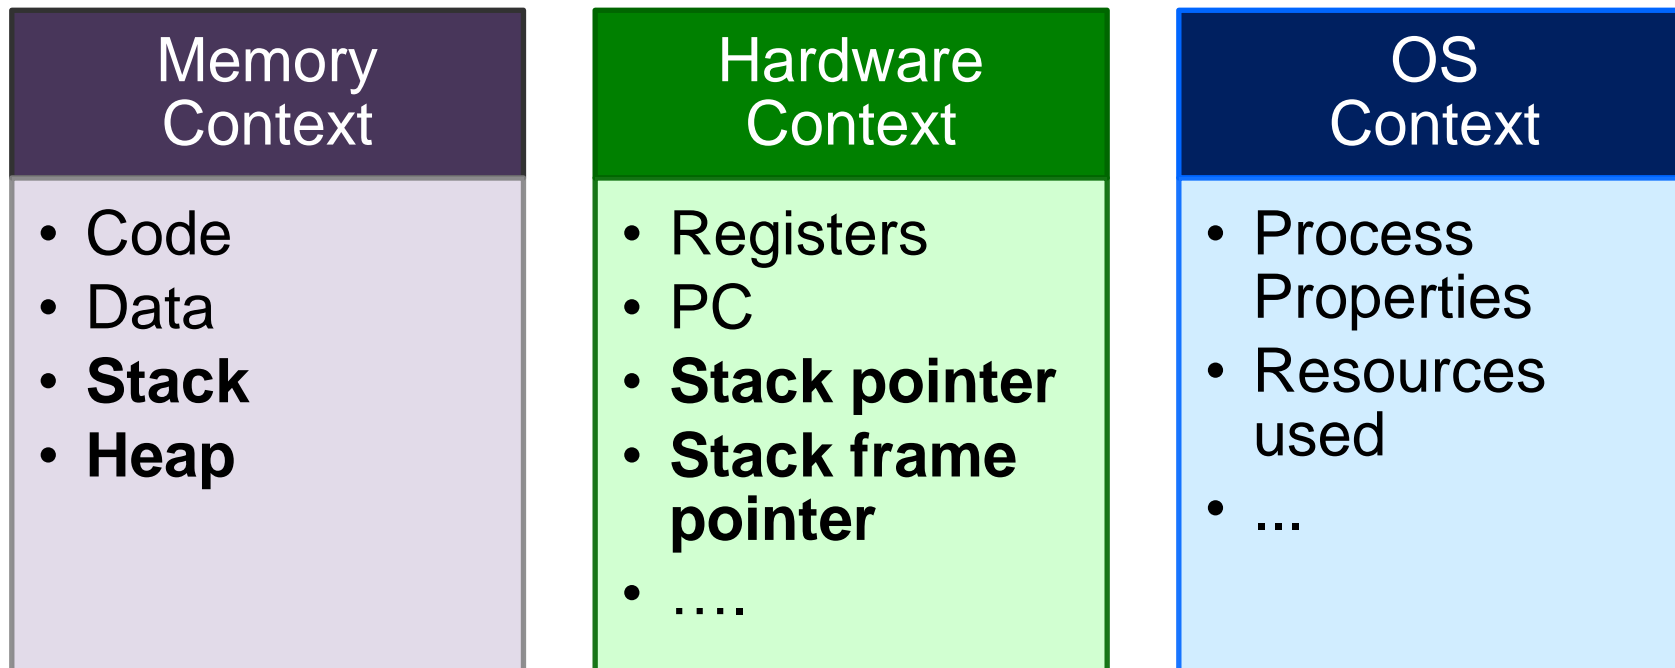
- Passing information between processes

## Alternative to Process

- Light-weight process aka Thread

# Process Abstraction

- **(Process / Task / Job)** is a dynamic abstraction for executing program
  - **information required** to describe a **running program**



OS Context

# Process Id & Process State

---

Your ID? Give me a status report!

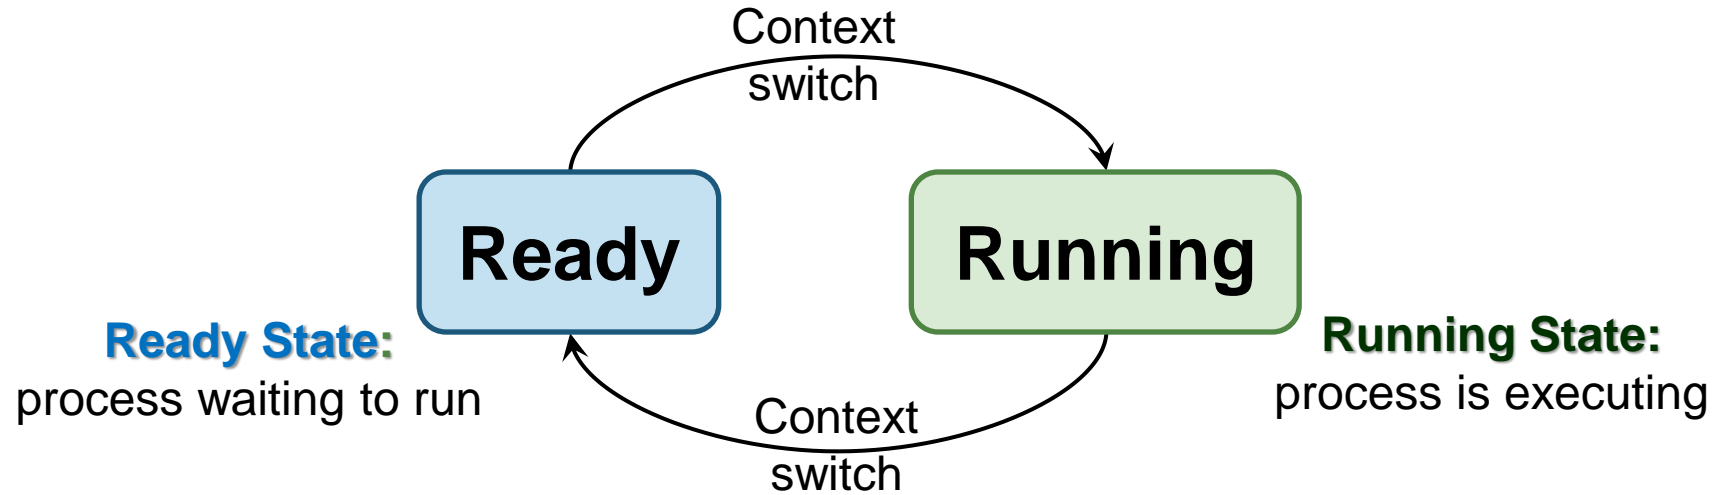
# Process Identification

- To distinguish processes from each other
  - ❑ Common approach is to use process ID (**PID**)
    - Just a number
  - ❑ Unique among processes
- There are a couple of OS dependent issues:
  - ❑ Are PIDs reused?
  - ❑ Does it limit the maximum no. of processes?
  - ❑ Are there reserved PIDs?

# Introducing Process State

- In the multitasking scenario:
  - A process can be:
    - Running OR
    - Not-running, e.g., another process running
- A process can be **ready to run**
  - But not actually executing
  - E.g., waiting for its turn to use the CPU
- Hence, each process should have a **process state**:
  - As an indication of the execution status

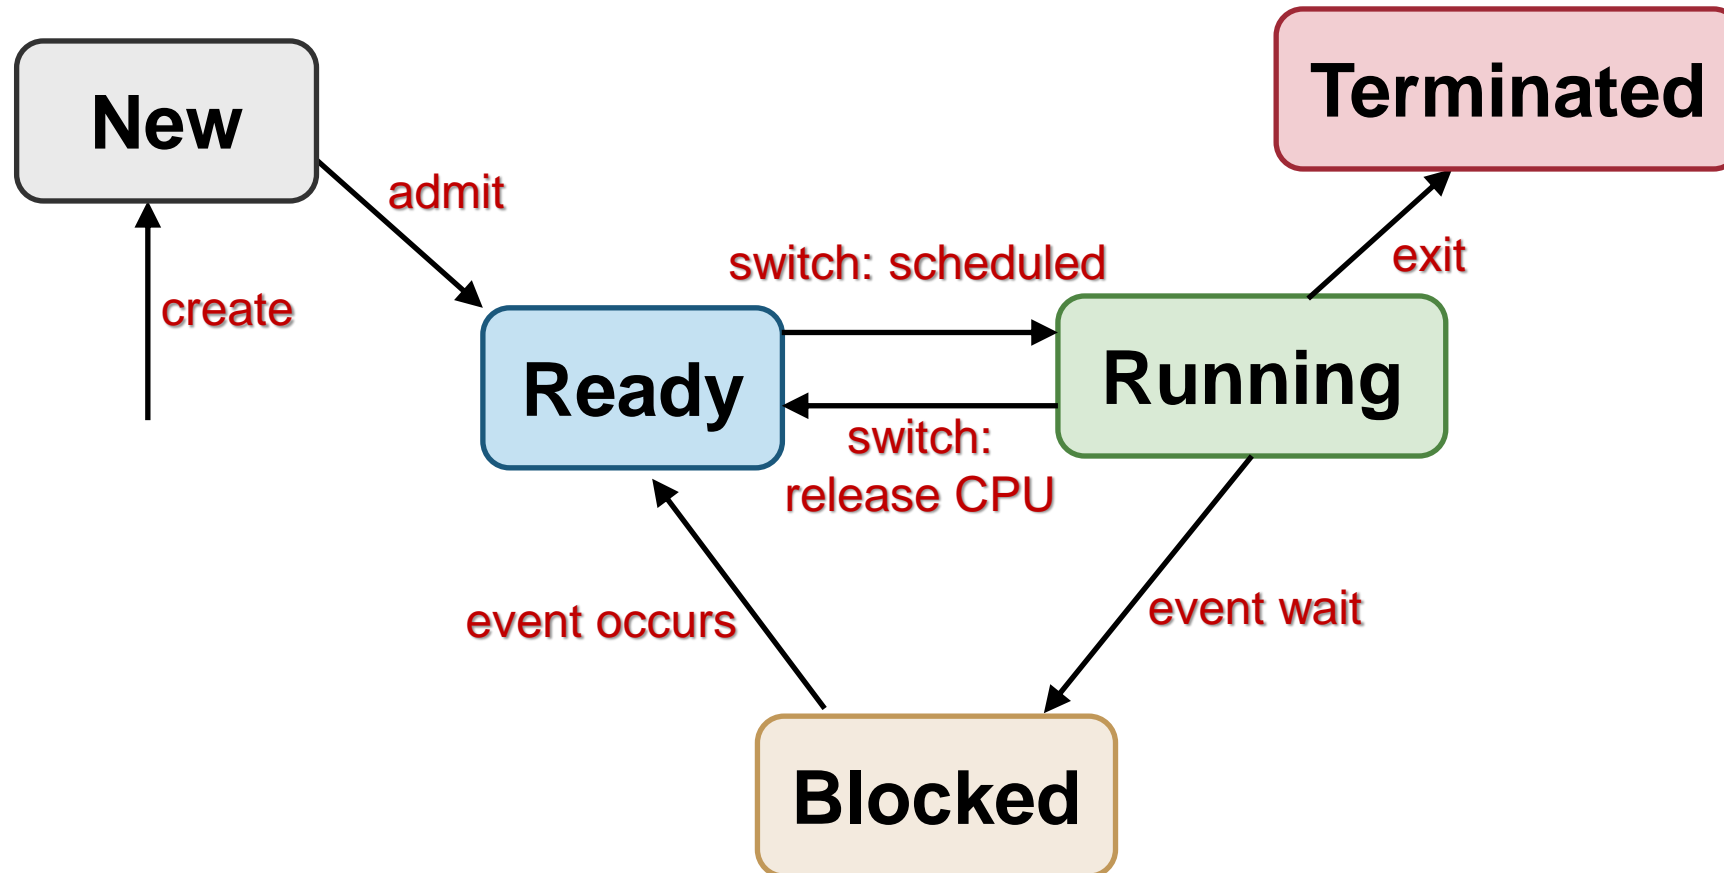
# (Simple) Process Model State Diagram



- The set of states and transitions are known as **process model**
  - Describes the behaviors of a process



# Generic 5-State Process Model



Notes: generic process states, details vary in actual OS

# Process States for 5-Stage Model

## ■ **New:**

- ❑ New process created
- ❑ May still be under initialization → not yet ready

## ■ **Ready:**

- ❑ process is waiting to run

## ■ **Running:**

- ❑ Process being executed on CPU

## ■ **Blocked:**

- ❑ Process waiting (sleeping) for event
- ❑ Cannot execute until event is available

## ■ **Terminated:**

- ❑ Process has finished execution, may require OS cleanup

# Process State Transitions in 5-Stage Model

- **Create** (nil  $\rightarrow$  New):
  - New process is created
- **Admit** (New  $\rightarrow$  Ready):
  - Process ready to be scheduled for running
- **Switch** (Ready  $\rightarrow$  Running):
  - Process selected to run
- **Switch** (Running  $\rightarrow$  Ready):
  - Process gives up CPU voluntarily or *preempted* by scheduler

# Process State Transitions

## ■ **Event wait** (Running → Blocked):

- ❑ Process requests event/resource/service which is not available/in progress
- ❑ Example events:
  - Acquiring lock, waiting for I/O, (*more later*)

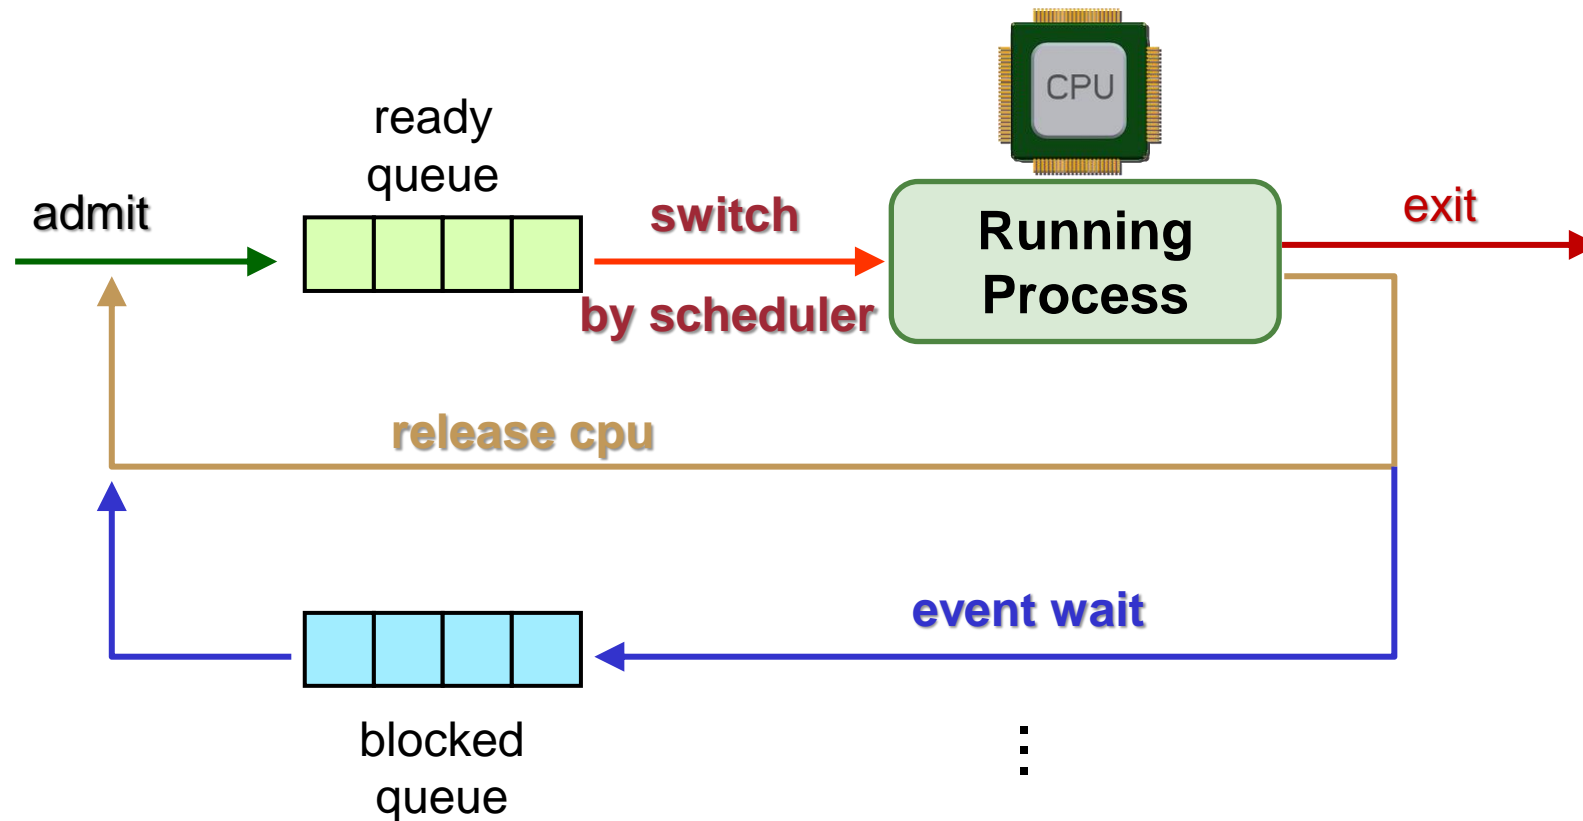
## ■ **Event occurs** (Blocked → Ready):

- ❑ Event occurs → process can continue

# Global View of Process States

- Given n processes:
  - With 1 CPU (core):
    - $\leq 1$  process in running state
    - conceptually 1 transition at a time
  - With m CPUs (cores):
    - $\leq m$  process in running state
    - possibly parallel transitions
- Different processes may be in different states
  - each process may be in different part of its state diagram
- **Assumption in CS2106: Our CPU has 1 core!**

# Queuing Model of 5 state transition



Notes:

- More than 1 process can be in ready + blocked queues
- May have separate event queues
- Queuing model gives global view of the processes, i.e. how the OS views them

# Checkpoint: Contexts updated

- When a program is **under execution**, there are **more information**:
  - Memory context:
    - Text and Data, Stack and Heap
  - Hardware context:
    - General purpose registers, Program Counter, Stack pointer, Stack frame pointer, ...
  - **OS context**:
    - **Process ID, Process State, ...**

# Process Table & Process Control Block

---

Putting it together



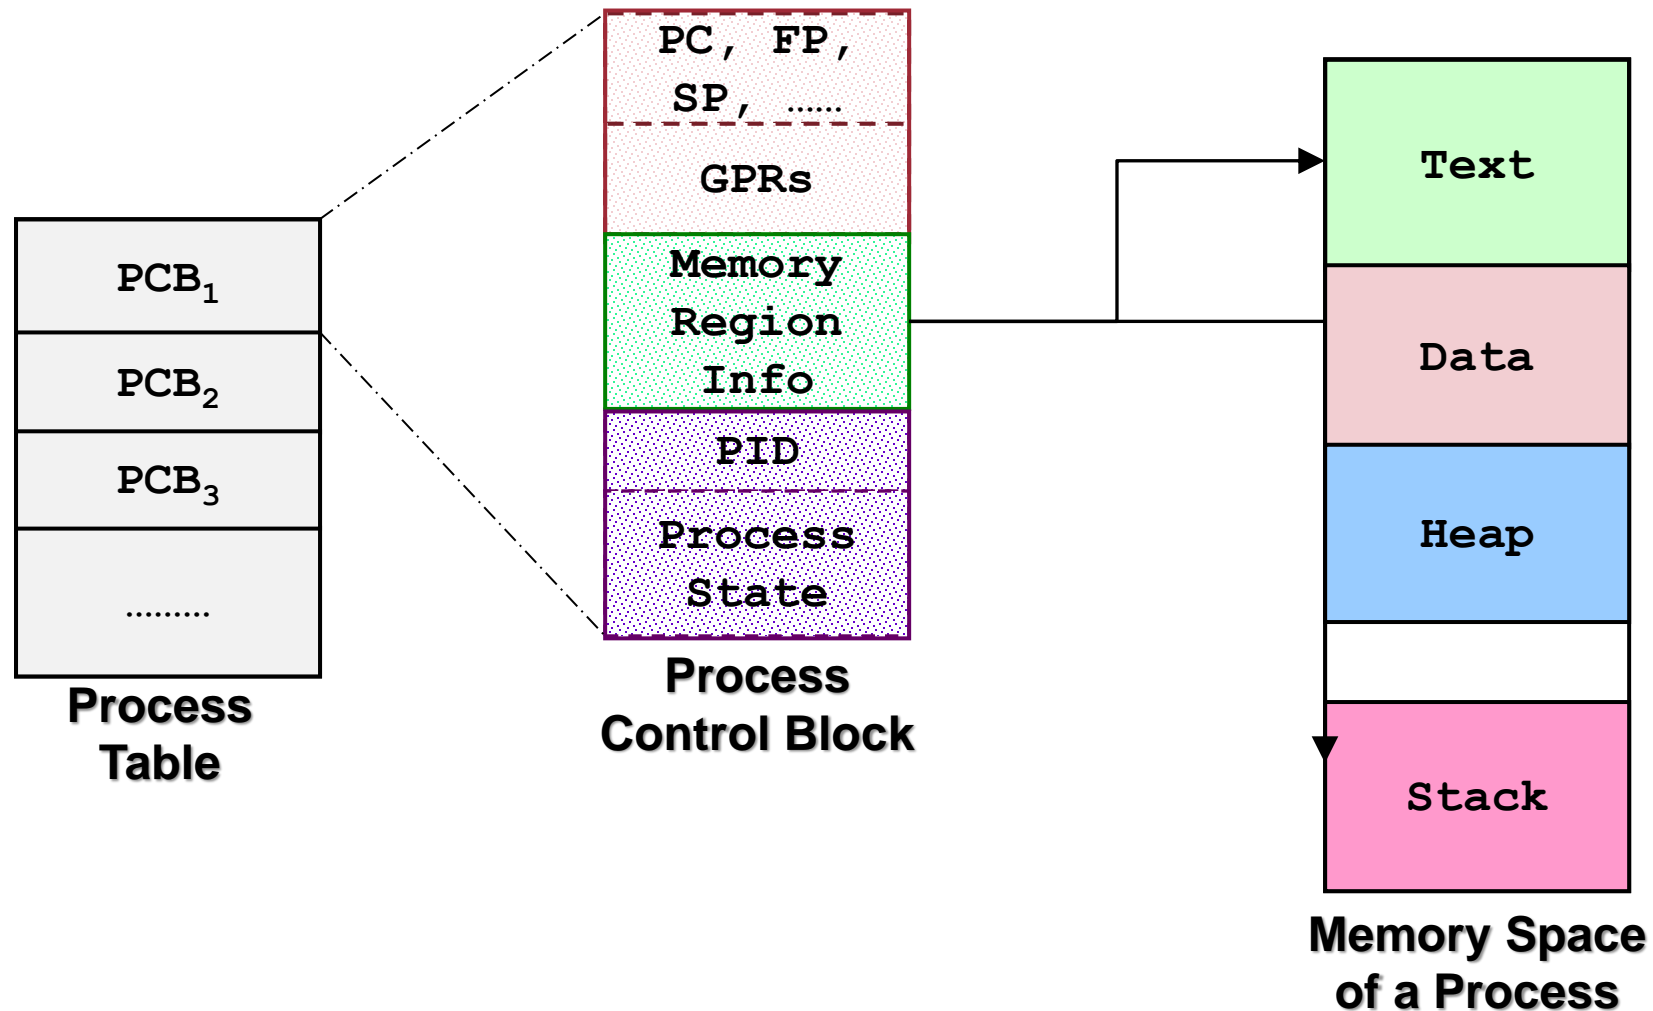
# Process Control Block & Table

- The entire execution context for a process
  - Traditionally called **Process Control Block** (PCB) or **Process Table Entry**
- Kernel maintains PCB for all processes
  - Conceptually stored as one table representing all processes

## Interesting Issues:

- Scalability
  - How many concurrent processes can you have?
- Efficiency
  - Should provide efficient access with minimum space wastage

# Illustration of a Process Table



---

Process interaction with OS

# System Calls

---

Can you please do this for me?

# System Calls

- Application Program Interface (API) to OS
  - Provides way of calling facilities/services in kernel
  - **NOT** the same as normal function call
    - have to change from user mode to kernel mode
- Different OS have different APIs:
  - Unix Variants:
    - Most follows **POSIX** standards
    - Small number of calls: ~100
  - Windows Family:
    - Uses **Win API** across different Windows versions
    - New version of windows usually adds more calls
    - Huge number of calls:~1000

# Unix System Calls in C/C++ program

- In C/C++ program, system call can be invoked *almost directly*
  - Majority of the system calls have a library version with the **same name** and the same parameters
    - The library version act as a **function wrapper**
  - Other than that, a few library functions present a more user friendly version to the programmer
    - E.g. lesser number of parameters, more flexible parameter values etc
    - The library version acts as a **function adapter**

# Example

```
#include <unistd.h>
#include <stdio.h>

int main()
{
    int pid;

    /* get Process ID */
    pid = getpid();

    printf("process id = %d\n", pid);

    return 0;
}
```

Library call that  
has the same  
name as a  
system call

Library call that  
make a system  
call

- System Calls invoked in this example:
  - ❑ `getpid()`
  - ❑ `write()` – made by `printf()` library call

# General System Call Mechanism

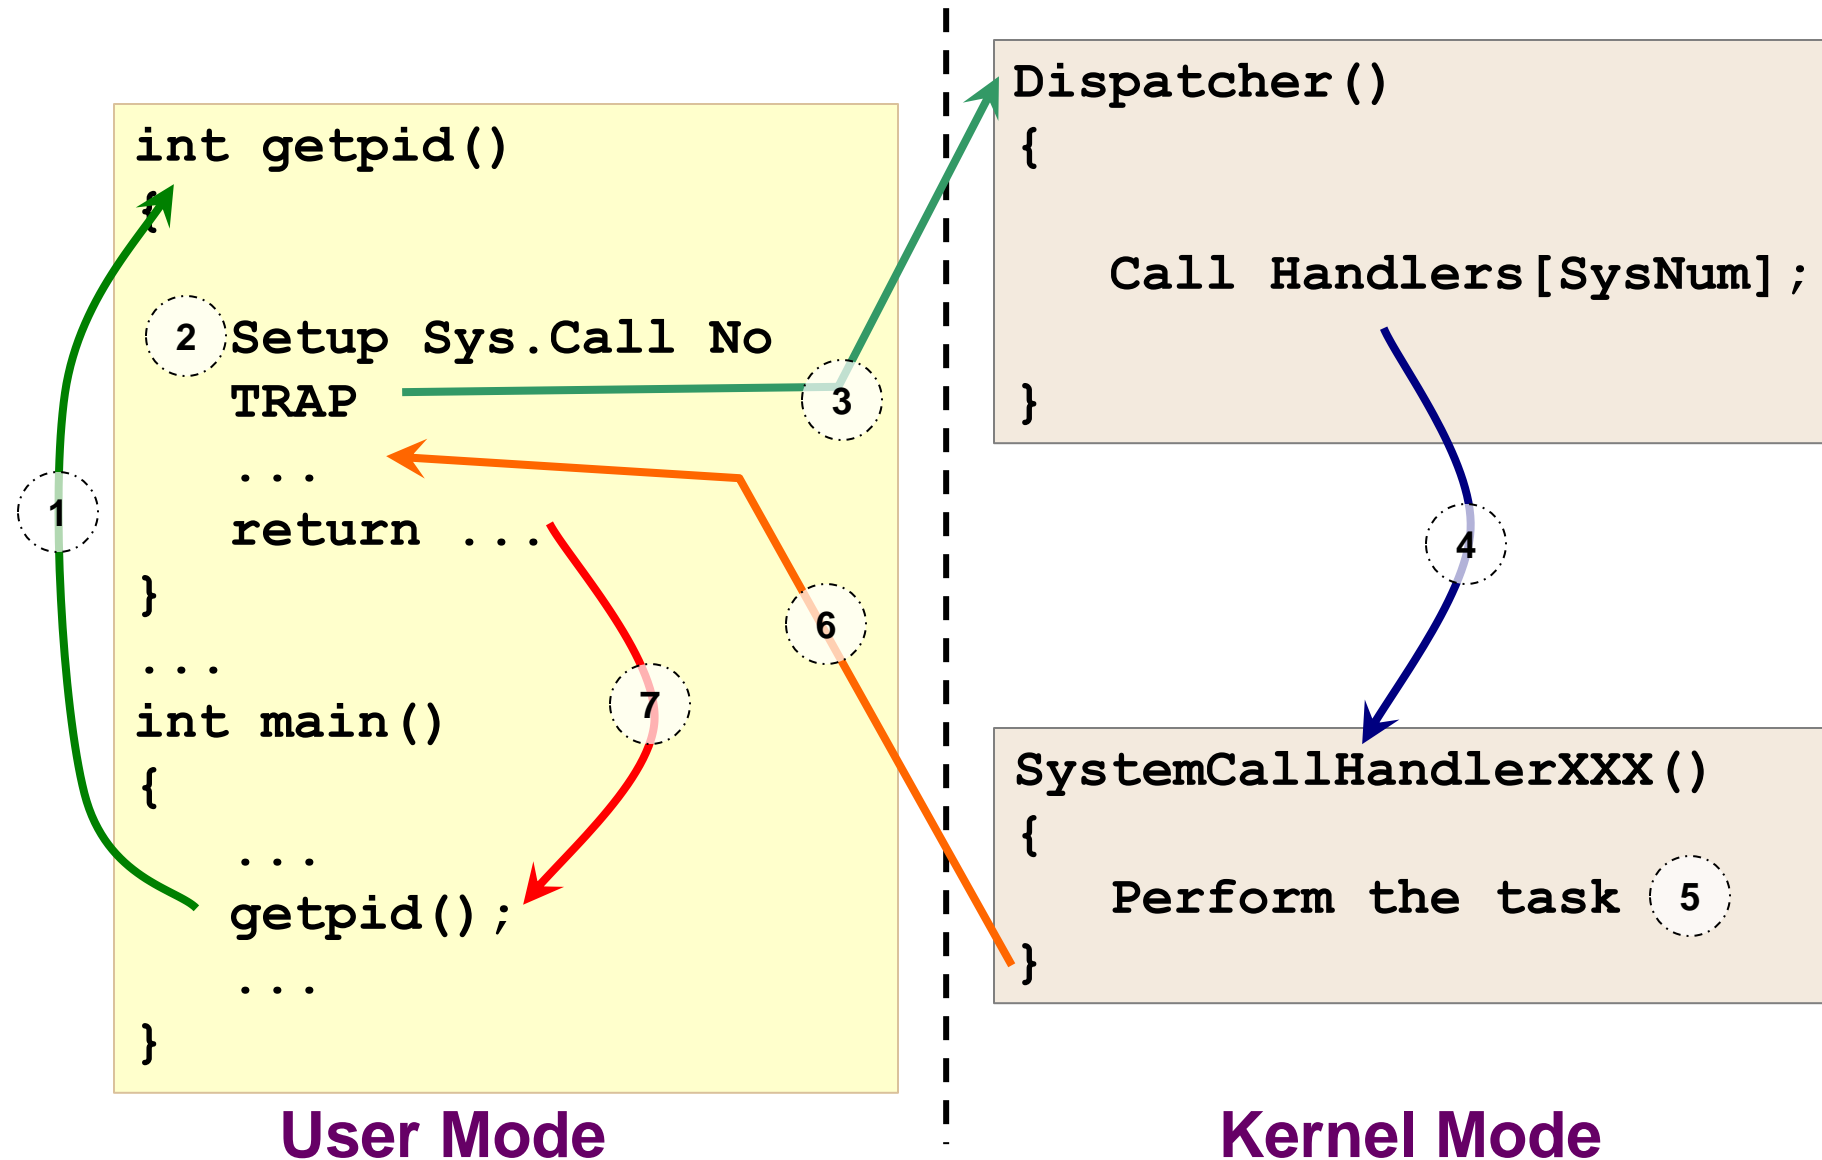
1. User program invokes the library call
  - Using the normal function call mechanism as discussed
2. Library call (usually in assembly code) places the **system call number** in a designated location
  - E.g., Register
3. Library call executes a special instruction to switch from user mode to kernel mode
  - That instruction is commonly known as **TRAP**

# General System Call Mechanism (cont)

4. Now in kernel mode, the appropriate system call handler is determined:
  - Save CPU state
  - Using the system call number as index
  - This step is usually handled by a **dispatcher**
5. System call handler is executed:
  - Carry out the actual request
6. System call handler ended:
  - Restore CPU state, and return to the library call
  - Switch from kernel mode to user mode
7. Library call return to the user program:
  - via normal function return mechanism



# Illustration: System Call Mechanism



Process interaction with OS

# Exception and Interrupt

---

Ops!

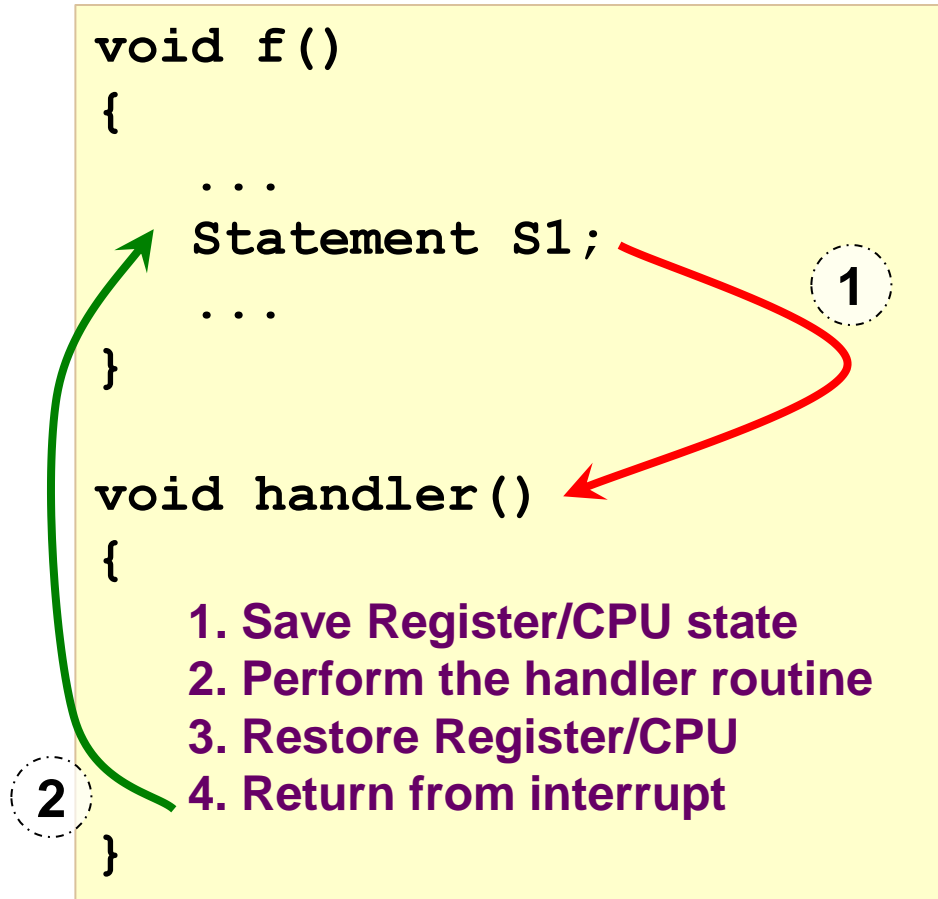
# Exception

- Executing a **machine level instruction** can cause exception
- For example:
  - Arithmetic Errors
    - Overflow, Division by Zero
  - Memory Accessing Errors
    - Illegal memory address, Misaligned memory access
  - etc.
- Exception is **Synchronous**
  - occur due to program execution
- Effect of exception:
  - Have to execute an **exception handler**
  - Similar to a **forced function call**

# Interrupt

- External events can interrupt the execution of a program
- Usually hardware related, e.g.:
  - ❑ Timer, Mouse Movement, Keyboard Pressed, etc.
- Interrupt is **asynchronous**
  - ❑ Events that occurs **independent** of program execution
- Effect of interrupt:
  - ❑ Program execution is suspended
  - ❑ Have to execute an **interrupt handler**

# Exception/Interrupt Handler: Illustration



1. Exception/Interrupt occurs:

- Control transfer to a handler routine **automatically**

2. Return from handler routine:

- Program execution resume
- **May** behave as if nothing happened

# Summary

- Using process as an abstraction of running program:
  - Necessary information (environment) of execution
  - Memory, Hardware and OS contexts
- Process from OS perspective:
  - PCB and process table
- OS  $\leftrightarrow$  Process interactions
  - System calls
  - Exception / Interrupt

# References

- Modern Operating System (3<sup>rd</sup> Edition)
  - Section 2.1
- Operating System Concepts (8<sup>th</sup> Edition)
  - Section 3.1