

---

Process Management

# Process Scheduling

---

Lecture 3

# Overview

- Concurrent Execution
- Process Scheduling
  - Definition
  - Process behavior
  - Processing environment
  - Criteria for good scheduling
  - Procedure of process scheduling
- Scheduling Algorithms
  - For batch processing systems
  - For interactive systems

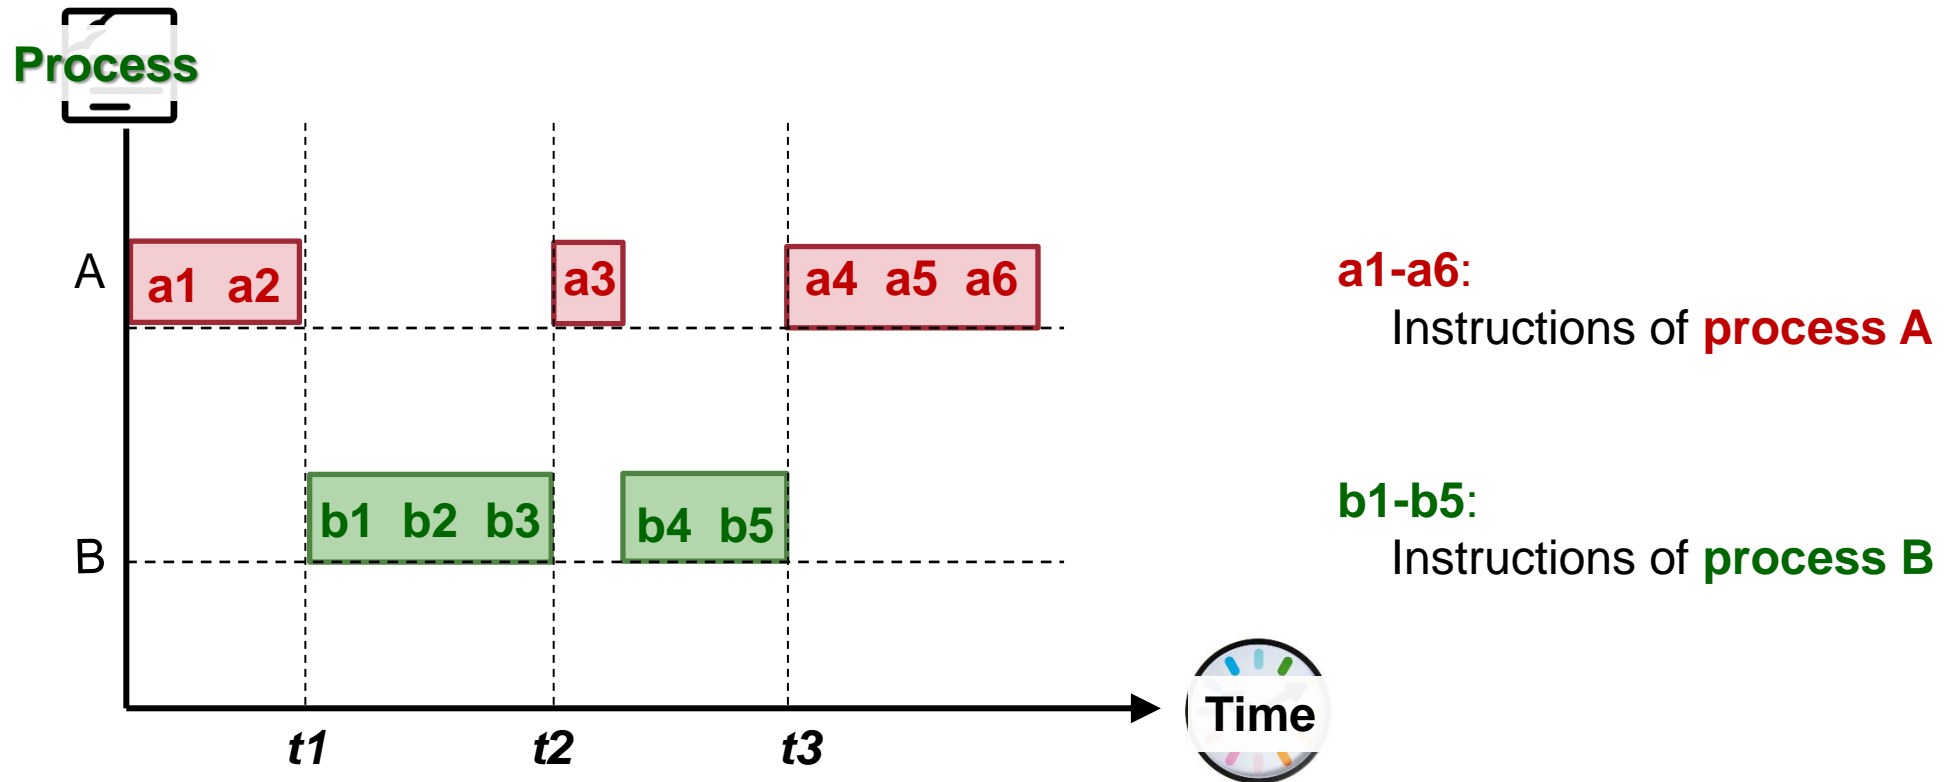
# Concurrent Execution

## ■ Concurrent processes:

- ❑ Logical concept meaning that multiple processes **progress** in execution (at the same time)
- ❑ Could be virtual parallelism:
  - illusion of parallelism (*pseudo-parallelism*)
- ❑ Could be physical parallelism
  - E.g. Multiple CPUs / Multi Core CPU to allow parallel execution of multiple processes


## ■ You can assume the two forms of parallelisms are not distinguished in the following discussion

# Concurrency Example (Simplistic)



Concurrent execution on 1 CPU (core):  
Interleave instructions from both processes  
Also called **timeslicing**


# Interleaved Execution (context switch)




- Multitasking needs to change context between A and B:
  - OS incurs overhead in switching processes

# Multitasking OS

- 1 core (CPU): timesliced execution of tasks



- Multiprocessor: timeslicing on  $n$  CPUs



# Scheduling in OS: A definition

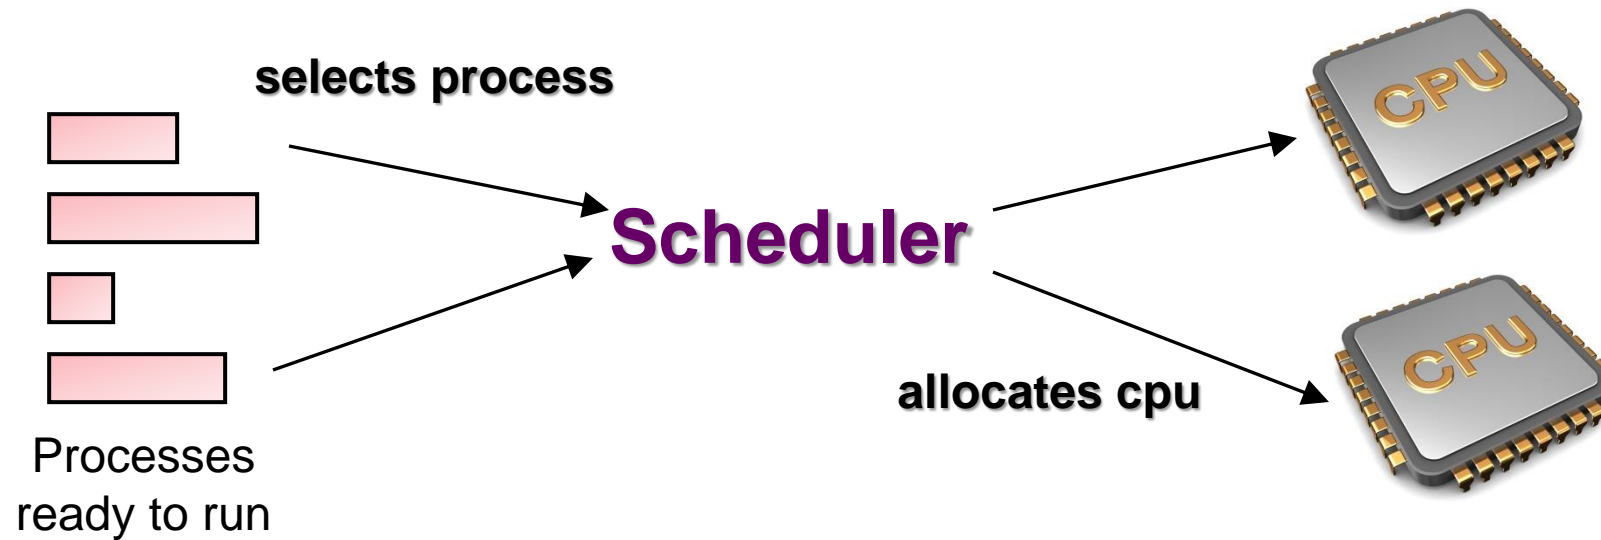
- Problems with having multiple processes:

- If ready-to-run process is more than available CPUs, which should be chosen to run?
  - Similar idea in thread-level scheduling
- Known as the **scheduling problem**

- **Terminology:**

- Scheduler
  - Part of the OS that makes scheduling decision
- Scheduling algorithm
  - The algorithm used by scheduler

# Scheduling: Illustration



- Each process has different requirement of CPU time
  - ❑ **Process behavior**
- Many ways to allocate
  - ❑ Influenced by the **process environment**
  - ❑ Known as **scheduling algorithms**
- A number of **criteria to evaluate the scheduler**



# Process Behavior

- A typical process goes through phases of:

## CPU-Activity:

- Computation
- E.g. Number crunching
- **Compute-Bound Process** spends majority of its time here

## IO-Activity:

- Requesting and receiving service from I/O devices
- E.g., Print to screen, read from file, etc.
- **IO-Bound Process** spends majority of its time here

# Processing Environment

## ■ Three categories:

### 1. **Batch Processing:**

- No user interaction required, No need to be responsive

### 2. **Interactive** (or Multiprogramming):

- With active user interacting with system
- Should be responsive: low and consistent in response time

### 3. **Real time processing:**

- Have deadline to meet
- Usually periodic process

# Criteria for Scheduling Algorithms

- Many criteria to evaluate scheduling algorithms:
  - ❑ Largely influenced by the processing environment
  - ❑ May be conflicting

## Criteria for **all** processing environments:

- **Fairness:**
  - ❑ Should get a fair share of CPU time
    - On a per process basis OR
    - On a per user basis
  - ❑ Also means **no starvation**
- **Utilization:**
  - ❑ All parts of the computing system should be utilized

# When to perform scheduling?

- Two types of scheduling policies

- Defined by **when** scheduling is triggered

- **Non-preemptive (Cooperative)**

- A process stayed scheduled (in running state) until it blocks or gives up the CPU voluntarily

- **Preemptive**

- A process is given a fixed time quota to run
    - possible to block or give up early
  - At the end of the time quota, the running process is suspended
    - Another ready process gets picked if available

# Scheduling a Process: Step-by-Step

1

- Scheduler is triggered (OS takes over)

2

- If context switch is needed:
  - Context of current running process is saved and placed on blocked queue / ready queue

3

- Pick a suitable process **P** to run base on scheduling algorithm

4

- Setup the context for **P**

5

- Let process **P** run

# SCHEDULING FOR BATCH PROCESSING

# Overview

- On batch processing system:
  - ❑ No user interaction
  - ❑ Non-preemptive scheduling is predominant
- Scheduling algorithms are generally easier to understand and implement
  - ❑ Commonly resulted in variants/improvements that can be used for other type of systems
- Three algorithms covered:
  - ❑ **F**irst-**C**ome **F**irst **S**erved (**FCFS**)
  - ❑ **S**hortest **J**ob **F**irst (**SJF**)
  - ❑ **S**hortest **R**emaining **T**ime (**SRT**)

# Criteria for batch processing

## ■ Turnaround time:

- ❑ Total time taken, i.e., finish time - arrival time
- ❑ Related to **waiting time**: time spent waiting for CPU

## ■ Throughput:

- ❑ Number of tasks finished per unit time
- ❑ i.e., Rate of task completion

## ■ CPU utilization:

- ❑ Percentage of time when CPU is working on a task



# First-Come First-Served: **FCFS**


## ■ General Idea:

- ❑ Tasks are stored on a First-In-First-Out (FIFO) queue based on arrival time
- ❑ Pick the first task in the queue to run until:
  - Task is done OR task is blocked
- ❑ Blocked task is removed from the FIFO queue
  - When it is ready again, it is placed at the back of queue
  - i.e., just like a newly arrive task

## ■ Guaranteed to have no **starvation**:

- ❑ The number of tasks in front of task X in FIFO is always decreasing  
➔ task X will get its chance eventually

# First-Come First-Served: Illustration



- The average total waiting time for 3 tasks
  - $(0 + 8 + 13)/3 = 7$  Time Units

# First-Come First-Served: Shortcomings

- Simple reordering can reduce the average waiting time!
- Also, consider this scenario:
  - ❑ First task (task **A**) is CPU-Bound and followed by a number of IO-Bound tasks **X**
  - ❑ Task **A** running
    - All tasks **X** waiting in ready queue (I/O device idling)
  - ❑ Task **A** blocked on I/O
    - All tasks **X** execute quickly and blocked on I/O (CPU idling)
  - ❑ known as **Convoy Effect**

# Shortest Job First: SJF


## ■ General Idea:

- ❑ Select task with the smallest total CPU time

## ■ Notes:

- ❑ Need to know **total CPU time** for a task in advance
  - Have to "guess" if this info is not available
- ❑ Given a fixed set of tasks:
  - Minimizes average waiting time
- ❑ Starvation is possible:
  - Biased towards short jobs
  - Long jobs may never get a chance!

# Shortest Job First: Illustration



- The average total waiting time for 3 tasks
  - $(0 + 3 + 8)/3 = 3.66$  Time Units
- Can be shown that SJF **guarantees** smallest average waiting time

# Shortest Job First: Predicting CPU Time

- A task usually goes through several phases of CPU-Activity:
  - Possible to guess the future CPU time requirement by the previous CPU-Bound phases
- Common approach (Exponential Average):
$$\text{Predicted}_{n+1} = \alpha \text{Actual}_n + (1-\alpha) \text{Predicted}_n$$
  - $\text{Actual}_n$  = The most recent CPU time consumed
  - $\text{Predicted}_n$  = The past history of CPU Time consumed
  - $\alpha$  = Weight placed on recent event or past history
  - $\text{Predicted}_{n+1}$  = Latest prediction

# Shortest Remaining Time: **SRT**

## ■ General Idea:

### □ Variation of SJF:


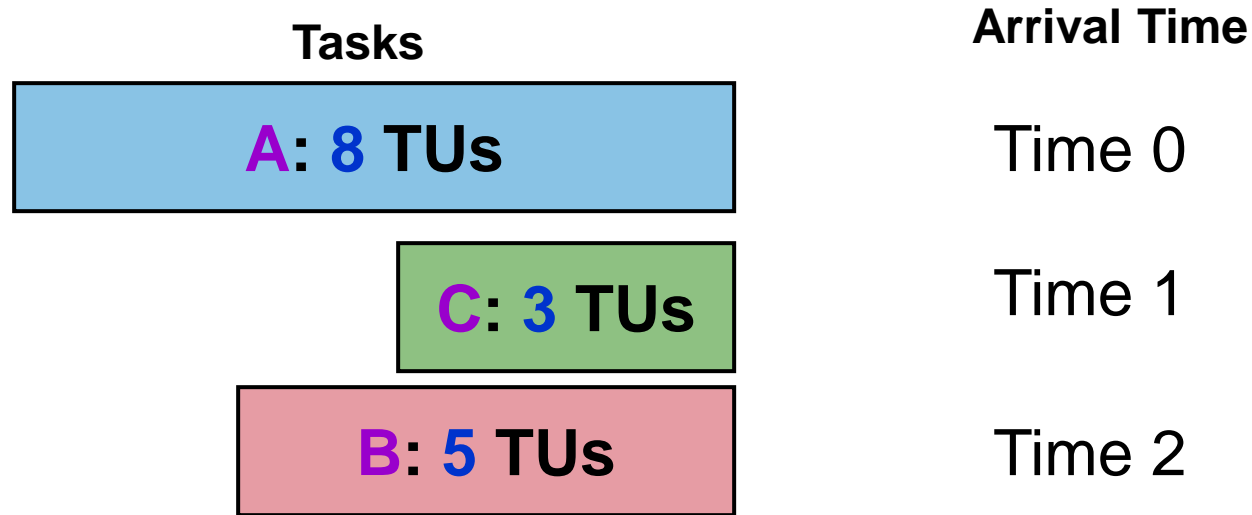
- Use remaining time
- Preemptive

### □ Select job with shortest remaining (or expected) time

## ■ Notes:

- New job with shorter remaining time can preempt currently running job
- Provide good service for short job even when it arrives late

# Shortest Remaining Time First: Illustration





# SCHEDULING FOR INTERACTIVE SYSTEMS

# Criteria for **interactive environment**

## ■ **Response time:**

- Time between request and response by system

## ■ **Predictability:**

- Variation in response time, lesser variation == more predictable

**Preemptive** scheduling algorithms are used to ensure good response time

→ Scheduler needs to run **periodically**

# Ensuring Periodic Scheduler

## ■ Questions:

- ❑ How can the scheduler "take over" the CPU periodically?
- ❑ How can we ensure the user program can never stop the scheduler from executing?



## ■ Ingredients for answer:

- ❑ Timer interrupt = Interrupt that goes off periodically (based on hardware clock)
  - ❑ OS ensure timer interrupt cannot be intercepted by any other program
- ➔ Timer interrupt handler **invokes scheduler**

# Terminology: Timer & Time Quantum

- Interval of Timer Interrupt (ITI):
  - ❑ OS scheduler is invoked on every timer interrupt
  - ❑ Typical values (1ms to 10ms )
- Time Quantum:
  - ❑ Execution duration given to a process
  - ❑ Could be constant or variable among the processes
  - ❑ Must be multiples of interval of timer interrupt
  - ❑ Large range of values (commonly 5ms to 100ms)

# Illustration: ITI vs Time Quantum



# Scheduling Algorithms:

- Algorithms covered:

1. Round Robin (RR)
2. Priority Based
3. Multi-Level Feedback Queue (MLFQ)
4. Lottery Scheduling

# Round Robin: **RR**

## ■ General Idea:

- ❑ Tasks are stored in a FIFO queue
- ❑ Pick the first task from queue front to run until:
  - A fixed **time slice** (**quantum**) elapsed, or
  - The task gives up the CPU voluntarily, or
  - The task blocks
- ❑ The task is then placed at the end of queue to wait for another turn
  - Blocked task will be moved to other queue to wait for its requested resource
- ❑ When blocked task is ready again, it is placed at the end of queue

# Round Robin: **RR**

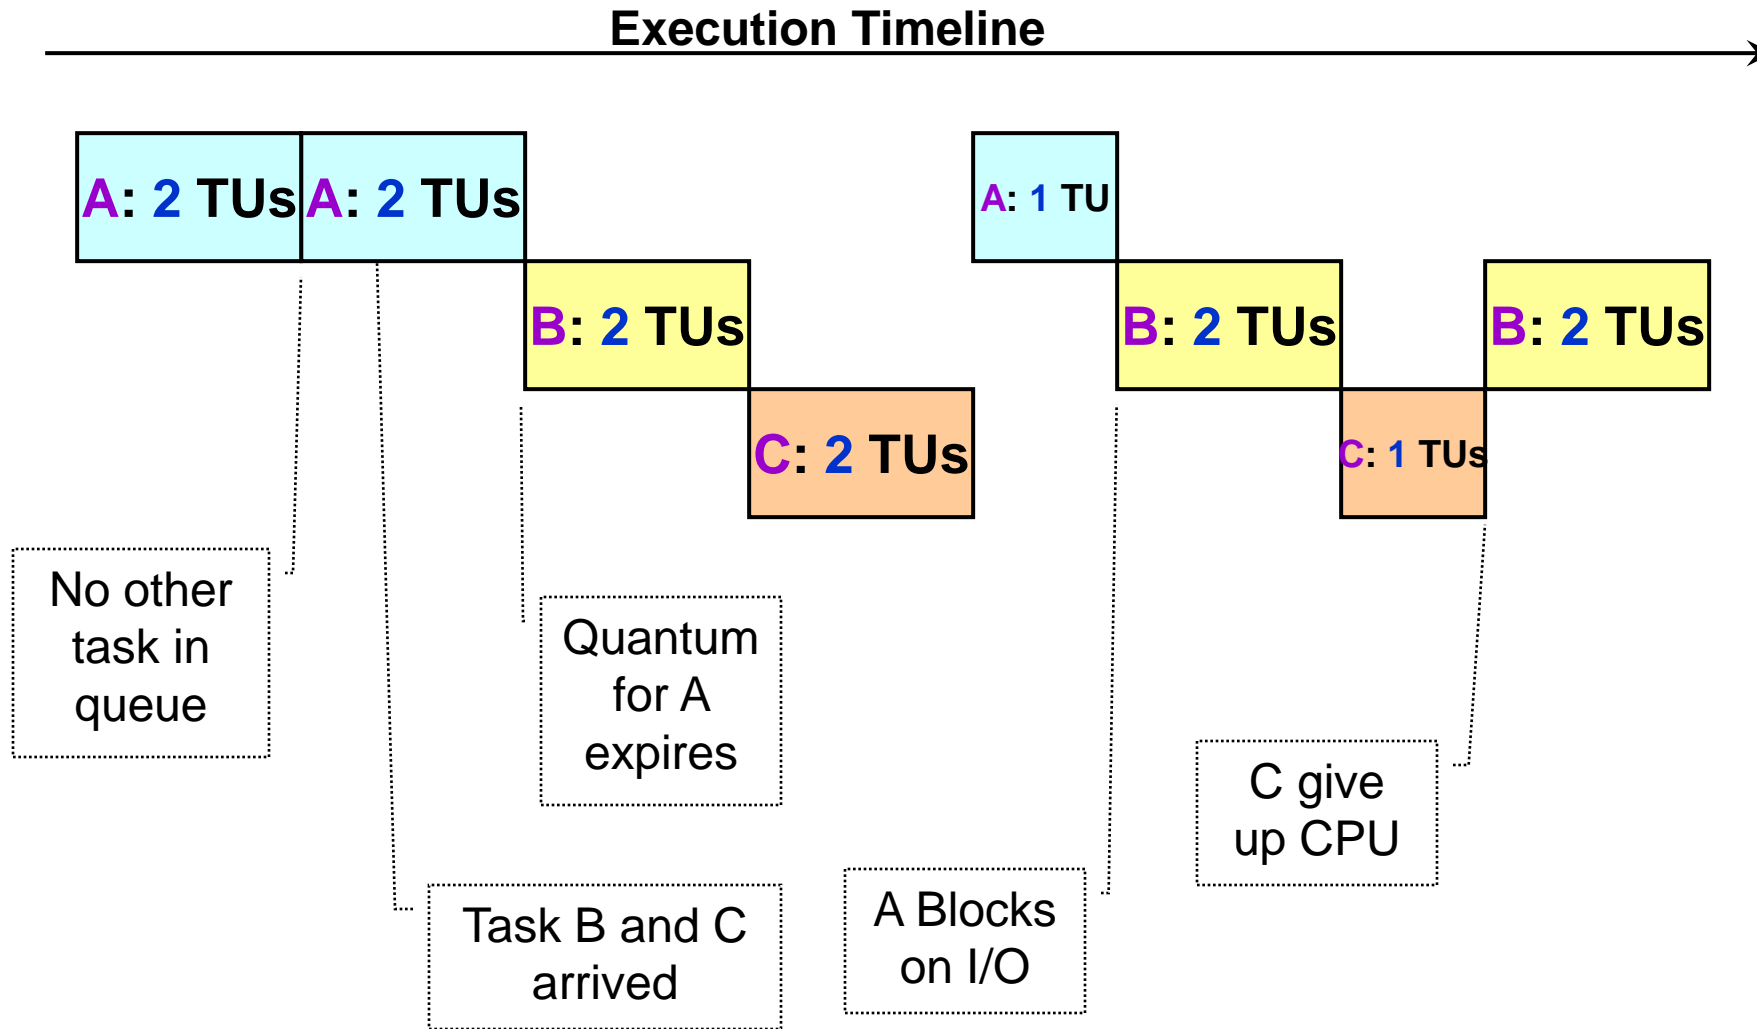
(cont)

## ■ Notes:

- ❑ Basically a preemptive version of FCFS
- ❑ **Response time guarantee:**
  - Given  $n$  tasks and quantum  $q$
  - Time before a task get CPU is bounded by  $(n-1)q$
- ❑ **Timer interrupt needed:**
  - For scheduler to check on quantum expiry
- ❑ The **choice of time quantum** duration is important:
  - Big quantum: Better CPU utilization but longer waiting time
  - Small quantum: Bigger overhead (worse CPU utilization) but shorter waiting time



# Round Robin: Illustration



# Priority Scheduling

## ■ General Idea:


- ❑ Some processes are more important than others
  - Cannot treat all process as equal
- ❑ Assign a priority value to all tasks
- ❑ Select task with highest priority value

## ■ Variants:

- ❑ Preemptive version:
  - Higher priority process can preempts running process with lower priority
- ❑ Non-preemptive version:
  - Late coming high priority process has to wait for next round of scheduling

# Priority Scheduling: Illustration

Tasks	Arrival Time	Priority (1=highest)
<b>A: 8 TUs</b>	Time 0	3
<b>C: 3 TUs</b>	Time 1	1
<b>B: 5 TUs</b>	Time 1	5



# Priority Scheduling: Shortcomings

- Low priority process can starve:
  - High priority process keep hogging the CPU
  - Even worse in preemptive variant
- Possible solutions:
  - Decrease the priority of currently running process after every time quantum
    - Eventually dropped below the next highest priority
  - Give the current running process a time quantum
    - This process is not considered in the next round of scheduling
- Generally, it is hard to guarantee or control the exact amount of CPU time given to a process using priority

# Priority Scheduling: Priority Inversion

- Consider the scenario:
  - Priority: {A = 1, B=3, C= 5} (1 is highest)
  - Task **C** starts and locks a resource (e.g., file)
  - Task **B** preempts **C**
    - **C** is unable to unlock the resource
  - Task **A** arrives and needs the same resource as **C**
    - but the resource is locked!
  - ➔ Task **B** continues execution even if Task **A** has higher priority
- Known as **Priority Inversion**:
  - Lower priority task preempts higher priority task

# Multi-level Feedback Queue (MLFQ)

- Designed to solve one BIG + HARD issue:
  - ❑ How do we schedule without perfect knowledge?
  - ❑ Most algorithms require certain information (process behavior, running time, etc.)
- MLFQ is:
  - ❑ Adaptive: "Learn the process behavior automatically"
  - ❑ Minimizes both:
    - Response time for IO bound processes
    - Turnaround time for CPU bound processes

# MLFQ: Rules

## ■ Basic rules:

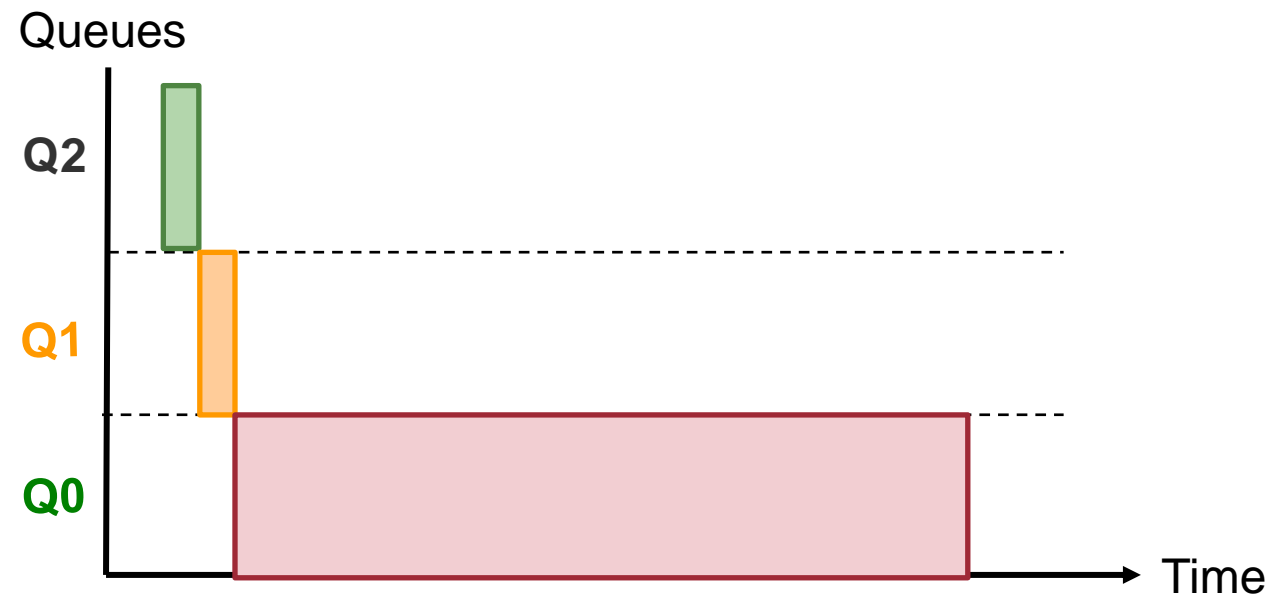
1. If  $\text{Priority}(A) > \text{Priority}(B) \rightarrow A$  runs
2. If  $\text{Priority}(A) == \text{Priority}(B) \rightarrow A$  and  $B$  runs in RR

## ■ Priority Setting/Changing rules:

1. New job  $\rightarrow$  Highest priority
2. If a job fully utilized its time quantum  $\rightarrow$  priority reduced
3. If a job gives up / blocks before finishes its time quantum  $\rightarrow$  priority retained

# MLFQ: Example 1


- 3 Queues: Q2 (highest priority), Q1, Q0
- A single long running job
  - Try to apply the rules and check your understanding





# MLFQ: Example 2


- Example 1 + a short job in the middle
  - A short job appears sometime in the middle



# MLFQ: Example 3

## ■ Two jobs:

- **A** = CPU bound (already in the system for quite some time)
- **B** = I/O bound



# MLFQ: Questions to ponder

- Can you think of a way to abuse the algorithm? 😊
  - Equivalent question: MLFQ does not work well for what kind combination of jobs?
- What are the ways to rectify the above?

# Lottery Scheduling

## ■ General Idea:

- ❑ Give out “lottery tickets” to processes for various system resources
  - E.g., CPU time, I/O devices, etc.
- ❑ When a scheduling decision is needed:
  - A lottery ticket is chosen randomly among eligible tickets
  - The winner is granted the resource
- ❑ In the long run, a process holding **X%** of tickets
  - Can win **X%** of the lottery held
  - Use the resource **X%** of the time

# Lottery Scheduling: Properties

- **Responsive:**
  - ❑ A newly created process can participate in the next lottery
- **Provides good level of control:**
  - ❑ A process can be given  $Y$  lottery tickets
    - It can then distribute to its child process
  - ❑ An important process can be given more lottery tickets
    - Can control the proportion of usage
  - ❑ Each resource can have its own set of tickets
    - Different proportion of usage per resource per task
- **Simple Implementation**

# Summary

## ■ Scheduling in OS:

- ❑ Basic definition
- ❑ Factors that affect scheduling
  - Process, Environment
- ❑ Criteria of good scheduling

## ■ Scheduling Algorithms:

- ❑ FCFS, SJF, SRT for batch processing systems
- ❑ RR, Priority base, Multi-Level Queues, MLFQ and Lottery scheduling for interactive systems

# Reference

- Modern Operating System (3<sup>rd</sup> Edition)
  - ❑ By Andrew S.Tanenbaum
  - ❑ Published by Pearson
  - ❑ Chapter 2.4
- Operating System Concepts (7<sup>th</sup> Edition)
  - ❑ By Silberschatz, Galvin, and Gagne
  - ❑ Published by Wiley Brothers
  - ❑ Chapter 5
- Operating Systems: Three Easy Pieces
  - ❑ By Arpaci-Dusseau and Arpaci-Dusseau
  - ❑ <http://pages.cs.wisc.edu/~remzi/OSTEP/>