

Memory Management

Virtual Memory Management

Lecture 9

Overview

■ **Virtual Memory:**

- ❑ Motivation
- ❑ Basic Idea
- ❑ Page Fault

■ **Common Applications of Virtual Memory:**

- ❑ Demand Paging

■ **Aspects of Virtual Memory Management:**

- ❑ Page Table Structure
- ❑ Page Replacement Algorithms
- ❑ Frame Allocation

Virtual Memory: Motivation

- Our last assumption of memory usage:
 - Physical memory is large enough to hold one or more process logical memory space completely
- This assumption is too restrictive:
 - What if the logical memory space of process is \gg than physical memory?
 - What if the same program is executed on a computer with lesser physical memory?

Virtual Memory: Basic Idea

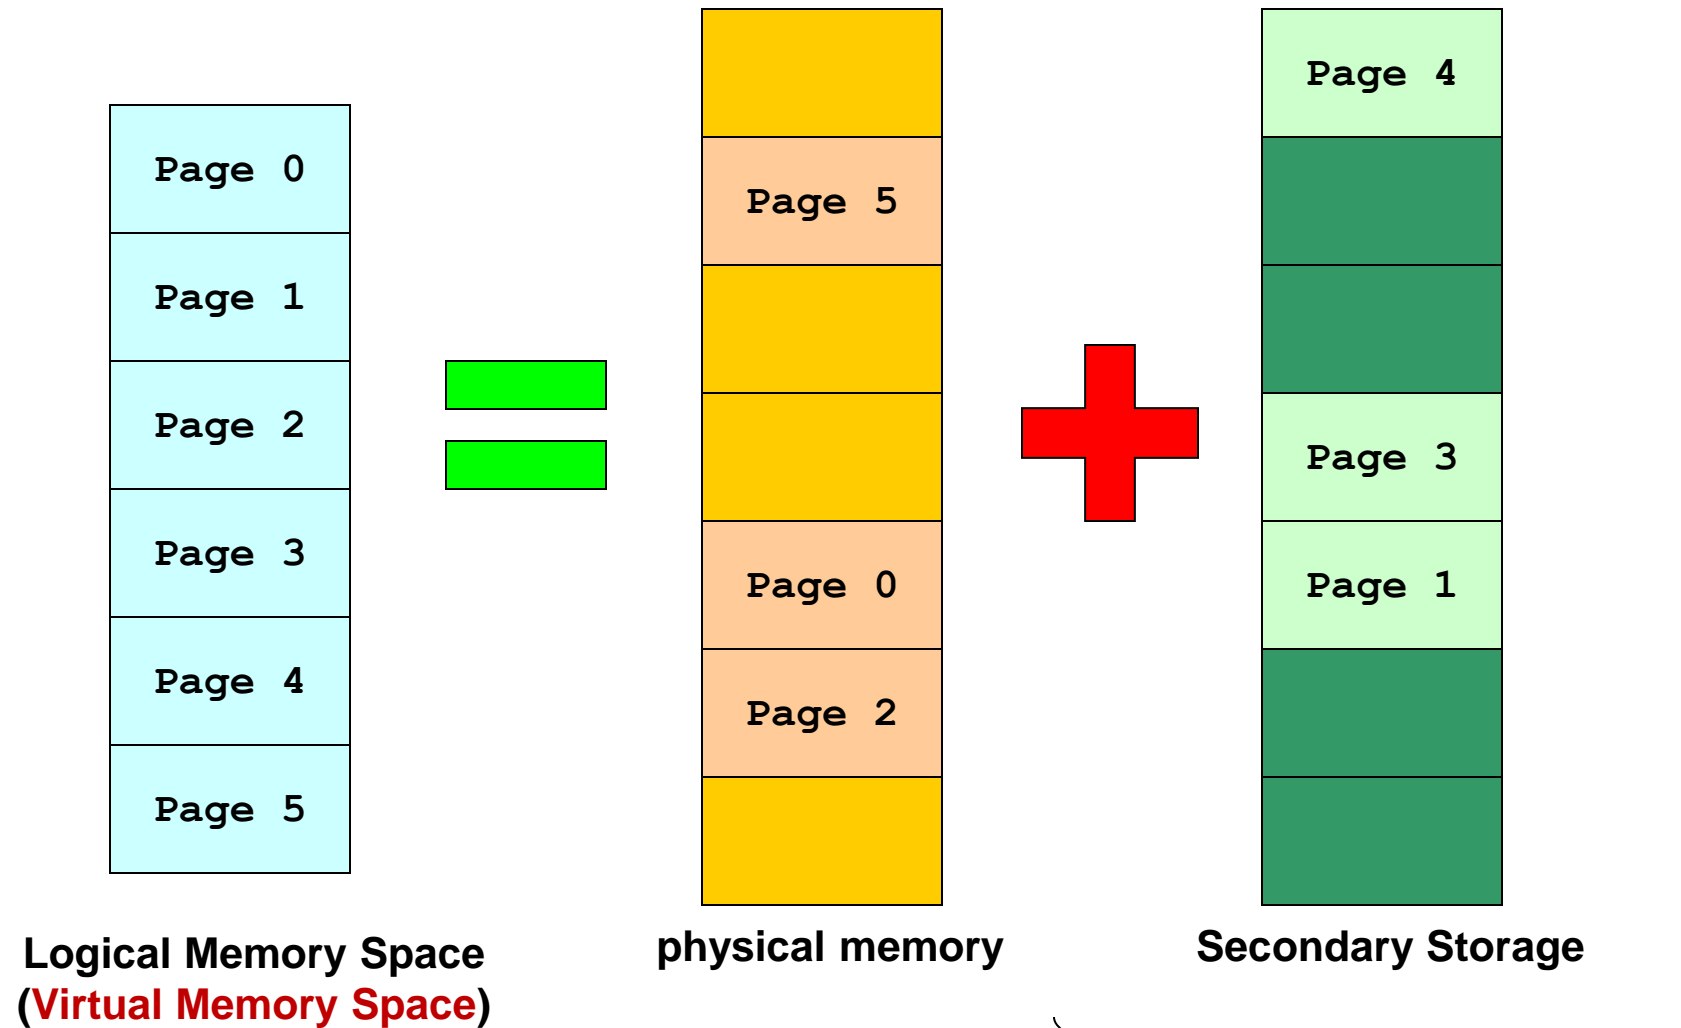
■ Observation:

- ❑ The logical memory space of a process >> physical memory
- ❑ Secondary storage has much larger capacity compared to physical memory

■ Basic Idea:

- ❑ Split the logical address space into pages:
 - Some pages reside in physical memory
 - **Other are stored on secondary storage**

Virtual Memory: Paging Illustration



Extended Paging Scheme

- Basic idea remains unchanged:
 - Use page table to translate **virtual** address to physical address
- New addition:
 - To distinguish between two pages types
 - **memory resident** (pages in physical memory)
 - non-memory resident (pages in secondary storage)
 - Use a (***is memory resident?***) bit in page table entry
 - CPU can only access memory resident pages:
 - **Page Fault**: When CPU tries to access non-memory resident page
 - OS need to bring a non-memory resident page into physical memory

Accessing Page X: General Steps

By Hardware

1. Check **page table**:

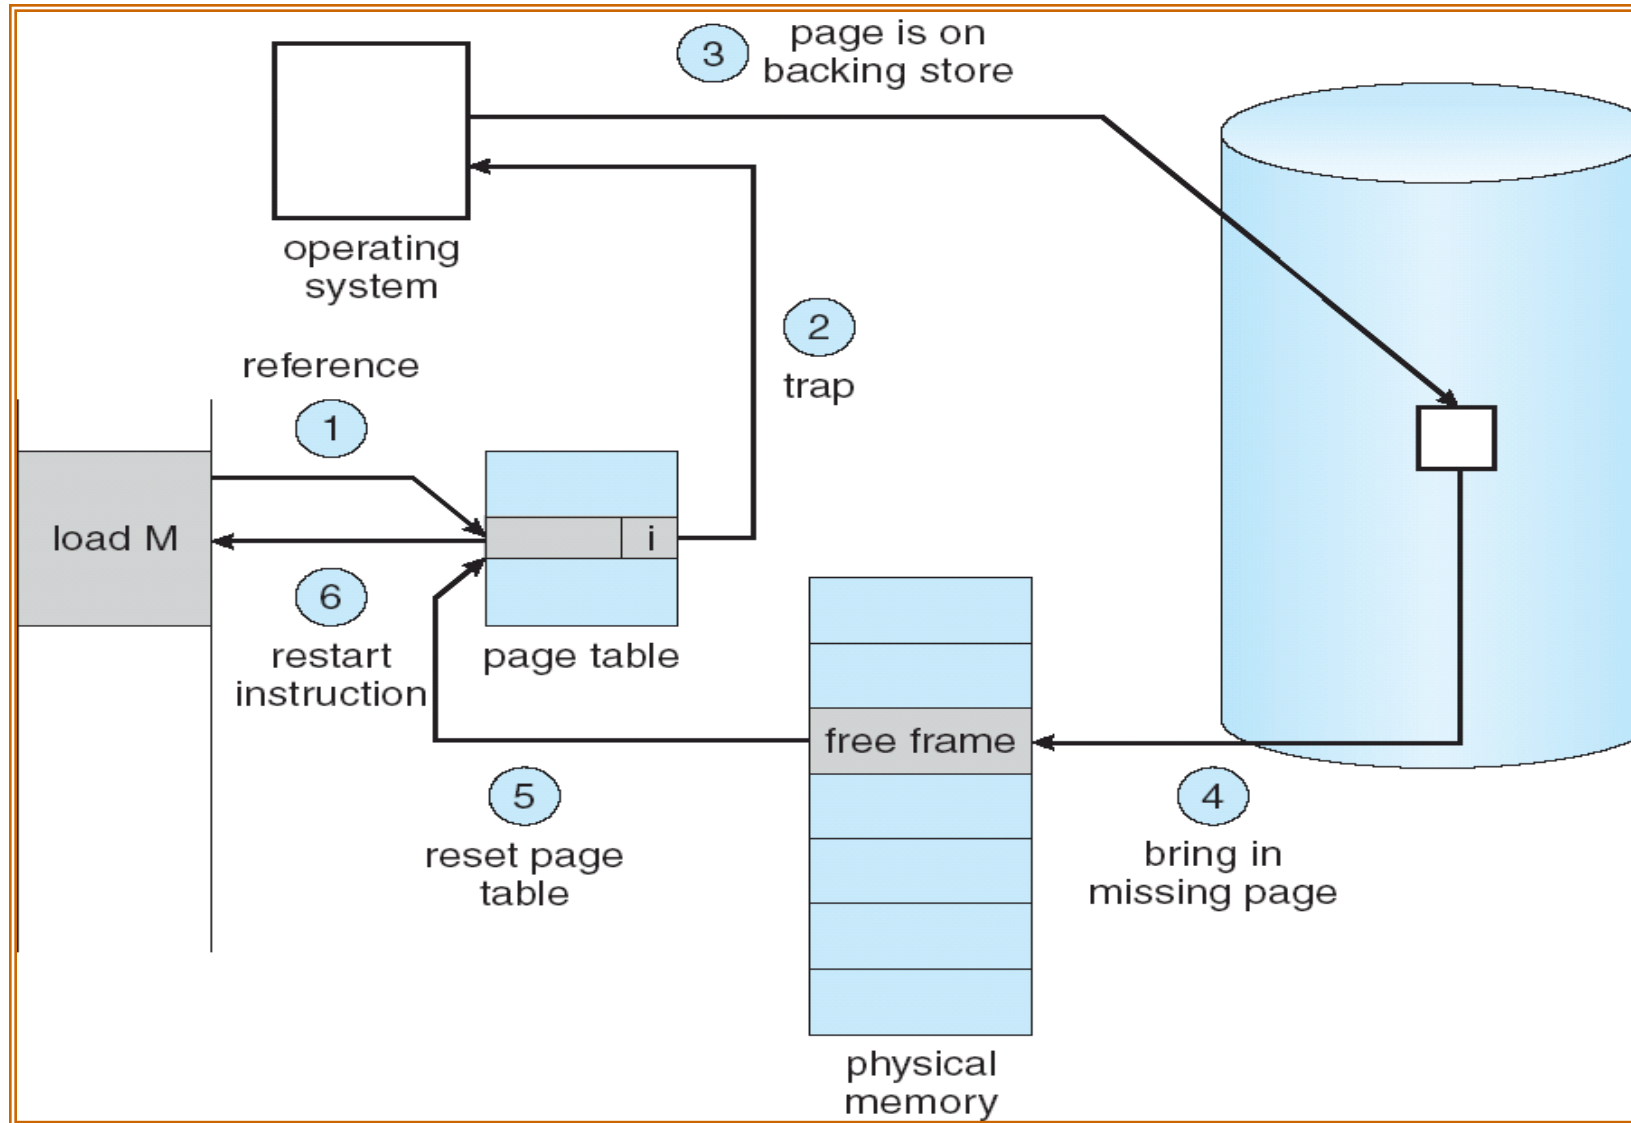
- Is page ***X*** ***memory resident***?
 - Yes: Access physical memory location. Done.
 - No: Continue to the next step

2. **Page Fault: Trap** to OS

By OS

3. Locate page X in secondary storage
4. Load page X into a physical memory frame
5. Update page table
6. Go to step 1 and retry

Virtual Memory Accessing: Illustration



Virtual Memory: Justification

- Observation:
 - ❑ Secondary Storage access time >> Physical memory access time
- If memory access results in page fault most of the time:
 - ❑ Non-memory resident pages need to be loaded
 - ❑ Known as **thrashing**
- How do we know that **thrashing** is **unlikely** to happen?
 - ❑ Related: How do we know that after a page is loaded, it is likely to be useful for future accesses?

Recap: Locality Principles

- *Most programs* exhibit these behaviors:
 - Most time are spent on a relatively small region of the code
 - Within a time period, accesses are made to a relatively small part of the data only
- Formalized as **locality principles**:
 - **Temporal Locality**:
 - Memory address which is referenced *is likely to be referenced again*
 - **Spatial Locality**:
 - Memory addresses close to a referenced address is likely to be referenced

Virtual Memory and Locality Principle

■ Exploiting **Temporal Locality**:

- ❑ After a page is loaded to physical memory, it is likely to be accessed in near future
 - Cost of loading the page is **amortized**

■ Exploiting **Spatial Locality**:

- ❑ A page contains contiguous addresses that are likely to be accessed in near future
 - Later access to nearby addresses will not cause page fault

■ However, there are always exceptions 😊

- ❑ Programs that behave badly due to poor design or with malicious intention

Virtual Memory: Summary

- Completely separate logical memory addresses from physical memory
 - Amount of physical memory no longer restrict the size of logical memory space
- More efficient use of physical memory
 - Page currently not needed can be on secondary storage
- Allow more processes to reside in memory
 - Improve CPU utilization as there are more processes to choose to run

More on Virtual Memory Management

- More in-depth looks on several aspects:
 - Large page table with big logical memory space → How to structure the page table for efficiency?
 - **Page Table Structures**
 - Each process has limited number of resident memory pages
→ Which page should be replaced when needed?
 - **Page Replacement Algorithms**
 - Limited physical memory frames
→ How to distribute among the processes?
 - **Frame Allocation Policies**

Waste not, want not

PAGE TABLE STRUCTURE

Page Table Structure

- Page table information is kept together with the process information and takes up physical memory space
- Modern computer systems provide huge logical memory space
 - 4GiB(32bit) was common, 8TiB or (much) more is possible now
 - Huge logical memory space → Large number of pages
 - Each page has a page table entry → Large page tables
- Problems with large page table
 - High overhead
 - Fragmented page table:
 - Page table occupies several memory pages

Page Table Structure: **Direct Paging**

- Direct Paging: keep all entries in a single table
- With 2^p pages in logical memory space
 - p bits to specify one unique page
 - 2^p page table entries (PTE), each contains:
 - physical frame number
 - additional information bits (valid/invalid, access right, etc.)
- Example:
 - Virtual Address: 32 bits, Page Size = 4KiB
 - $P = 32 - 12 = 20$
 - Size of PTE = 2 bytes
 - Page Table Size = $2^{20} * 2 \text{ bytes} = 2\text{MiB (!)}$

2-Level Paging: Basic Idea

■ Observation:

- ❑ Process may not use the entire virtual memory space → Full page table is a waste!

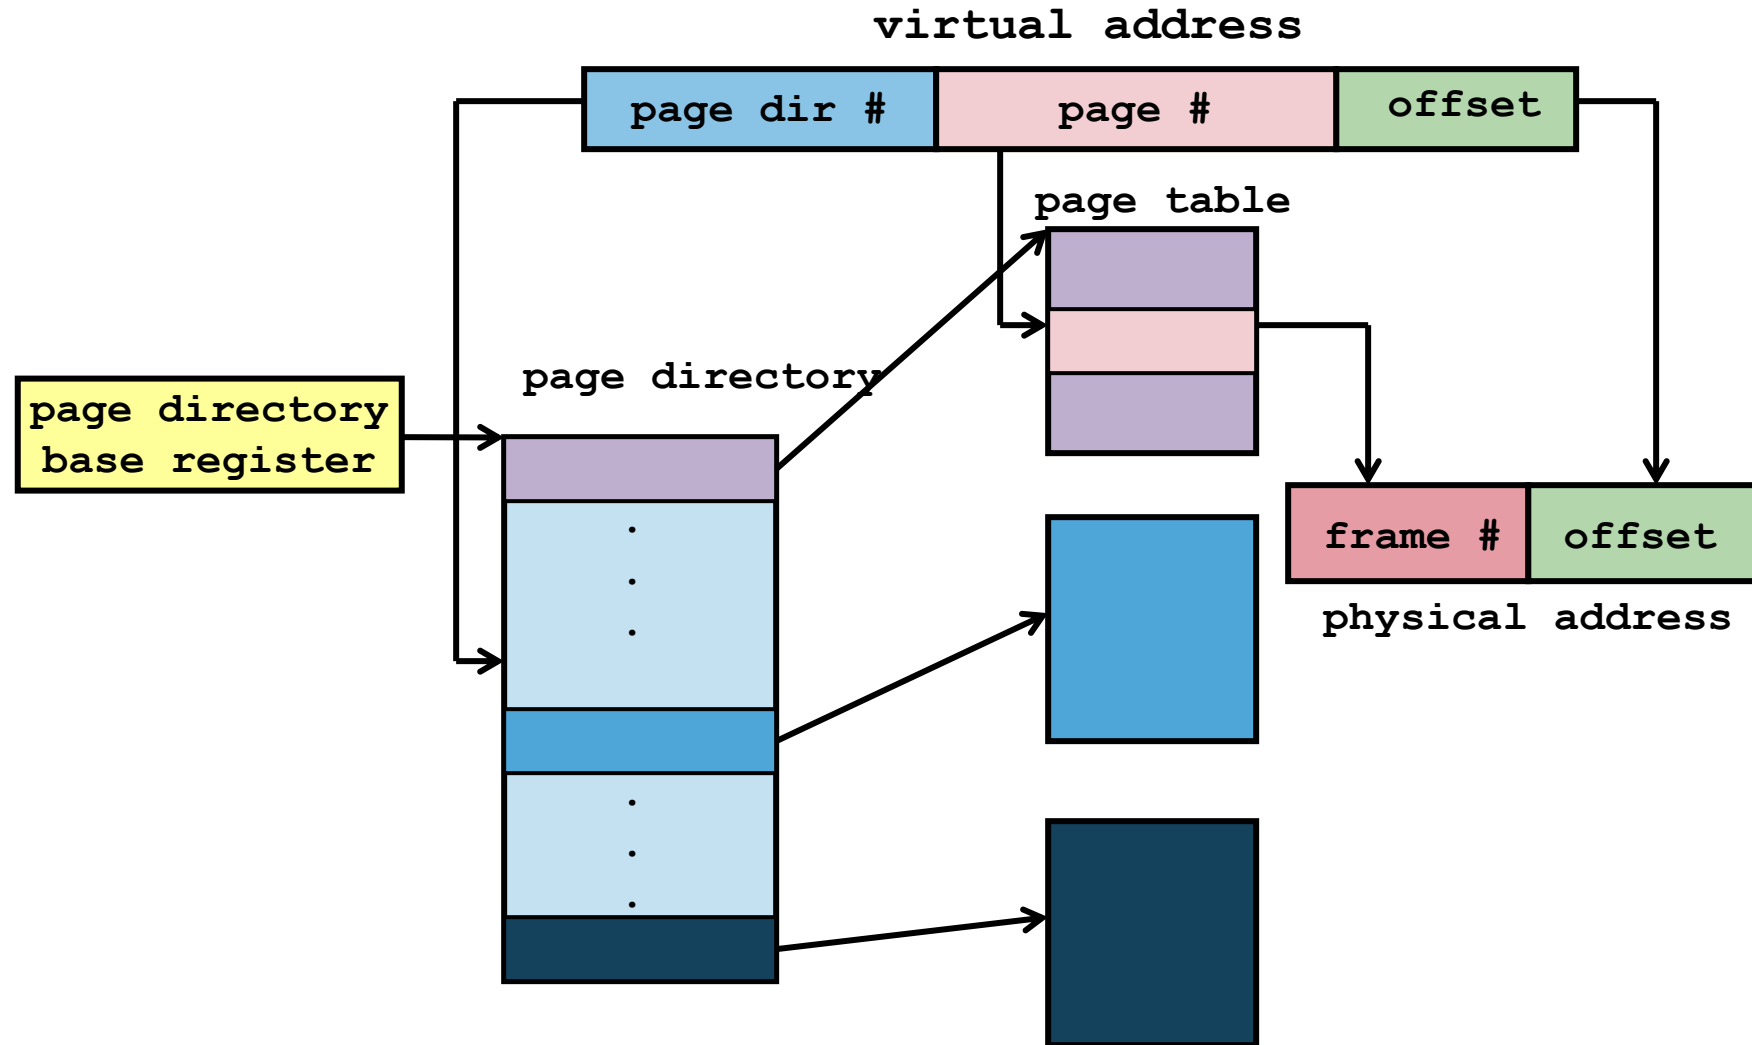
■ Basic Idea:

- ❑ Split the full page table into regions
- ❑ Only a few regions are used
 - As memory usage grows, new regions can be allocated
- ❑ This idea is similar to split logical memory space into pages 😊
- ❑ Need a directory to keep track of the regions
 - Analogues of page table ↔ pages

2-Level Paging: Description

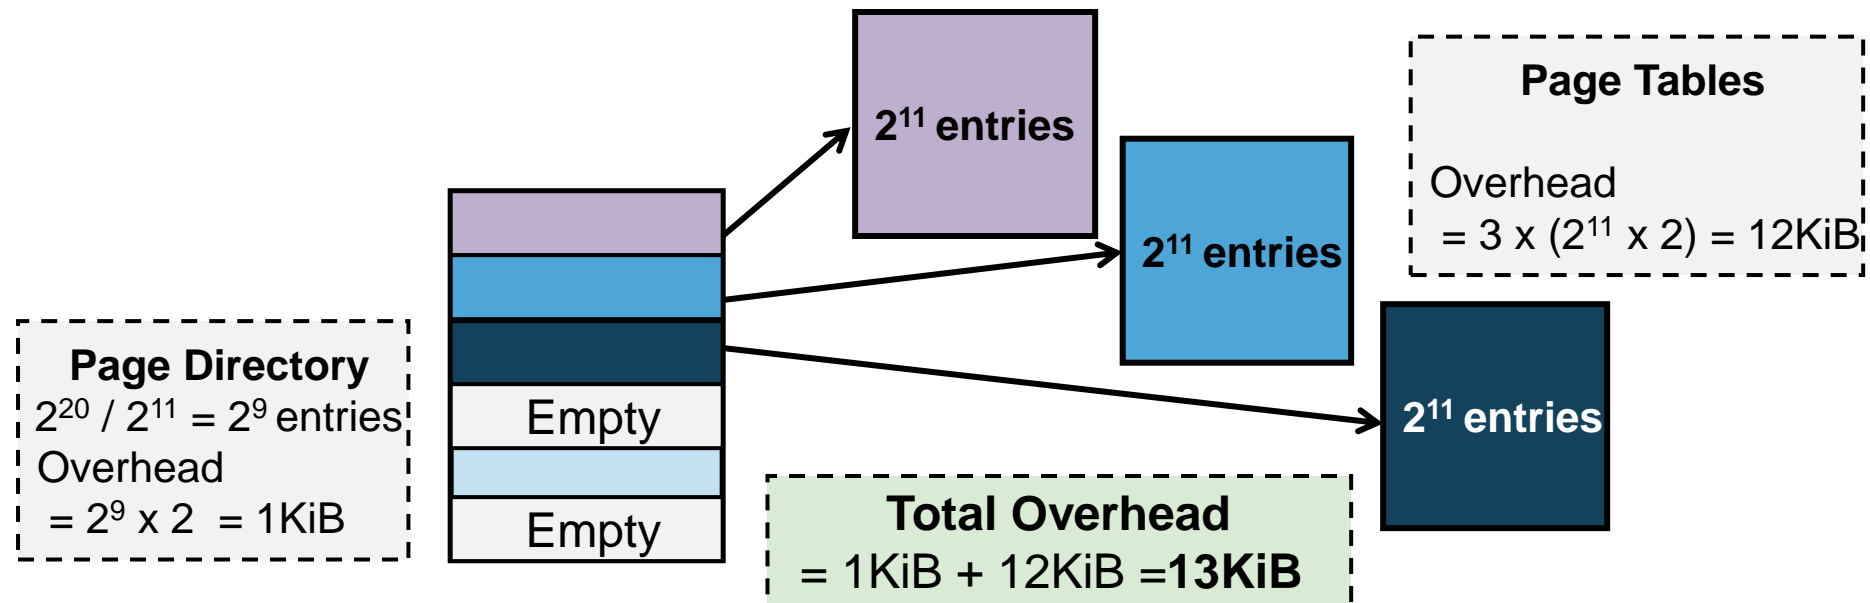
- Split page table into *smaller page tables*
 - Each with a **page table number**
 - Each **page table size** is equal to the **page size**
- If the original page table has 2^P entries:
 - With 2^M smaller page tables, M bits are needed to uniquely identify one page table
 - Each smaller page table contains $2^{(P-M)}$ entries
- To keep track of the smaller page tables
 - A single **page directory** is needed
 - Page directory contains 2^M indices to locate each of the smaller page table

2-level Paging: Illustration



2-Level Paging: Advantages

- We can have empty entries in the page directory
 - ➔ The corresponding page tables need not be allocated!
- Using the same setting as the previous example:
 - ❑ Assume only 3 page tables are in use
 - ❑ Overhead = 1 page directory + 3 smaller page tables



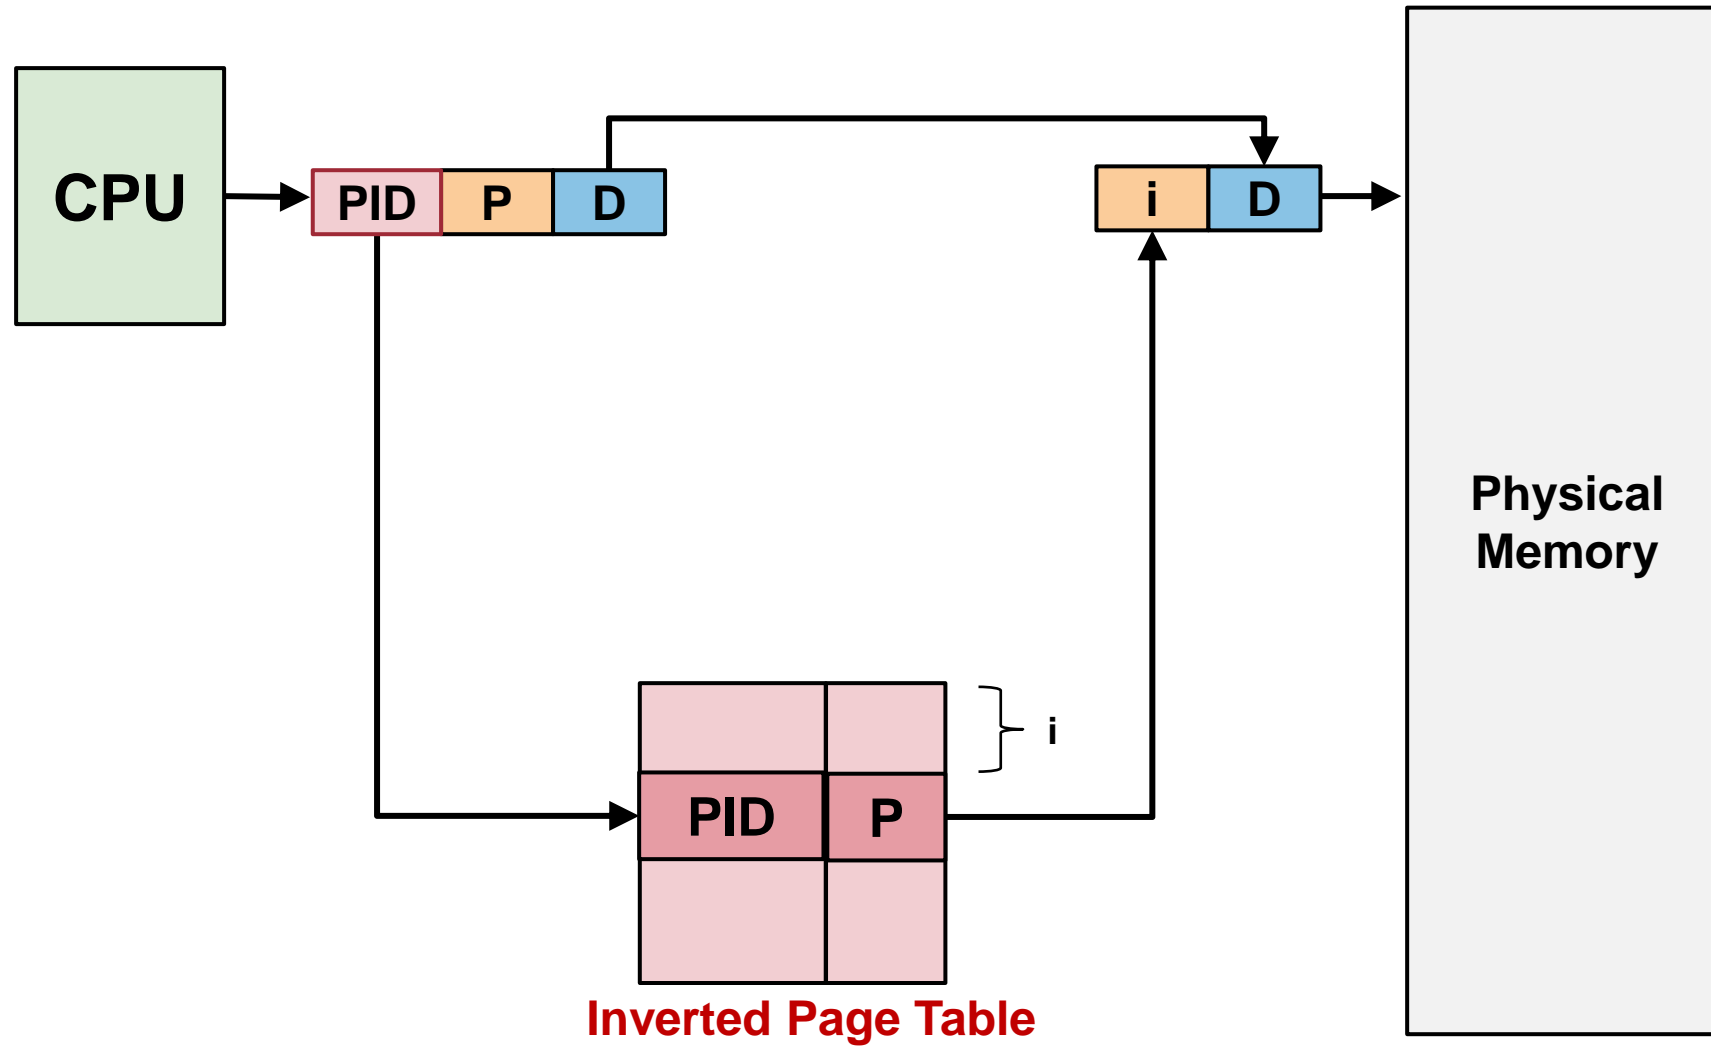
Inverted Page Table: Basic Idea

- Page table is a per-process information
 - With M processes in memory, there are M independent page tables
- Observation:
 - Difficult to find out which frames are occupied, and by which processes
 - Only N physical memory frames can be occupied
 - Out of the M page tables, only N entries are valid
 - Huge waste: $N \ll \text{Overhead of } M \text{ page tables}$
- Idea:
 - Keep a **single** mapping of physical frame to <pid, page#>
 - pid = process id, page# = page number
 - page# is not unique among processes
 - pid + page# can uniquely identify a memory page

Inverted Page Table: Basic Idea (cont)

- In a normal page table, entries are ordered by page number
 - To lookup page X , simply access the X^{th} entry
- In an inverted page table, entries are ordered by frame number
 - To lookup page X , need to search the whole table
- **Advantage:**
 - Huge saving: One table for all processes
 - Frame management is easier and faster
- **Disadvantage:**
 - Slow translation

Inverted Table: Illustration



Who should I kick out next?

PAGE REPLACEMENT ALGORITHMS

Page Replacement Algorithms

- Suppose there is no free physical memory frame during a page fault:
 - ❑ Need to evict (free) a memory page
- When a page is evicted:
 - ❑ Clean page: not modified → no need to write back to storage
 - ❑ Dirty page: modified → need to write back to storage
- Algorithms to find a suitable replacement page
 - ❑ **Optimum (OPT)**
 - ❑ **FIFO**
 - ❑ **Least Recently Used**
 - ❑ **Second-Chance (Clock)**
 - ❑ etc.

Modeling Memory References

- In actual memory references:
 - $\text{Logical Address} = \text{Page Number} + \text{Offset}$
 - However, to study page replacement algorithms
 - Only **page number** is important
- ➔ To simplify discussion, memory references are often modeled as **memory reference strings**, i.e., a sequence of page numbers

Page Replacement Algorithms: Evaluation

Memory access time:

$$T_{access} = (1 - p) * T_{mem} + p * T_{page_fault}$$

- ❑ p = probability of page fault
- ❑ T_{mem} = access time for memory resident page
- ❑ T_{page_fault} = access time if page fault occurs
- Since $T_{page_fault} \gg T_{mem}$
 - ❑ Need to reduce p to keep T_{access} reasonable
- See for yourself, try to find p if:
 - ❑ $T_{mem} = 100\text{ns}$, $T_{page_fault} = 10\text{ms}$, $T_{access} = 120\text{ns}$

Good algorithm should **reduce the total number of page faults**

Optimal Page Replacement (OPT)

- General Idea:

- Replace the page that **will not** be used again for the **longest period of time**
- **Guarantees** minimum number of page faults

- Unfortunately, not feasible:

- Need **future knowledge** of memory references

- Still useful:

- As a base of comparison for other algorithms
- The closer to OPT == better algorithm

Example: OPT (6 Page Faults)

Time	Memory Reference	Frame			Next Use Time			Fault?
		A	B	C				
1	2	2			3	?	?	Y
2	3	2	3		3	9	?	Y
3	2	<u>2</u>	3		6	9	?	
4	1	2	3	1	6	9	?	Y
5	5	2	3	5	6	9	8	Y
6	2	<u>2</u>	3	5	10	9	8	
7	4	4	3	5	?	9	8	Y
8	5	4	3	<u>5</u>	?	9	11	
9	3	4	<u>3</u>	5	?	?	11	
10	2	<u>2</u>	3	5	12	?	11	Y
11	5	2	3	<u>5</u>	?	?	11	
12	2	<u>2</u>	3	5	?	?	?	

FIFO Page Replacement Algorithm

■ General Idea:

- ❑ Memory pages are evicted based on their loading time
- ➔ Evict the oldest memory page

■ Implementation:

- ❑ OS maintains a queue of resident page numbers
 - Remove the first page in queue if replacement is needed
 - Update the queue during page fault trap
- ❑ Simple to implement
 - No hardware support needed

Example: FIFO (9 Page Faults)

Time	Memory Reference	Frame			Loaded at Time			Fault?
		A	B	C				
1	2	2			1			Y
2	3	2	3		1	2		Y
3	2	<u>2</u>	3		1	2		
4	1	2	3	1	1	2	4	Y
5	5	5	3	1	5	2	4	Y
6	2	5	2	1	5	6	4	Y
7	4	5	2	4	5	6	7	Y
8	5	<u>5</u>	2	4	5	6	7	
9	3	3	2	4	9	6	7	Y
10	2	3	<u>2</u>	4	9	6	7	
11	5	3	5	4	9	11	7	Y
12	2	3	5	2	9	11	12	Y

FIFO: Problems

- If number of physical frames increases (e.g., more RAM)
 - Number of page faults should decrease
- FIFO violates this simple intuition!
 - Use 3 / 4 frames to try: **1 2 3 4 1 2 5 1 2 3 4 5**
- Opposite behavior (\uparrow frames \rightarrow \uparrow page faults)
 - Known as **Belady's Anomaly**
- Reason:
 - FIFO does not exploit **temporal locality**

Least Recently Used Page Replacement (**LRU**)

■ General Idea:

- ❑ Make use of temporal locality:
 - Replace the page that has not been accessed in the longest time
- ❑ Expect a page to be reused in a short time window
 - Have not accessed for some time → most likely will not be accessed again

■ Notes:

- ❑ Attempts to approximate the OPT algorithm
 - Gives good results generally
- ❑ Does not suffer from Belady's Anomaly

Example: LRU (7 Page Faults)

Time	Memory Reference	Frame			Last Use Time			Fault?
		A	B	C				
1	2	2			1			Y
2	3	2	3		1	2		Y
3	2	<u>2</u>	3		3	2		
4	1	2	3	1	3	2	4	Y
5	5	2	5	1	3	5	4	Y
6	2	<u>2</u>	5	1	6	5	4	
7	4	2	5	4	6	5	7	Y
8	5	2	<u>5</u>	4	6	8	7	
9	3	3	5	4	9	8	7	Y
10	2	3	5	2	9	8	10	Y
11	5	3	<u>5</u>	2	9	11	10	
12	2	3	5	<u>2</u>	9	11	12	

LRU: Implementation Details

- Implementing LRU is not easy:
 - ❑ Need to keep track of the "last access time" somehow
 - ❑ Need substantial hardware support
- 1. Approach A - **Use a Counter**:
 - ❑ A logical "time" counter, which is incremented for every memory reference
 - ❑ Page table entry has a "time-of-use" field
 - Store the time counter value whenever reference occurs
 - Replace the page with smallest "time-of-use"
 - ❑ Problems:
 - Need to search through all pages
 - "Time-of-use" is forever increasing (overflow possible!)

LRU: Implementation Details (cont)

2. Approach B - Use a "Stack":

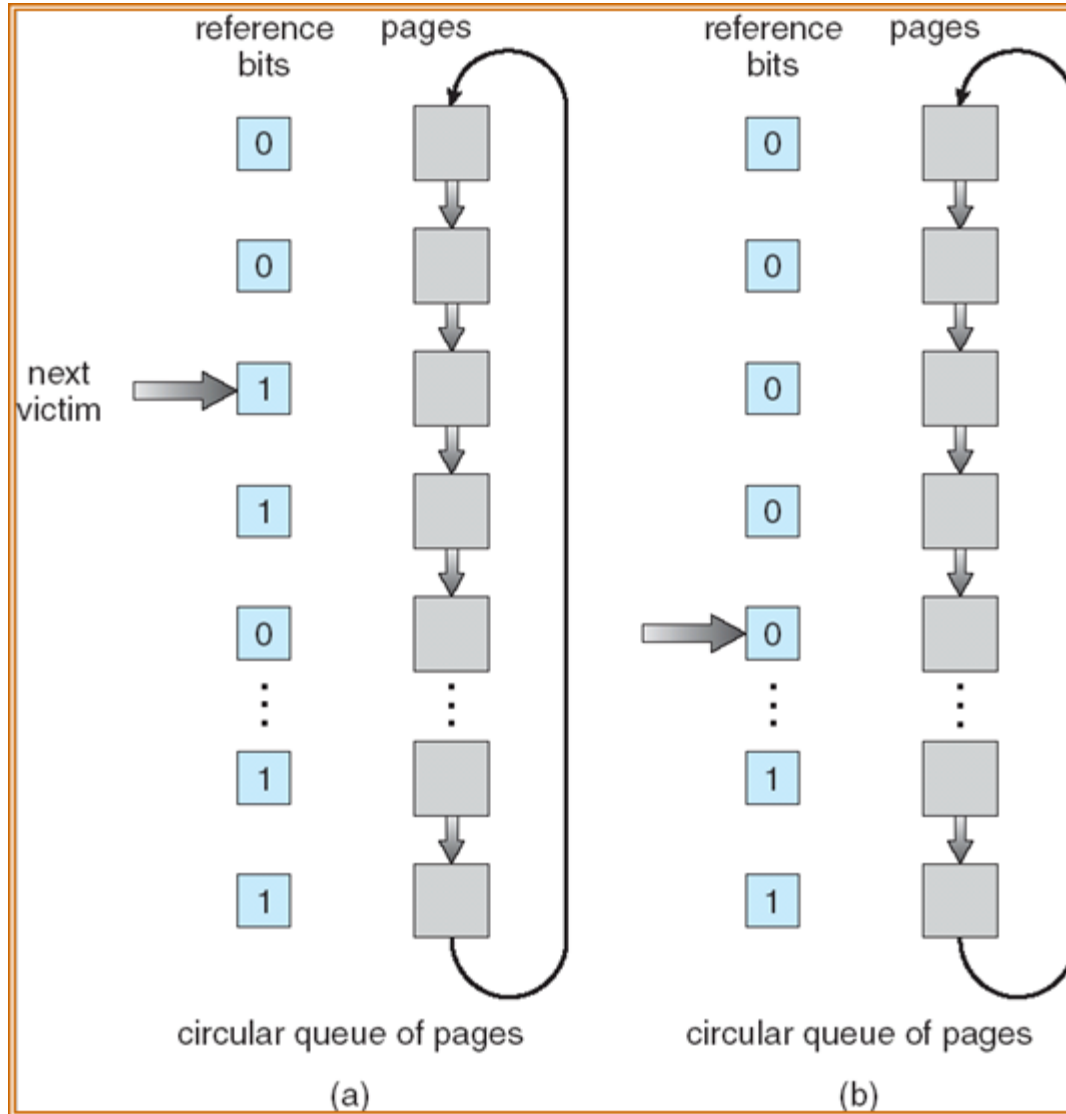
- ❑ Maintain a stack of page numbers
- ❑ If page **X** is referenced
 - Remove from the stack (if entry exists)
 - Push on top of stack
- ❑ Replace the page at the bottom of stack
 - No need to search through all entries
- ❑ Problems:
 - Not a pure stack: Entries can be removed from anywhere in the stack
 - Hard to implement in hardware

Second-Chance Page Replacement (CLOCK)

■ General Idea:

- ❑ Modified FIFO to give a second chance to pages that are accessed
- ❑ Each PTE now maintains a "reference bit":
 - 1 = Accessed, 0 = Not accessed
- ❑ Algorithm:
 1. The oldest FIFO page is selected
 2. If reference bit == 0 → Page is replaced
 3. If reference bit == 1 → Page is given a 2nd chance
 - ❑ Reference bit cleared to 0
 - ❑ Load time reset → page taken as newly loaded
 - ❑ Next FIFO page is selected, go to Step 2
- ❑ Degenerate into FIFO algorithm
 - When all pages have reference bit == 1 (or all == 0)

Second-Chance: Implementation Details



- Use **circular queue** to maintain the pages:
 - With a pointer pointing to the oldest page (the **victim page**)
- To find a page to be replaced:
 - Advance to a page with '0' reference bit
 - Clear the reference bit as pointer passes through

Example: CLOCK(6 Page Faults)

Time	Memory Reference	Frame (with Ref Bit)			Fault?
		A	B	C	
1	2	▶ 2 (0)			Y
2	3	▶ 2 (0)	3 (0)		Y
3	2	▶ <u>2</u> (1)	3 (0)		
4	1	▶ 2 (1)	3 (0)	1 (0)	Y
5	5	2 (0)	5 (0)	▶ 1 (0)	Y
6	2	<u>2</u> (1)	5 (0)	▶ 1 (0)	
7	4	▶ 2 (1)	5 (0)	4 (0)	Y
8	5	▶ 2 (1)	<u>5</u> (1)	4 (0)	
9	3	▶ 2 (0)	5 (0)	3 (0)	Y
10	2	▶ <u>2</u> (1)	5 (0)	3 (0)	
11	5	▶ 2 (1)	<u>5</u> (1)	3 (0)	
12	2	▶ <u>2</u> (1)	5 (1)	3 (0)	

▶ Victim Page

Which process should I favor?

FRAME ALLOCATION

Frame Allocation

■ Consider:

- ❑ There are N physical memory frames
- ❑ There are M processes competing for frames
- ❑ What is the best way to distribute the N frames among M processes?

■ Simple Approaches:

❑ **Equal Allocation:**

- Each process gets N / M frames

❑ **Proportional Allocation:**

- Processes are different in size (memory usage)
- Let size_p = size of process p , $\text{size}_{\text{total}}$ = total size of all processes
- Each process gets $\text{size}_p / \text{size}_{\text{total}} * N$ frames

Frame Allocation and Page Replacement

- The implicit assumption for page replacement algorithms discussed:
 - ❑ Victim pages are selected **among pages of the process** that causes page fault
 - ❑ Known as **local replacement**
- If victim page can be chosen **among all physical frames**:
 - ➔ Process P can take a frame from Process Q by evicting Q's frame during replacement!
 - ❑ Known as **global replacement**

Local vs Global Replacement

■ Local Replacement:

□ Pros:

- Frames allocated to a process remain constant → Performance is stable between multiple runs

□ Cons:

- If frames allocated to a process are not enough → hinder the performance of the process

■ Global Replacement:

□ Pros:

- Allow dynamic self-adjustment between processes
 - Process that needs more frames can get from other

□ Cons:

- Badly behaved process can negatively affect others
- Frames allocated to a process can be different from run to run

Frame Allocation and Thrashing

- Insufficient physical frames → Thrashing in process
 - Heavy I/O to bring non-resident pages into RAM
- Hard to find the right number of frames:
 - If global replacement is used:
 - A thrashing process "steals" page from other processes
 - causes other processes to thrashing (**Cascading Thrashing**)
 - If local replacement is used:
 - Thrashing can be limited to one process
 - But that single process can hog the I/O and degrades the performance of other processes

Finding the right number of frames...

■ Observation:

- ❑ The **set** of pages referenced by a process is relatively constant in a period of time
 - Known as **locality**
- ❑ However, as time passes, the set of pages can change

■ Example:

- ❑ When a function is executing, memory references are likely on:
 - *local variables, parameters, code in that function*
 - these pages define the locality for the function
- ❑ After the function terminates, references will change to another set of pages

Working Set Model

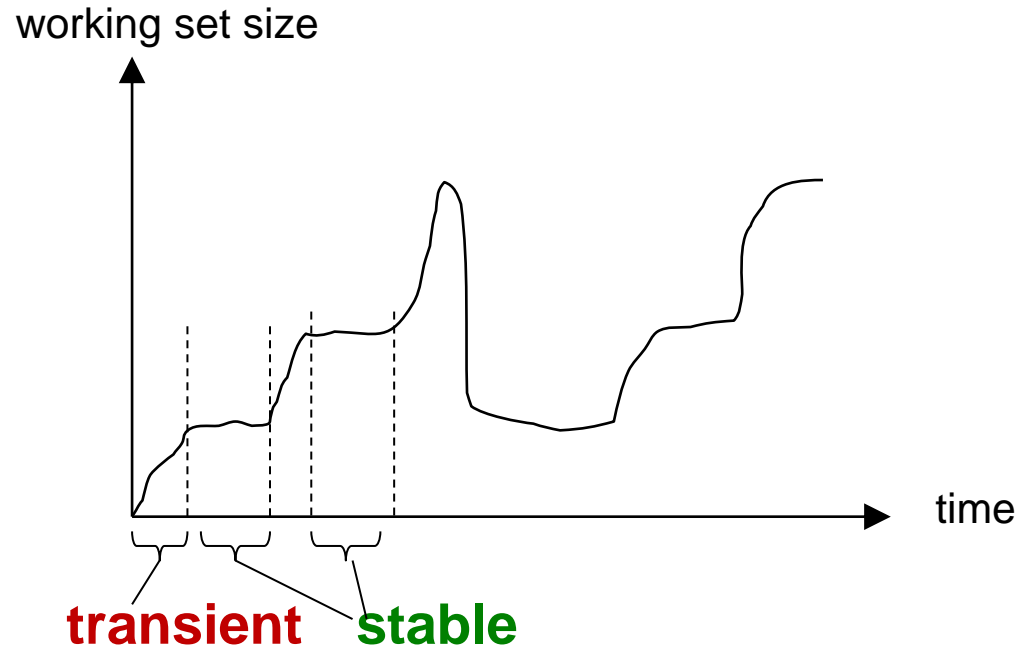
■ Using the observation on locality:

- In a new “locality”:
 - A process will cause page fault for the set of pages
- With the set of pages loaded in frames:
 - No/few page faults until process transits to new locality

■ Working Set Model:

- Defines Working Set Window Δ
 - An interval of time
- $W(t, \Delta)$ = active pages in the interval at time t
- Allocate enough frames for pages in $W(t, \Delta)$ to reduce the number of page faults

Working Set Model: Illustration




Transient region:
working set changing
in size

Stable region:
working set about the
same for a long time

- Accuracy of working set model is directly affected by the choice of Δ
 - Too small: May miss pages in the current locality
 - Too big: May contains pages from different locality

Working Set Model: Illustration

■ Example memory reference strings



■ Assume

□ Δ = an interval of 5 memory references

■ $W(t1, \Delta) = \{1, 2, 5, 6, 7\}$ (5 frames needed)

■ $W(t2, \Delta) = \{3, 4\}$ (2 frames needed)

■ Try using different Δ values

Summary

- Virtual memory
 - The "why" and "how"
- Discussed different aspects of virtual memory management
 - Use different page table structure to reduce page table overhead
 - Use different page replacement algorithms to reduce page fault
 - How frame allocation affects page fault of a process