
File System Management

File System Introduction

Lecture 10

Overview

■ File System

- Definition
- Vs Memory Management
- Motivation

■ File

- Metadata
- Operations

■ Directory

- Directory Structure

■ I/O Scheduling

File System: Motivation

- Physical memory is **volatile**
 - Use external storage to store **persistent** information
- Direct access to the storage media is **not portable**:
 - Dependent on hardware specification and organization
- File System provides:
 - An abstraction on top of the physical media
 - A high level resource management scheme
 - Protection between processes and users
 - Sharing between processes and users

File System: General Criteria

■ **Self-Contained:**

- ❑ Information stored on a media is enough to describe the entire organization
- ❑ Should be able to "plug-and-play" on another system

■ **Persistent:**

- ❑ Beyond the lifetime of OS and processes

■ **Efficient:**

- ❑ Provides good management of free and used space
- ❑ Minimum overhead for bookkeeping information

Memory Management vs File Management

	Memory Management	File System Management
Underlying Storage	RAM	Disk
Access Speed	Constant	Variable disk I/O time
Unit of Addressing	Physical memory address	Disk sector
Usage	Address space for process Implicit when process runs	Non-volatile data Explicit access
Organization	Paging/Segmentation: determined by HW & OS	Many different FS: ext* (Linux), FAT* (Windows), HFS* (Mac OS)etc.

Key Topics

File System Abstraction

- Discuss the logical entities present in file system
- E.g. Files / Directories

File System Implementation

- Common implementation schemes
- Discuss pros/cons
- Case studies

You mean files and folders are not real?

FILE SYSTEM **ABSTRACTIONS**

File System Abstraction

■ File System:

- ❑ Consists of a collection of **files** and **directory structures**
 - **File**: An abstract storage of data
 - **Directory (Folder)**: Organization of files
- ❑ Provides an abstraction of accessing and using the above

■ Look at the two abstractions closely next:

- ❑ **File**
- ❑ **Directory (Folder)**

File: Overview

- Basic Definition
- File Metadata
- File Data
 - File structure
 - Access Methods
- File Operations



File: Basic Description

- Represent a logical unit of information created by process
- An ***abstraction***
 - Essentially an **Abstract Data Type**:
 - A set of common operations with various possible implementation
- Contains:
 - **Data**: Information structured in some ways
 - **Metadata**: Additional information associated with the file
 - Also known as **file attributes**

File Metadata

Name:	A human readable reference to the file
Identifier:	A unique id for the file used internally by FS
Type:	Indicate different type of files E.g. executable, text file, object file, directory etc
Size:	Current size of file (in bytes, words or blocks)
Protection:	Access permissions, can be classified as reading, writing and execution rights
Time, date and owner information:	Creation, last modification time, owner id etc
Table of content:	Information for the FS to determine how to access the file

File Name

- Different FS has different **naming rule**
 - To determine valid file name
- Common naming rule:
 - Length of file name
 - Case sensitivity
 - Allowed special symbols
 - File extension
 - Usual form **Name**.**Extension**
 - On **some** FS, extension is used to indicate **file type**

File Type

- An OS commonly supports a number of **file types**
- Each file type has:
 - An associated set of operations
 - Possibly a specific program for processing
- Common file types:
 - **Regular files:** contains user information
 - **Directories:** system files for FS structure
 - **Special files:** character/block oriented

Two Major Types of Regular Files

■ ASCII files:

- ❑ Example: text file, programming source codes, etc
- ❑ Can be displayed or printed **as is**

■ Binary files:

- ❑ Example: executable, Java class file, pdf file, mp3/4, png/jpeg/bmp etc
- ❑ Have a predefined internal structure that can be processed by specific program
 - JVM to execute Java class file
 - PDF reader for pdf file etc

Distinguishing File Type

1. **Use file extension** as indication:
 - ❑ Used by Windows OS
 - ❑ e.g. **xxx.docx** → Words document
 - ❑ Change of extension implies a change in file type!

2. **Use embedded information** in the file:
 - ❑ Used by Unix
 - ❑ Usually stored at the beginning of the file
 - ❑ Commonly known as **magic number**

File Protection


- Controlled access to the information stored in a file
- **Type of access:**
 - ❑ **Read:** Retrieve information from file
 - ❑ **Write:** Write/Rewrite the file
 - ❑ **Execute:** Load file into memory and execute it
 - ❑ **Append:** Add new information to the end of file
 - ❑ **Delete:** Remove the file from FS
 - ❑ **List:** Read metadata of a file

File Protection: How?

- Most common approach:
 - Restrict access base on the user identity
- Most general scheme:
 - **Access Control List**
 - A list of user identity and the allowed access types
 - **Pros:** Very customizable
 - **Cons:** Additional information associated with file
- A common condensed file protection scheme is discussed next

File Protection: Permission Bits

- Classified the users into three classes:
 1. **Owner:** The user who created the file
 2. **Group:** A set of users who need similar access to a file
 3. **Universe:** All other users in the system
- Example (Unix)
 - ❑ Define permission of three access types (**R**ead/**W**rite/**E**xecute) for the 3 classes of users
 - ❑ Use "ls -l" to see the permission bits for a file



File Protection: Access Control List

- In Unix, Access Control List (ACL) can be:
 - ❑ Minimal ACL (the same as the permission bits)
 - ❑ Extended ACL (added **named users / group**)

```
$ getfacl exampleDir
```

```
# file: exampleDir
```

```
# owner: ccris
```

```
# group: compsc
```

```
user::rwx
```

```
user:sooyj:rwx
```

```
group::r-x
```

```
group:cohort21:rwx
```

```
mask::rwx
```

```
other:---
```

"getfacl" is the command
to get ACL information

Permission for
Specific User

Permission for
Specific Group

Permission
"upperbound"

Operations on File Metadata

■ **Rename:**

- ❑ Change filename

■ **Change attributes:**

- ❑ File access permissions
- ❑ Dates
- ❑ Ownership
- ❑ etc

■ **Read attribute:**

- ❑ Get file creation time

File Data: **Structure**

■ **Array of bytes:**

- ❑ The traditional Unix view
- ❑ No interpretation of data: **just raw bytes**
- ❑ Each byte has a unique **offset (distance)** from the file start

■ **Fixed length records:**

- ❑ Array of records, can grow/shrink
- ❑ Can jump to any record easily:
 - Offset of the **Nth** record = size of Record * (N-1)

■ **Variable length records**

- ❑ Flexible but harder to locate a record

File Data: **Access Methods**

■ **Sequential Access:**

- ❑ Data read in order, starting from the beginning
- ❑ Cannot skip but can be rewound

■ **Random Access:**

- ❑ Data can be read in any order
- ❑ Can be provided in two ways:
 1. **Read (Offset)** : Every read operation explicitly state the position to be accessed
 2. **Seek (Offset)** : A special operation is provided to move to a new location in file
 - E.g. Unix and Windows uses (2)

File Data: Access Methods (cont)

■ Direct Access:

- ❑ Used for file contains fixed-length records
- ❑ Allow ***random access to any record directly***
- ❑ Very useful where there is a large amount of records
 - e.g. In database
- ❑ The basic random access method can be view as a special case:
 - Where each record == one byte

File Data: Generic Operations

Create:	New file is created with no data
Open:	Performed before further operations To prepare the necessary information for file operations later
Read:	Read data from file, usually starting from current position
Write:	Write data to file, usually starting from current position
Repositioning:	Also known as seek Move the current position to a new location No actual Read/Write is performed
Truncate:	Removes data between specified position to end of file

File Operations as System Calls

- OS provides file operations as **system calls**:
 - ❑ Provide protection, concurrent and efficient access
 - ❑ Maintain information
- Information kept for an opened file:
 - ❑ **File Pointer**: Current location in file
 - ❑ **Disk Location**: Actual file location on disk
 - ❑ **Open Count**: How many times has this file opened?
 - Useful to determine when to remove the entry in table

File Information in the OS

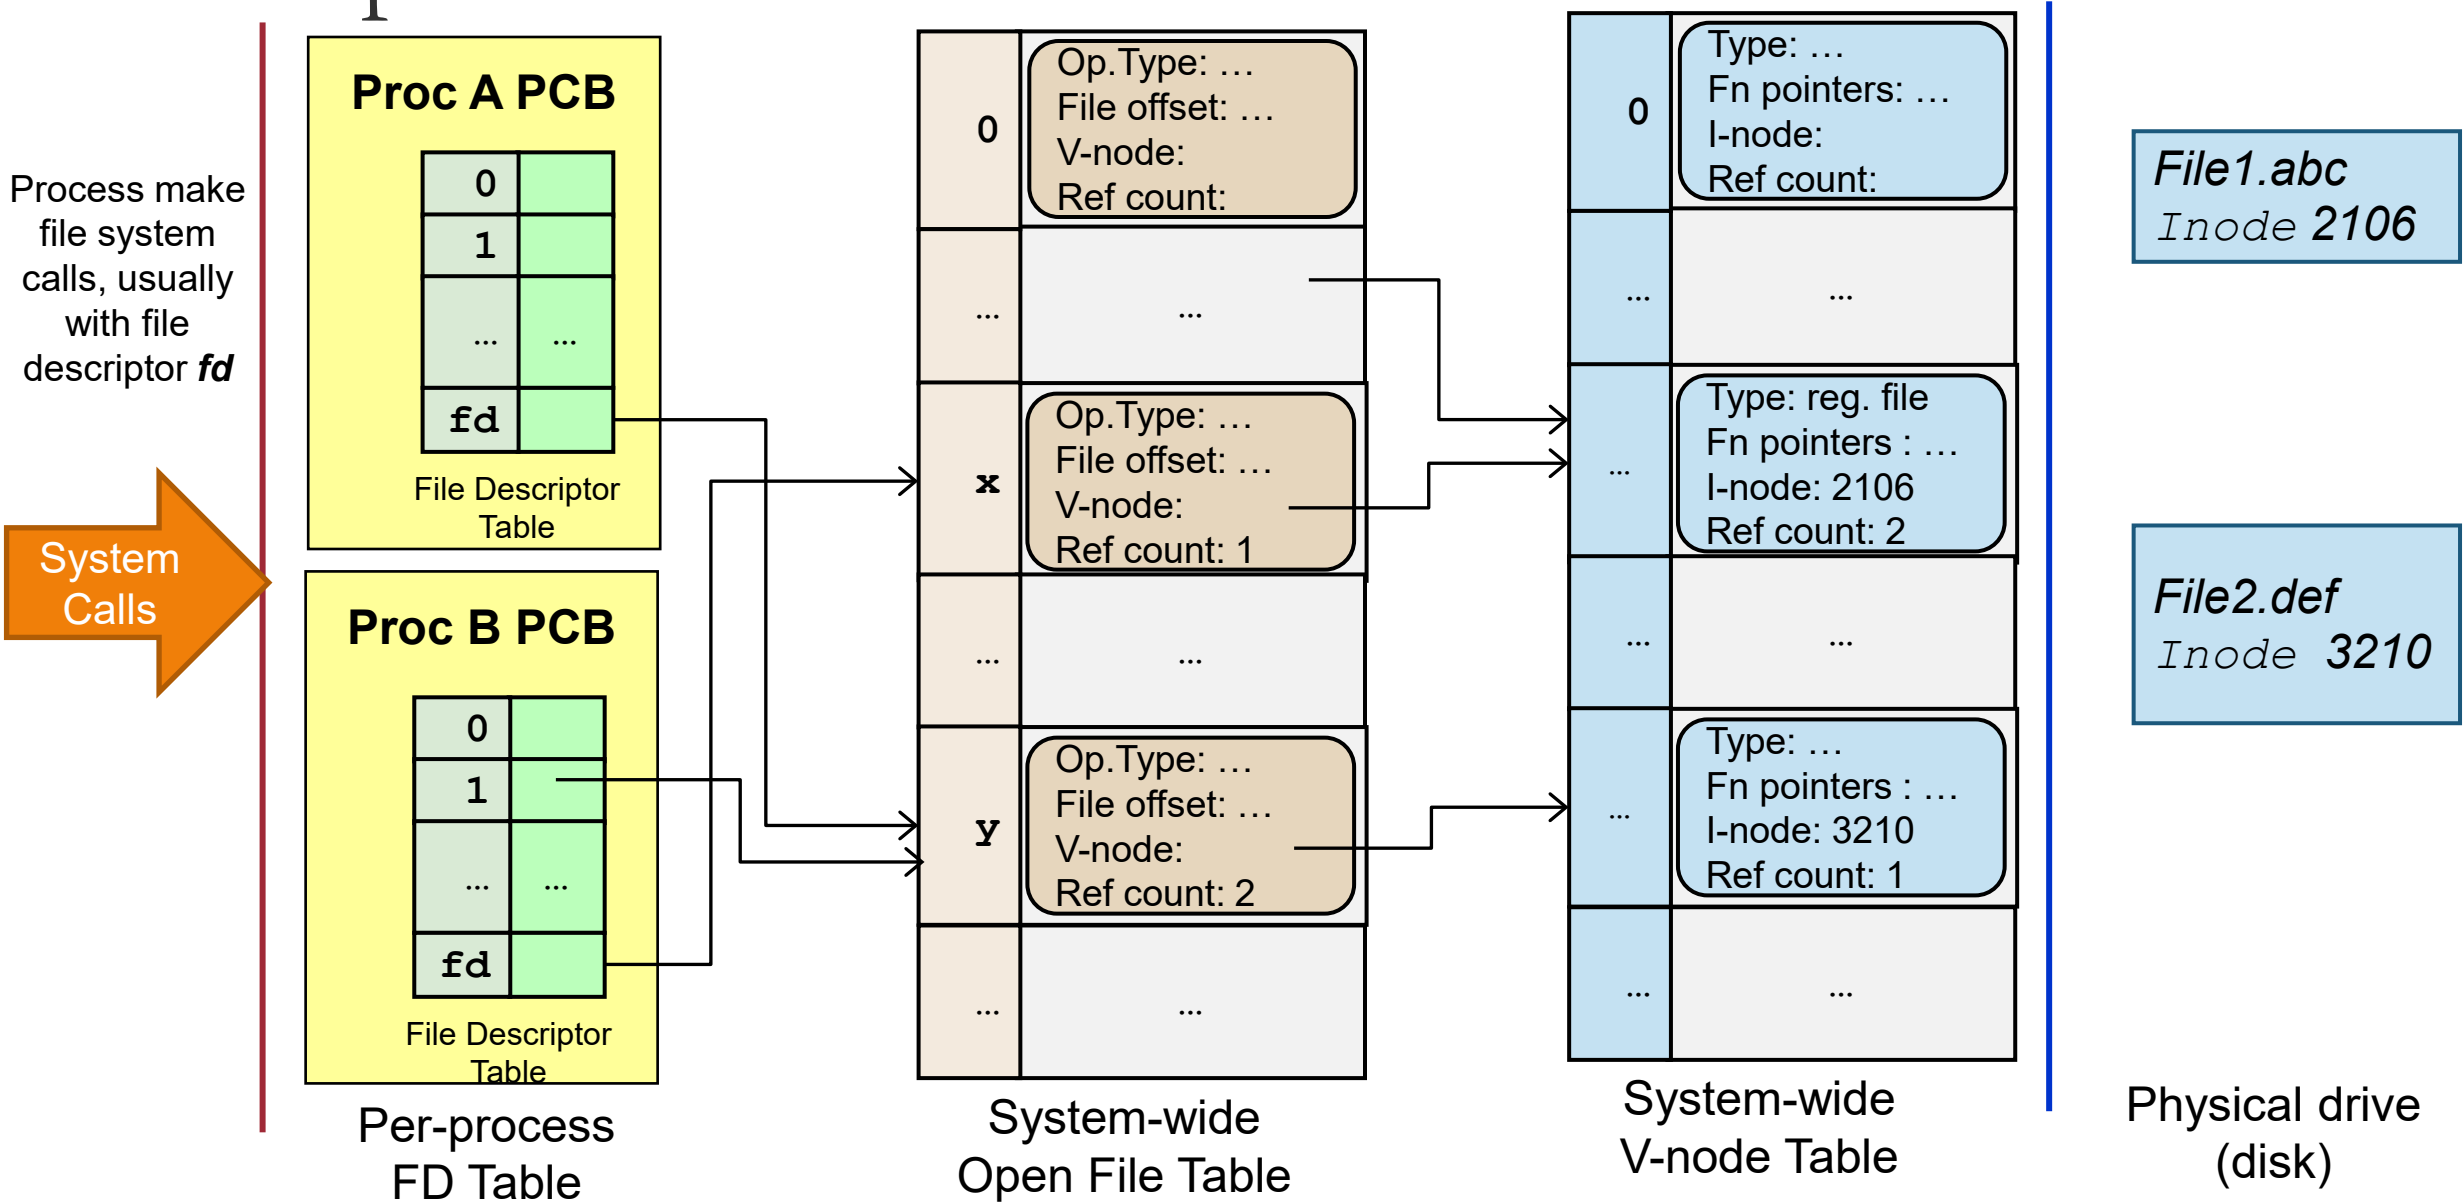
■ Consider:

- ❑ Several processes can open the same file
- ❑ Several different files can be opened at any time

■ Common approach – uses 3 tables:

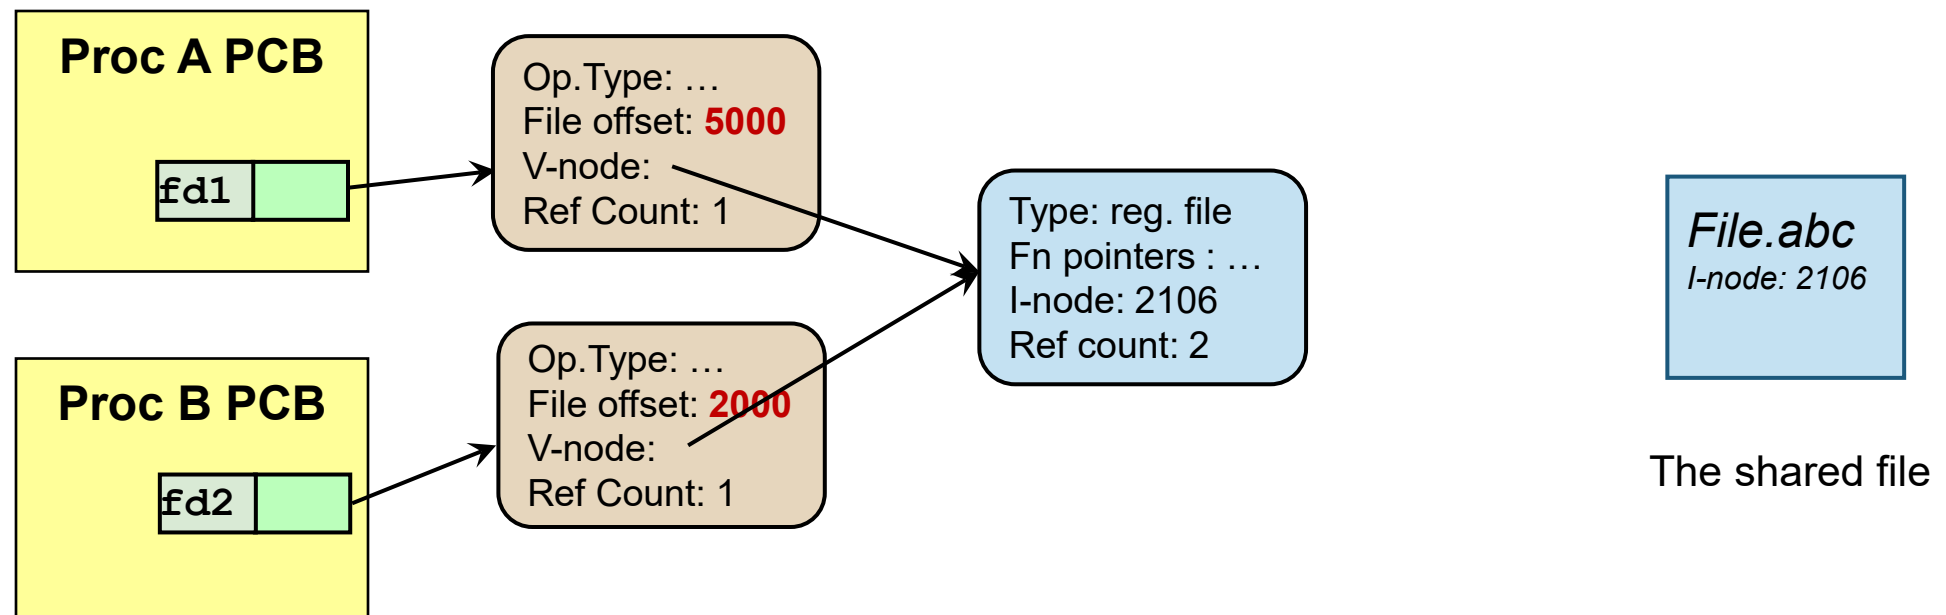
- ❑ Per-process open-file table:
 - To keep track of the open files for a process
 - Each entry points to the **system-wide open-file table** entries
- ❑ System-wide open-file table:
 - To keep track of all the open files in the system
 - Each entry points to a **V-node** entry
- ❑ System-wide V-node(virtual node) table
 - To link with the file on physical drive
 - Contains the information about the physical location of the file.

File Operations: Unix Illustration



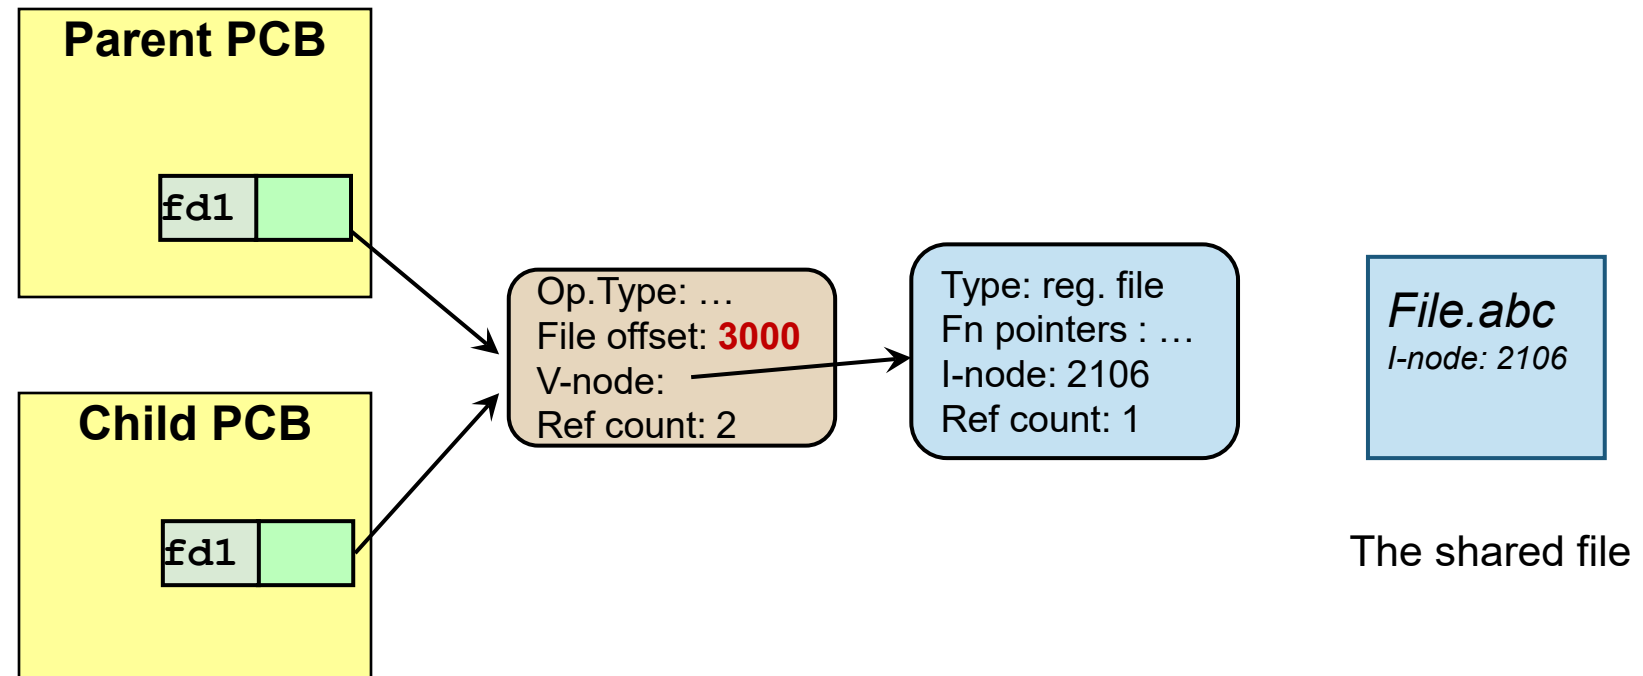
Process Sharing File in Unix: Case 1

- A file is opened twice from two processes:
 - ❑ 2 file descriptors
 - ❑ 2 entries in the system-wide open file table
 - ❑ I/O can occur at independent offsets
- When:
 - ❑ Two process open the same file
 - ❑ Same process open the file twice



Process Sharing File in Unix: Case 2

- Two file descriptors pointing to the same entry in the system-wide open file table
 - Only one offset → I/O changes the offset for the other process
- When:
 - fork() after file is opened
 - dup () within the same process



Just your regular folders

DIRECTORY

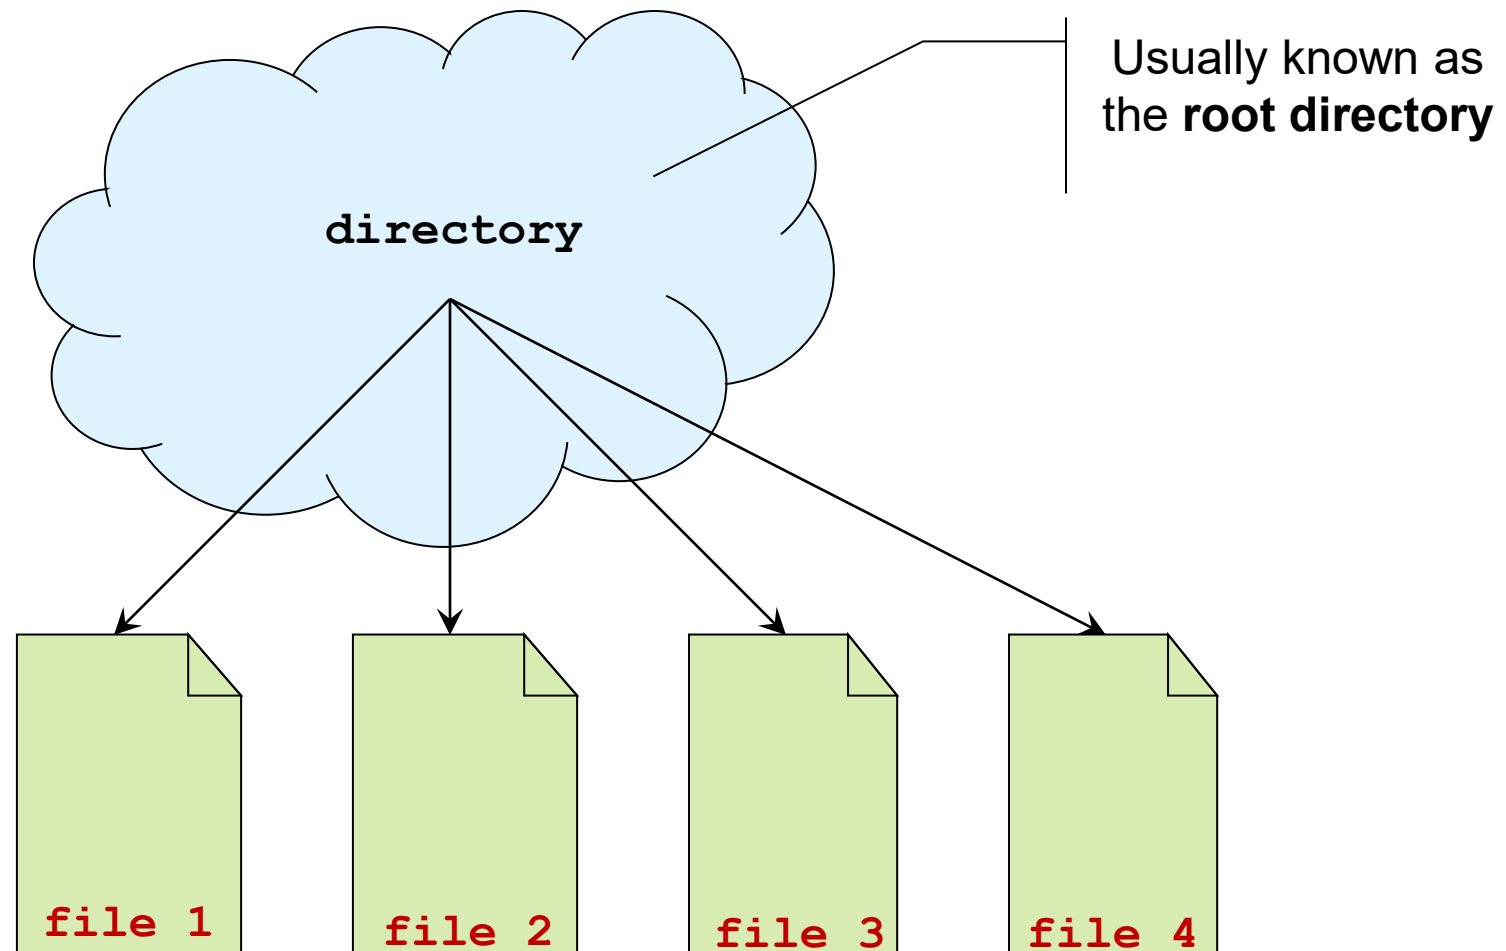


Directory: Basics

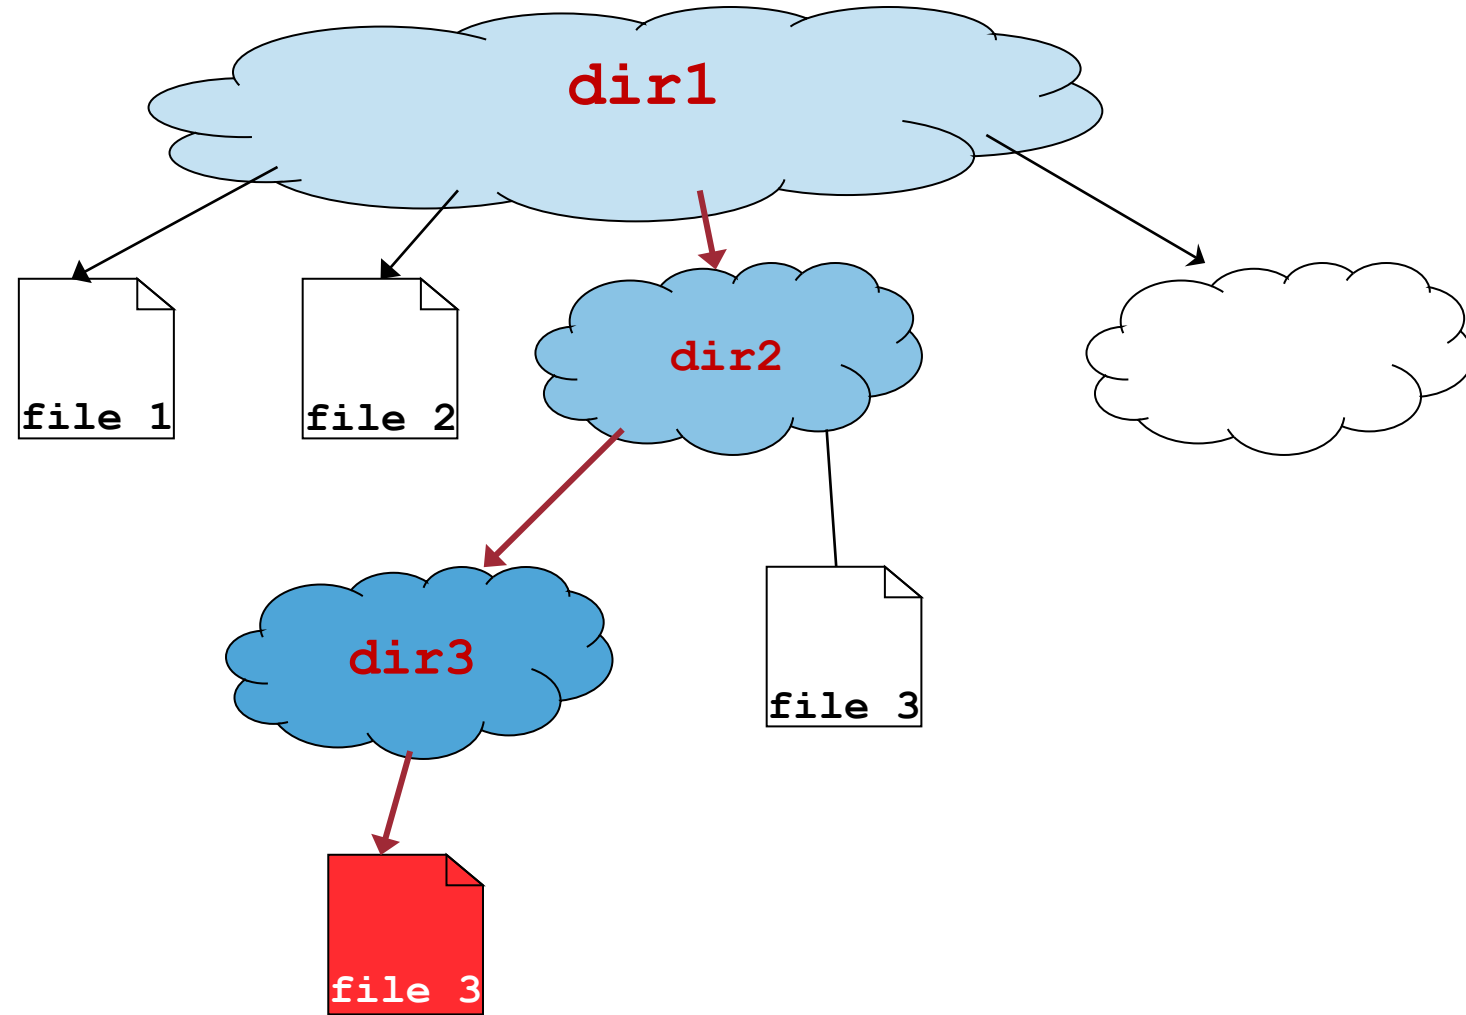
- **Directory** (**folder**) is used to:
 1. Provide a logical grouping of files
 - The user view of directory
 2. Keep track of files
 - The actual system usage of directory

- Several ways to structure directory:
 - ❑ Single-Level
 - ❑ Tree-Structure
 - ❑ Directed Acyclic Graph (DAG)
 - ❑ General Graph

Directory Structure: **Single-Level**



Directory Structure: **Tree-Structured**



Directory Structure: Tree-Structured

■ **General Idea:**

- ❑ Directories can be recursively embedded in other directories
- ❑ Naturally forms a tree structure

■ Two ways to refer to a file:

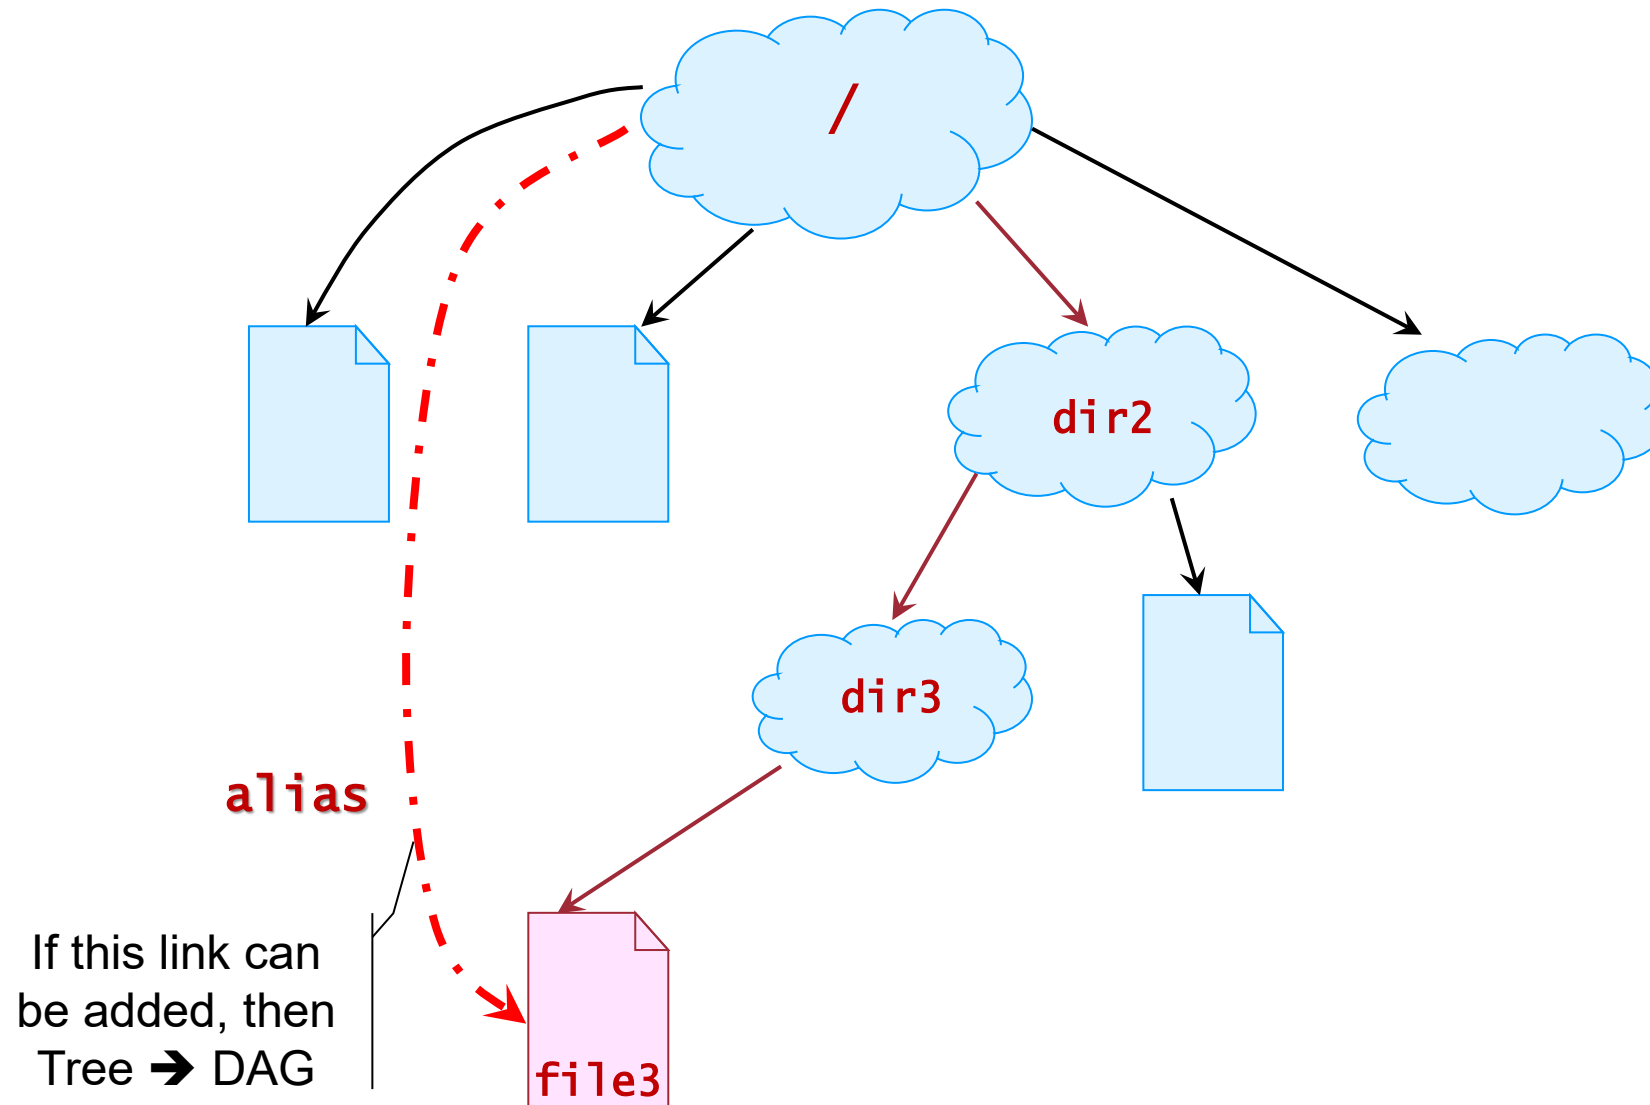
❑ **Absolute Pathname:**

- Directory names followed from root of tree + final file
- i.e. the Path from root directory to the file

❑ **Relative Pathname:**

- Directory names followed from the **current working directory (CWD)**
- CWD can be set explicitly or implicitly changed by moving into a new directory under shell prompt

Directory Structure: **DAG**



Directory Structure: **DAG**

- If a **file** *can be shared*:
 - Only one copy of actual content
 - "Appears" in multiple directories
 - With different path names
- Then tree structure → DAG
- Two implementations in Unix:
 - **Hard Link**
 - Not allowed for directories
 - **Symbolic Link**
 - This has an "interesting" effect....

DAG: **Unix Hard Link**

■ **Consider:**

- ❑ Directory **A** is the owner of file **F**
- ❑ Directory **B** wants to share **F**

■ **Hard Link:**

- ❑ **A** and **B** has **separate pointers** point to the actual file **F** in disk

- ❑ **Pros:**

- Low overhead, only pointers are added in directory

- ❑ **Cons:**

- Deletion problems:

- ❑ e.g. If B deletes F? If A deletes F?
- ❑ Ref. count is needed

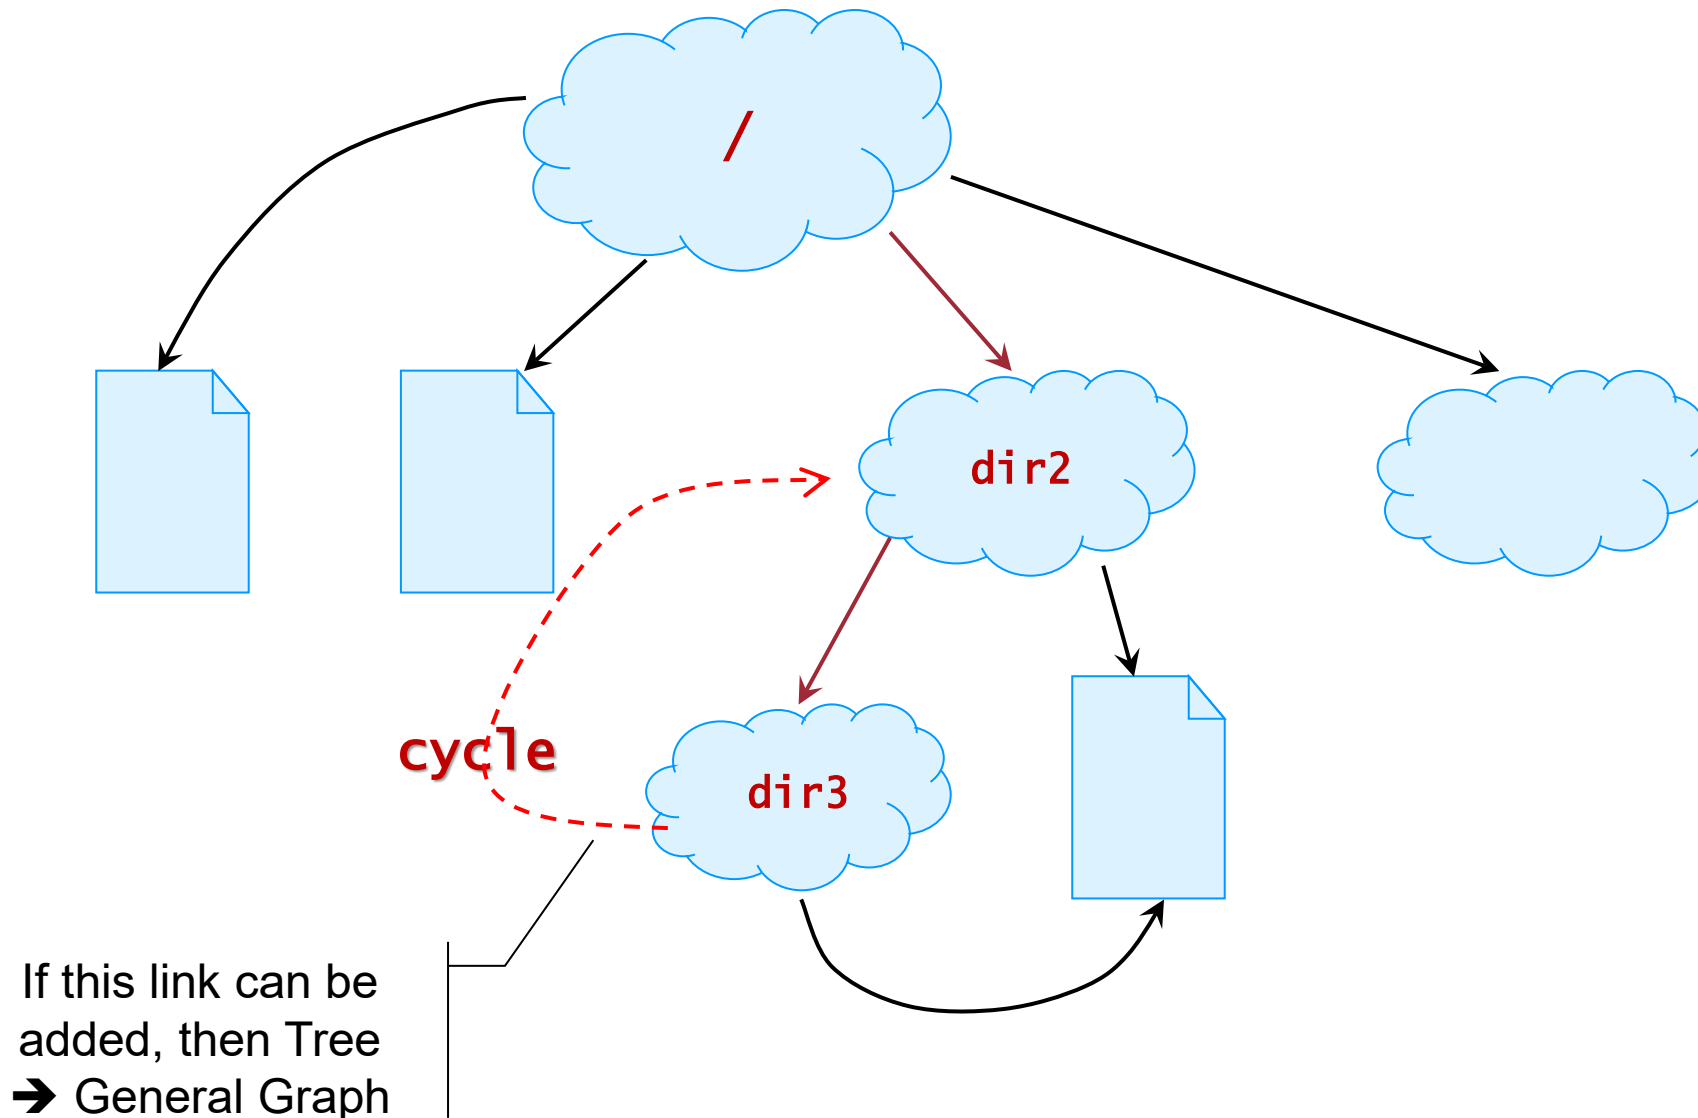
- ❑ Unix Command: " **ln** "

DAG: Unix Symbolic Link

■ Symbolic Link:

- ❑ The symbolic link is a ***special link file***, **G**
 - **G** contains the path name of **F**
- ❑ When **G** is accessed:
 - Find out where is **F**, then access **F**
- ❑ **Pros:**
 - Simple deletion:
 - ❑ If the symbolic link is deleted: **G** deleted, not **F**
 - ❑ If the linked file is deleted: **F** is gone, **G** remains (but not working)
- ❑ **Cons:**
 - Larger overhead:
 - ❑ Special link file take up actual disk space
- ❑ Unix Command: "**ln -s**"

Directory Structure: General Graph




Directory Structure: General Graph

- General Graph Directory Structure is ***not desirable***:
 - **Hard to traverse**
 - Need to prevent infinite looping
 - **Hard to determine when to remove a file/directory**
- **In Unix:**
 - Symbolic link is allowed to link to directory
 - General Graph **can be created**

I'm afraid you have to wait.....

I/O SCHEDULING

Magnetic Disk in One Glance



Disk Scheduling: The Problem

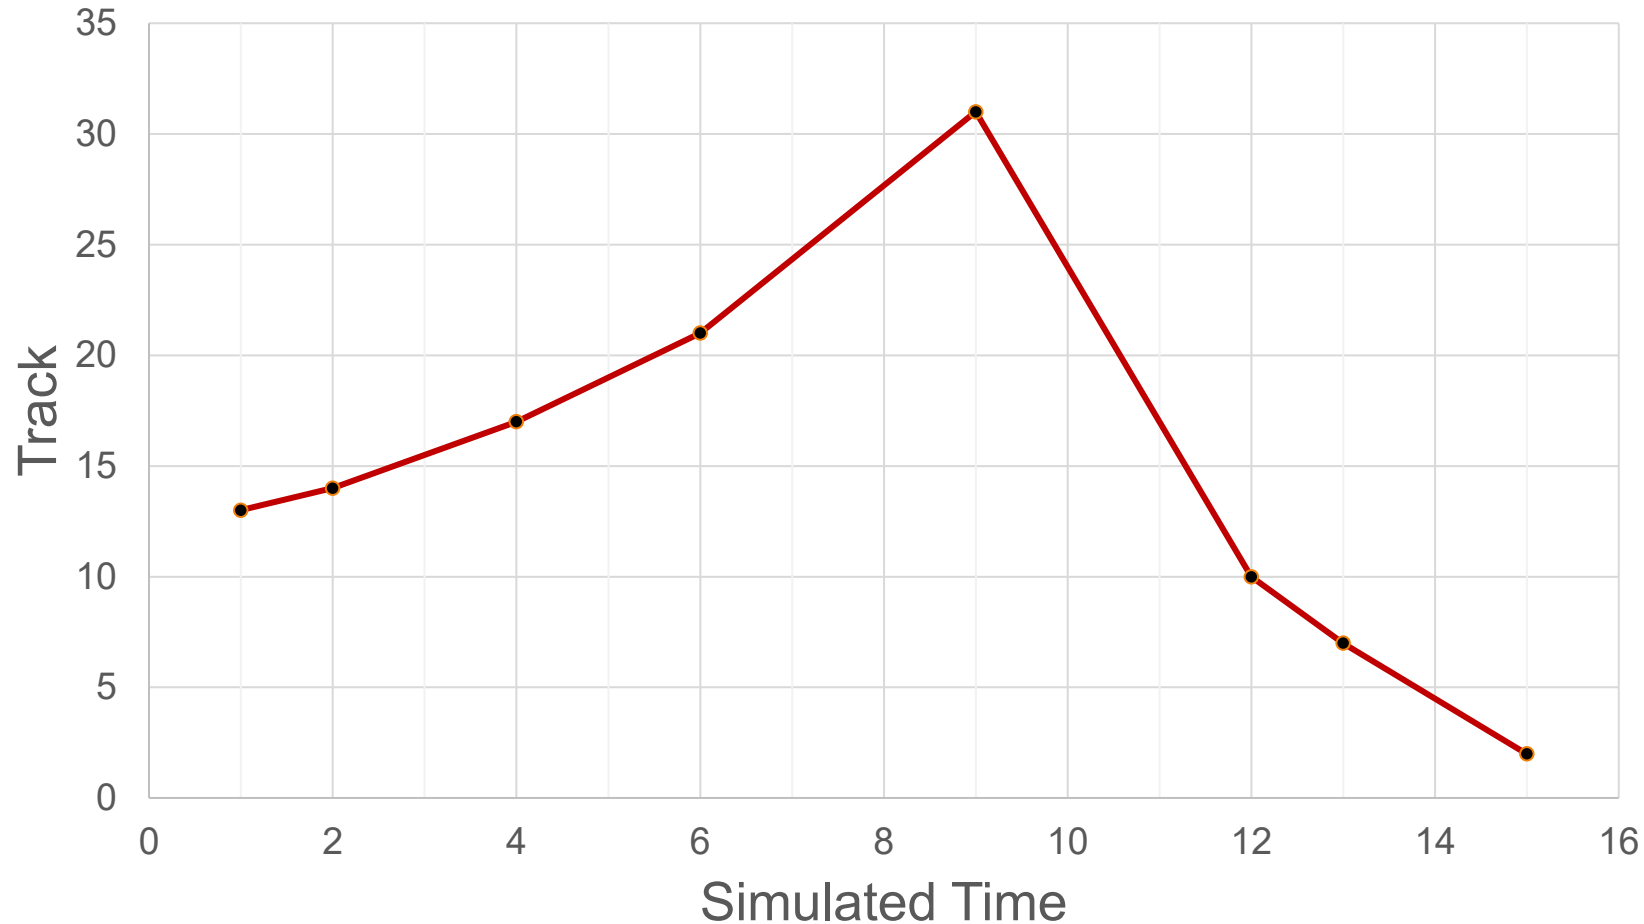
- Due to the significant seek and rotational latency, OS should schedule the disk I/O requests
- I/O (disk) scheduling:
 - ❑ Intention of reducing **overall waiting time**
 - ❑ As rotational latency is hard to mitigate, we focus on reducing the **seeking time**
 - ❑ Balance the need for high throughput while trying to fairly share I/O requests amongst processes

Disk Scheduling: **Algorithms**

- Consider the following disk I/O requests indicated by only the **track number (magnetic disks)**:
 - ❑ 13, 14, 2, 10, 17, 31, 21, 7
- A few obvious candidates:
 - ❑ **FCFS**
 - ❑ **SSF (Shortest Seek First)**
 - "SJF" modified for the disk context
 - ❑ The **SCAN** family (aka **Elevator**):
 - Bi-Direction [Innermost \leftrightarrow Outermost] (SCAN)
 - 1-Direction [Outermost \rightarrow Innermost] (C-SCAN)
- Very intuitive: Imagine the tracks are floors in a building, and the disk head is the elevator servicing the floors (Figure out the algorithm before lecture 😊)

SCAN: Disk Head Movement

- disk I/O requests indicated by only the **track number** :
[13, 14, 2, 10, 17, 31, 21, 7]



I/O Scheduling: **Newer Algorithms**

- **Deadline** - 3 queues for I/O requests:
 - Sorted
 - Read FIFO - read requests stored chronologically
 - Write FIFO - write requests stored chronologically
- **noop (No-operation)** - no sorting
- **cfq (Completely Fair Queueing)** - time slice and per-process sorted queues
- **bfq (Budget Fair Queuing) (Multiqueue)** - fair sharing based on the number of sectors requested

Summary

- Covered basics of file system from a user point of view
- Understand the basic requirements of a FS
- Understand the components of a FS:
 - File and Directory
- Discussed OS responsibility in I/O scheduling

For your reference only

UNIX FILE OPERATIONS

File Operations Example: Unix System Calls

■ Header Files:

- ❑ `#include <sys/types.h>`
- ❑ `#include <sys/stat.h>`
- ❑ `#include <fcntl.h>`

■ File related Unix System Calls

- ❑ `open()`, `read()`, `write()`, `lseek()`, `close()`

■ General Information:

- ❑ Opened file has an identifier
 - **File Descriptor:** Integer
 - Used for other operations
- ❑ File is access on a byte-by-byte basis
 - No interpretation of data

Opening Files: `open ()`

- Function Call:

```
int open( char *path, int flags )
```

- Return:

- ❑ `-1` : Failed to open file
- ❑ `>=0` : **file descriptor**, a unique index for opened file

- Parameters:

- ❑ `path`: File path
- ❑ `flags`: Many options can be set using bit-wise-OR
 - Read, Write or Read+Write mode
 - Truncation, Append mode
 - Create file if no exists
 - ... Many many more 😊

Opening Files: `open()` (cont)

- Example:

```
int fd;    //file descriptor
```

```
//Open an existing file for read only
```

```
fd = open( "data.txt", O_RDONLY );
```

```
//Create the file if not found, open for read + write
```

```
fd = open("data.txt", O_RDWR | O_CREAT );
```

- By convention:

- Default file descriptors:

- **STDIN** (0), **STDOUT** (1), **STDERR** (2)

Read Operation: `read()`

- Function Call:

```
int read(int fd, void *buf, int n)
```

- Purpose:

- reads up to *n* bytes from current offset into buffer *buf*

- Return:

- number of bytes read, can be $0 \dots n$
- $< n$: end of file is reached

- Parameters:

- *fd*: file descriptor (must be opened for read)
- *buf*: An array large enough to store *n* bytes

- `read()` is ***sequential read***:

- starts at current offset and increments offset by bytes read

Write Operation: `write()`

- Function Call:

```
int write(int fd, void *buf, int n)
```

- Purpose:

- writes up to `n` bytes from current offset from buffer `buf`

- Return:

- `-1`: Error
- `>= 0`: Number of bytes written

- Parameters:

- `fd`: file descriptor (must be opened for write)
- `buf`: An array of at least `n` bytes with values to be written

- Possible errors:

- exceeds file size limit, quota, disk space, etc.

- `write()` is ***sequential write***:

- starts at current offset and increments offset by bytes written
- can increase file size beyond EOF → append new data

Repositioning: `lseek()`

- Function Call:

`off_t lseek(int fd, off_t offset, int whence)`

- Purpose:

- Move current position in file by **offset**

- Return:

- **-1**: Error
- **>= 0**: Current offset in file

- Parameters:

- **fd**: file descriptor (must be opened)
- **offset**: positive = move forward, negative = move backward
- **whence**: Point of reference for interpreting the **offset**
 - **SEEK_SET**: absolute **offset** (count from the file start)
 - **SEEK_CUR**: relative **offset** from current position (+/-)
 - **SEEK_END**: relative **offset** from end of file (+/-)

- Can seek anywhere in file, even beyond end of existing data

Closing Files: `close()`

- Function Call:

```
int close( int fd )
```

- Return:

- ❑ `-1`: Error
- ❑ `0`: Successful

- Parameters:

- ❑ `fd`: file descriptor (must be opened)

- With `close()`:

- ❑ `fd` no longer used anymore
- ❑ Kernel can remove associated data structures
- ❑ The identifier `fd` can be reused later

- By default:

- ❑ Process termination automatically closes all open files