

## Section 1: MCQ ( 2 marks each )

MCQ 1 and 2 based on the following code fragment:

```
int main( )
{
    int cid[9]= {0}; //init the entire array to zeros

    for (i = 0; i < 9; i++){
        cid[i] = fork();
        <Point Alpha>
        if ( <Condition> ) {
            ...some statements...
            exit(0);
        }
    }

    return 0;
}
```


1. In loop `i = 3`, how many zeroes are there in the `cid[]` array for the child process at **<point alpha>**? The child process refers to the process just created by the `fork()` system call in that loop iteration.
- a. 6
  - b. 7
  - c. 8
  - d. 9
  - e. None of the above



2. Which of the following condition(s) can be used for <Condition> if we want to create a total of 9 additional processes (i.e. not counting the original process)?

- i. `cid[i] == 0`
- ii. `cid[i] != 0`
- iii. `cid[0] == 0`


- a. (i) only.
- b. (i) and (ii) only.
- c. (ii) and (iii) only.
- d. (i), (ii) and (iii).
- e. None of the above.



3. Which of the following statment(s) regarding **zombie** process is **TRUE**?

- i. Zombie process takes up a slot in the OS PCB table.
- ii. Zombie process is created so that `wait()` system call can be implemented properly.
- iii. A user command running in the **background** under a shell interpreter can become a zombie process.

- a. (i) only.
- b. (i) and (ii) only.
- c. (ii) and (iii) only.
- d. (i), (ii) and (iii).
- e. None of the above.



[REDACTED]

4. Which of the following statement(s) regarding **Unix Shared Memory IPC** is **TRUE**?

- i. Shared memory region created by program **P** can stay around after **P** exited.
  - ii. Shared memory region is identified by a pointer (memory address).
  - iii. Shared memory region can be accessed by any process (including process from other users).
- a. (i) only.
  - b. (i) and (iii) only.
  - c. (ii) and (iii) only.
  - d. (i), (ii) and (iii).
  - e. None of the above.

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

5. Given the following pseudo code:

Code A	Code B
<pre>wait( S ); &lt;do some work&gt; signal( S );</pre>	<pre>&lt;heavy computation&gt;</pre>

Which of the following setup can potentially cause **priority inversion**?

- A high priority task running code B and a lower priority task running code A.
- A high priority task running code A and a lower priority task running code B.
- The highest and lowest priority tasks running code A and a middle priority task running task B.
- The highest and lowest priority tasks running code B and a middle priority task running task A.
- None of the above



6. Ms. Raycond coded the following function:

```
int globalVar = 0;    //shared among all threads
void* doSum( void* arg)
{
    int i, localVar = 0;

    for (i = 0; i < 50000; i++){
        localVar++;
    }
    globalVar += localVar;
}
```

If we spawn **two threads** to work on the **doSum()** function and **wait for them to finish**, what is the **most accurate** description of the program behavior?

- The program is now deterministic with the globalVar equal to 100000 for all runs.
- The program still exhibits race condition. The globalVar value can be 0, 50000 or 100000.
- The program still exhibits race condition. The globalVar value can be 50000 or 100000.
- The program still exhibits race condition. The globalVar value can be any positive number.
- None of the above



## Section 2: Short Questions (28 marks)


### Question 7 (7 marks )

Consider an array **A** of **N integer values**, a task can execute two operations: i) **IN**: read and remove one of the N values and ii) **OUT**: write into one of the N values. Below is an attempt to use semaphore to synchronize the tasks in operating on the array values:

<pre>Semaphore mutex = 1;    //binary semaphore int A[N];    //shared array</pre>	
<pre>int IN( int idx ) {     int result;     wait( mutex );     result = A[idx];     //"remove" value     A[idx] = -1;     signal( mutex );     return result; }</pre>	<pre>void OUT( int idx, int newValue ) {     wait( mutex );     A[idx] = newValue;     signal( mutex ); }</pre>

- [2 marks] Briefly describe one shortcoming of this implementation.
- [5 marks] Give an implementation that solve the shortcoming in (a). Note that you can only:
  - Introduce / modify the semaphore declaration and initialization.
  - Add **only** wait / signal to the IN and OUT operation.





### Question 8 ( 6 marks )

The **responsiveness** of a scheduling algorithm refers to how soon can a newly created task receives its **first share of CPU time**. The following questions focus on a newly created task  $T_{\text{new}}$  added into an environment where **there are  $N$  ( $N > 0$ ) ready to run tasks**. Restrict your answer to scheduling algorithms discussed in the course so far.

- a. [4 marks] Give **two** algorithms that can be responsive. Briefly explain / describe how the algorithms enable responsiveness.
- b. [2 marks] Give **one** algorithm that is irresponsible. Similarly explain / describe how the algorithm prohibits responsiveness.

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

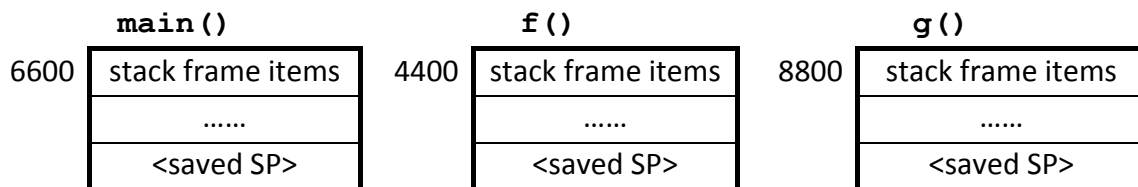


**Question 9 ( 7 marks )**


Instead of using the stack memory, Mr. S. Penn suggested the following alternative to support function invocation:


- During compilation, all functions will be allocated a **predetermined** memory location to store their stack frame. (Similar to how global variable has a fixed memory location). There is no change to the stack frame structure.
- The stack pointer (SP) / frame pointer (FP) can simply points to these predetermined locations during the function execution.

For example, suppose there is a program with three functions: **main()**, **f()** and **g()**. The compiler allocated the following locations for their respective stack frame:



- [1 mark] Suppose **main()** calls **g()**, show the value(s) of the **<saved SP>** in the relevant stack frame(s) when **g()** is still executing. Put a “---” for irrelevant **<saved SP>**. If you think it is not possible, please **cross out the stack frames** as an indication.
- [2 marks] Suppose **main()** calls **f()** which calls **g()**, show the value(s) of the **<saved SP>** in the relevant stack frame(s) when **g()** is still executing. Put a “---” for irrelevant **<saved SP>**. If you think it is not possible, please **cross out the stack frames** as an indication.
- [2 marks] What are the conditions for this implementation scheme to work?
- [2 marks] Give one example where this implementation scheme fails?





### Question 10 ( 8 marks )



Given the following two multi-threaded processes with their respective execution behavior:


$P_1$	$T_1$	C1, IO1, C1, IO1, C1, ..... repeats
	$T_2$	C20
$P_2$	$T_1$	C1, IO1, C1, IO1, C1, ..... repeats
	$T_2$	C20

If we use **the standard 3-level MLFQ** scheduling with a time quantum of **2 time units**, answer the following. Note that whenever there is a need to order the processes / threads, you can assume  $P_1$  is ordered before  $P_2$  and the respective  $T_1$  is ordered before  $T_2$ .

- a. [4 marks] Suppose the threads are implemented as **kernel threads**, give the first 8 time units of the CPU schedule. Remember to indicate both the process number and thread number, e.g.  $P_2T_1$ .
- b. [4 marks] Suppose the threads are implemented as **user threads**, give the first 8 time units of the CPU schedule. Remember to indicate both the process number and thread number, e.g.  $P_2T_1$ . You should try to give fair CPU share to the threads within the same process as much as possible.

For both questions, give **important assumptions you have made (if any)**. You may not receive any mark if key assumption (any assumption not stated in question) is missing.





### Section 3: Bonus Question (1 mark )

11. "Know what you don't know": Predict your score for this assessment (excluding this bonus question). If your prediction is within  $\pm 2$  marks of your actual score, you will get a bonus "true understanding" 1 mark. 😊

