

CS2106 Operating Systems

Semester 1 2021/2022

Tutorial 4

IPC & Threads

- 1) (Shared Memory) In this question, we are going to analyze the pitfalls of having multiple processes accessing and modifying data at the same time. Compile and run the code in **shm.c**. When running the executable, you can specify two command line arguments, **n** and **nChild**; if no command line arguments are given, the default values will be used: **n** is initialized to 100 and **nChild** is initialized to 1.

The code in **shm.c** does a simple job: the parent creates **nChild** processes and then each process, the parent included, increases a shared memory location by **n** times. After all processes have finished incrementing the shared value, the parent prints the shared result and exits. Because the value of the shared memory is initialized to 0, we expect the printed value to be equal to $(1 + \text{nChild}) * n$ at the end. Is this always the case?

Try running the program multiple times with increasingly higher values for **n** and **nChild** (**n** = 10000 and **nChild** = 100 should do the trick). Explain the results.

Answer:

When incrementing the value, one process must 1) read the value from the memory, 2) modify it, and 3) write the value back to the memory. If two or more processes have their executions interleaved, it is very likely that one process will overwrite the value written by another process. In this case, it is impossible to have a deterministic value.

The reason why the number of increments and/or the number of processes must be large enough is to ensure that the processes do run at the same time in the system. Note that it takes some time to setup the recently created process, and if this time exceeds the time required for the increment, then the processes won't try to simultaneously read and write at the shared memory.

(Message Passing) In the lecture we talked about message passing between processes. Here we are implementing **Direct Communication** between master and worker processes.

Pseudo message passing API

Send(Pid , Msg);

Send **Msg** to the process whose process id is **Pid**.

Receive(Pid , Msg);

Receive **Msg** from the process whose process id is **Pid**.

As a continuation of Q3 from Tutorial 2, let us investigate how to implement the same prime factorization program using message passing. Fill in the pseudo code for the master and worker processes below. Your pseudocode must ensure that the master process receives results from workers in the same order as the user input is sent to workers by the master.

For each of the message passing operations, indicate whether it is blocking or non-blocking.

For simplicity, you can assume the following: We have 1 master process and N worker processes. Master process is responsible for distributing workload among worker processes. Worker processes have been already spawned via a fork call by master process.

Master Process:
<pre>workerPid[] = PID array for 1..N worker processes N = number of inputs inputs[] = N user inputs //Give user input to the corresponding worker to work on For i = 0 to N send(workerPid[i], inputs[i]); //non-blocking //Wait for each worker to finish and get back the result For i = 0 to N receive(workerPid[i], result); //blocking print result</pre>
Each Worker Process:
<pre>//Get the user input to work on from master receive(masterPid, input); //blocking result = PrimeFactorization(input); //Send back the result to master send(masterPid, result);</pre>

3) (Thread vs Process)

- a) For each of the following scenarios, discuss whether **multithreading** or **multiprocessing** is the best implementation model. If you think that both models have merits for a particular scenario, briefly describe additional criteria you would use to choose between the models.
 - i) Implementing a command line shell.

Process model. Key points: Executing another program in the newly created process + Memory + Protection.

- ii) Implementing the “tabbed browsing” in a web browser, i.e., each tab visits an independent webpage.

Arguable. Thread model (used by older Firefox) has lower overhead and suitable for less memory intensive webpages. Process model (used by most browsers nowadays) provides protection and can support multiple resource heavy webpages. In process model, if one tab becomes unresponsive while rendering, it will not affect other tabs because each tab’s rendering is handled by a process (isolation).

- iii) Implementing a complex multi-player game with dynamic environments and sophisticated.

Thread model. Games have a lot of shared state, can be time sensitive, and the process model requires overheads that are not necessary in this context.

- b) Suppose task A is CPU-intensive and task B is I/O intensive (reading from a file), which of the following statement(s) is/are TRUE if the two implementations are executed on a single CPU (single core)?

- i) If the threads are implemented as user threads, then multithreaded implementation will finish execution in shorter amount of time compared to the sequential implementation.

False: When Task B blocks for I/O, Task A is also blocked as the OS scheduler is not aware of user threads presence. A greater overhead from user thread scheduling will increase execution time compared to sequential implementation.

- ii) If the threads are implemented as kernel threads, then multithreaded implementation will finish execution in shorter amount of time compared to the sequential implementation.

True: Task B can block while Task A still runs on CPU. This is because the OS scheduler is aware of the kernel threads. Execution will be faster than sequential implementation where Task A will block when Task B requests I/O

4) (Threads)

- a) Suppose we have the following multi-threaded processes on a **hybrid thread** model:

Process ID	Number of User Threads	Number of Kernel Threads
P1	U1	K1
P2	U2	K2
P3	U3	K3

- i) If all threads are CPU intensive and run forever, what is the approximate proportion of CPU time utilized by the process P3 after running for a long time? You can assume a pre-emptive and fair process scheduler is used by the OS. You can express your answer in terms of U1-U3 and K1-K3.

$$\text{CPU time of P3} = \frac{K3}{K1+K2+K3}$$

- ii) When [U1 = 4, K1 = 3] [U2 = 3, K2 = 1] [U3 = 5, K3 = 2] what is the CPU utilization by one user thread of process P1? Does the number of user threads of other process affect this value? State your assumptions.
 12.5%. 50% of CPU time is allocated to entire Process P1. If user thread library allocates resources equally between user threads, one thread receives $50\% / 4 = 12.5\%$ of CPU time

This question tests the distinction between user thread and kernel thread. Only the latter is visible to OS (i.e. schedulable). So, regardless of the binding (1-1 1-N N-M) between user and kernel thread, it is the number of kernel threads that determine the CPU received.

- b) Evaluate each of the following statements regarding Thread and POSIX Thread (pthread) as True or False.
- i) All pthreads must execute the same function when they start.
 False: Pthread can start on any function as long as the function signature is `void* f(void*)`
 - ii) Multi-threaded program using pure user threads can never exploit multi-core processors.
 True: This is the main short coming of user thread as OS is not aware of the threads.
 - iii) If we use a 1-to-1 binding in a hybrid thread model, the end result is the same as pure-user thread model.
 False: The result is a pure kernel-thread model as each user thread is bounded to a kernel thread that can be scheduled.
 - iv) On a single-core processor that supports simultaneous multithreading, we can execute more than one thread at the **same time**.
 True: This is a strength of SMT processors (briefly discussed in lecture, you can explain more if needed).

Questions for your own exploration

- 5) (Protecting the Shared Memory) Allowing processes to access and modify shared data whenever they please can be problematic! Therefore, we would like to modify the code in **shm.c** such that the output is deterministic regardless of how large `n` and `nChild` are. More precisely, we want the processes to take turns when modifying the result value that resides in the shared memory. To this end, we will add another field in the shared memory, called the order value, that specifies which process' turn it is to increment the shared result. If the order value is 0, then the parent should increase the shared result, if the order value is 1, then the first child should increase it, if the order value is 2, then the second child and so on. Each process has an associated `pOrder` and checks whether the value order is equal to its `pOrder`; if it is, then it proceeds to increment the shared result. Otherwise, it waits until the order value is equal to its `pOrder`. Your task is to modify the code in **shm_protected.c** to achieve the desired result. You will have to:

- a) Create a shared memory region with two locations, one for the shared result, and the other one for the order value;
- b) Write the logic that allows a process to modify the shared result only if the order value is equal to its pOrder (the skeleton already takes care of assigning the right pOrder to each process);
- c) Print the result value and cleanup the shared memory.

Why is **shm** faster than **shm_protected**? Why is running **shm_protected** with large values for nChild particularly slow? (Hint: You may want to take a look at the output of *htop*)

Answer:

To guarantee that two processes won't try to modify the shared memory location at the same time, we must serialize the access to it, i.e., make it impossible for two processes to have access to it at the same time. Serializing the code will, of course, make the execution time longer.

The technique we use to prevent simultaneous access is called **busy waiting**. Each process is continuously checking whether its time to modify the variable has come. This requires the process to run on the CPU and consume CPU cycles (thus the name). Note that at the moment when it's Child X's time to increment the value, there are $nChild - X$ processes doing busy waiting – each of these processes will get scheduled to run on the CPU but the execution of the program will not make any progress. Because of this, your program will take a long time to complete for large values of nChild.