

CS2106 Operating Systems
Semester 1 AY2021/22

Tutorial 8 (Solution)
Disjoint Memory Allocation

1. (Walkthrough of Paging/Segmentation/Hybrid Schemes) Let us use a tiny example to understand the various disjoint memory schemes. For simplicity, we assume there are only two types of memory usage in a program: `text` (instruction) and `data` (global variables).

The following questions assume that program P has:

- 6 instructions, each fitting in a processor word (instruction words #1 to #6)
- 5 data words (data words #1 to #5)

a. (Paging) Given the following parameters:

- Page Size = Frame Size = 4 words
- Largest logical memory size = 16 words
- Number of physical memory frames = 16

Assuming that P's data region is placed right after the instruction region in the logical memory space, fill in the following page table. Use frames 5, 2, 10, 9 for pages 0, 1, 2, 3 respectively (note: you may not need all frames). Indicate the value of the valid bit for all page table entries.

Page#	Frame#	Valid
0		
1		
2		
3		

Find out the logical address and the corresponding physical address for the following actions taken by the processor.

Processor Action	Logical Address	Physical Address
Fetch the 1 st Instruction		
Load the 2 nd Data word		
Load the 3 rd Data word		
Load the 6 th Data word (This is intentionally outside of the range)		

- b. (Segmentation) Assuming `text` and `data` are stored in segments 0 and 1 respectively, fill in the following segment table. Use addresses 50 and 23 as the starting addresses for the two segments, respectively.

Segment#	Base Address	Limit
0		
1		

Similar to (a), find out the logical address and physical address for the following processor actions:

Processor Action	Logical Address	Physical Address
Fetch the 1 st Instruction		
Load the 2 nd Data word		
Load the 3 rd Data word		
Load the 6 th Data word (This is intentionally outside of the range)		

c. (Segmentation with Paging) Assuming the following parameters:

- Page Size = Frame Size = 4 words
- Number of physical memory frames = 16
- Maximum size of each segment = 4 pages

Furthermore, assume the pages from the code segment are allocated to frames 7, 4, 1, and 2 and data segment allocated to frames 9, 3, 14, and 6 (note that you may not need all of them).

Draw the segment and page tables for this setup, then fill in the processor action table. For the logical addresses, use the notation of <segment id, page number, offset>.

Processor Action	Logical Address	Physical Address
Fetch the 1 st Instruction		
Load the 2 nd Data word		
Load the 3 rd Data word		
Load the 6 th Data word (This is intentionally outside of the range)		

ANS:

a.
Memory layout:

Page 0	0	1 st Instruction
	1	2 nd Instruction
	2	3 rd Instruction
	3	4 th Instruction
Page 1	4	5 th Instruction
	5	6 th Instruction
	6	1 st Data word
	7	2 nd Data word
Page 2	8	3 rd Data word
	9	4 th Data word
	10	5 th Data word
	11	<empty>
Page 3	12	<empty>
	13	<empty>
	14	<empty>
	15	<empty>

Page#	Frame#	Valid
0	5	T
1	2	T
2	10	T
3	--	F

Processor Action	Logical Address	Physical Address
Fetch the 1 st Instruction	0	$5 \times 4 + 0 = 20$
Load the 2 nd Data word	7	$2 \times 4 + 3 = 11$
Load the 3 rd Data word	8	$10 \times 4 + 0 = 40$
Load the 6 th Data word (This is intentionally outside of the range)	11	$10 \times 4 + 3 = 43$

Note 1: Observe that the consecutive logical addresses may not be consecutive in physical memory (e.g. 2nd and 3rd Data word).

Note 2: Incorrect memory access is not catchable if it is still within valid page boundary.

b.

Segment#	Base Address	Limit
0	50	6
1	23	5

Processor Action	Logical Address	Physical Address
Fetch the 1 st Instruction	<0, 0>	$50 + 0 = 50$
Load the 2 nd Data word	<1, 1>	$23 + 1 = 24$
Load the 3 rd Data word	<1, 2>	$23 + 2 = 25$
Load the 6 th Data word	<1, 5>	Triggers memory addressing error

(This is intentionally outside of the range)		
--	--	--

- c. Memory layout (code segment on the left, data segment on the right; empty pages not shown):

Page 0	0	1 st Instruction
	1	2 nd Instruction
	2	3 rd Instruction
	3	4 th Instruction
Page 1	4	5 th Instruction
	5	6 th Instruction
	6	<empty>
	7	<empty>

Page 0	0	1 st Data Word
	1	2 nd Data Word
	2	3 rd Data Word
	3	4 th Data Word
Page 1	4	5 th Data Word
	5	<empty>
	6	<empty>
	7	<empty>

Segment#	Page Limit	Page Table Base
0	2	
1	2	

Page#	Frame#	Valid
0	7	T
1	4	T
2	--	F
3	--	F

Page#	Frame#	Valid
0	9	T
1	3	T
2	--	F
3	--	F

Processor Action	Logical Address	Physical Address
Fetch the 1 st Instruction	<0, 0, 0>	$7 \times 4 + 0 = 28$
Load the 2 nd Data word	<1, 0, 1>	$9 \times 4 + 1 = 37$
Load the 3 rd Data word	<1, 0, 2>	$9 \times 4 + 2 = 38$
Load the 6 th Data word (This is intentionally outside of the range)	<1, 1, 1>	$3 \times 4 + 1 = 13$

Note: As the limit in segment table now refers to the limit of pages in the corresponding page table, memory error that is in page boundary is now "not catchable".

- (Paging and TLB) In this question, we attempt to quantify the benefit of using TLB by looking at the memory access time. Suppose the system uses the paging scheme with the page tables entirely stored in physical memory (DRAM). The page size is 4KB, and the logical addresses are 32-bit long. Answer the following:
 - Assuming the system does not use TLB and accessing DRAM takes 50ns (nanoseconds), what is the latency of accessing a global variable of type char?

- b. Assuming the system uses a TLB and 75% of all page-table references hit in the TLB. What is the average memory access time? You can assume that looking up a page table entry in TLB takes negligible time.
- c. How many entries does a TLB need to have to achieve a hit ratio of 75%? Assume the program generates logical memory addresses uniformly at random. Do you think a TLB in an actual machine is this large? If not, then how is it possible to achieve a high TLB-hit rate?

ANS:

a. $50\text{ns} (\text{access page table}) + 50\text{ns} (\text{access actual item}) = 100\text{ns}$

b. $0.75 * 50\text{ns} + (1-0.75)*100\text{ns} = 62.5\text{ns}$

c. Total number of page table entries = $2^{(32-12)} = 2^{20} = 1\text{M}$ entries

If we assume the memory access are **uniformly distributed** between all the possible pages, then we need at least $1\text{M} * 0.75 = 786,432$ entries in the TLB. Using this reasoning, it seems impossible to have a TLB large enough to guarantee a high hit-rate. You may be surprised that TLB is actually quite small in real-world machines (commonly in the range of 32 to 1024 entries) due to the high hardware complexity (associative search is expensive in term of hardware). However, TLB hit ratio is still very high in actual usage (commonly with 99% hit rate!). The main contributing reason is that memory access are **not uniformly distributed**. This can be understood with the aid of locality principle: it is more likely to have repeated memory accesses to same (temporal locality) or different (spatial locality) parts of the same memory page in a time interval rather than uniformly spread accesses.

3. (Segmentation) [Adapted from AY1617 Exam Paper] Suppose we have the following machine setup:

- Pure segmentation scheme, with four default segments: (0 = text, 1 = data, 2 = heap, 3 = stack).
- A special set of four hardware registers known as **segment base registers (SBR)**: **SBR0**, **SBR1**, **SBR2** and **SBR3**. Each **SBR** has two fields:
 - o **SAddr** = point to the starting address of a segment
 - o **Limit** = maximum valid offset of a segment.

For simplicity, you can use array indexing syntax to use one of the **SBRs**, e.g. **SBR[0]** refers to the **SBR0**. So, **SBR[1].SAddr** refers to the starting address as stored in **SBR1**.

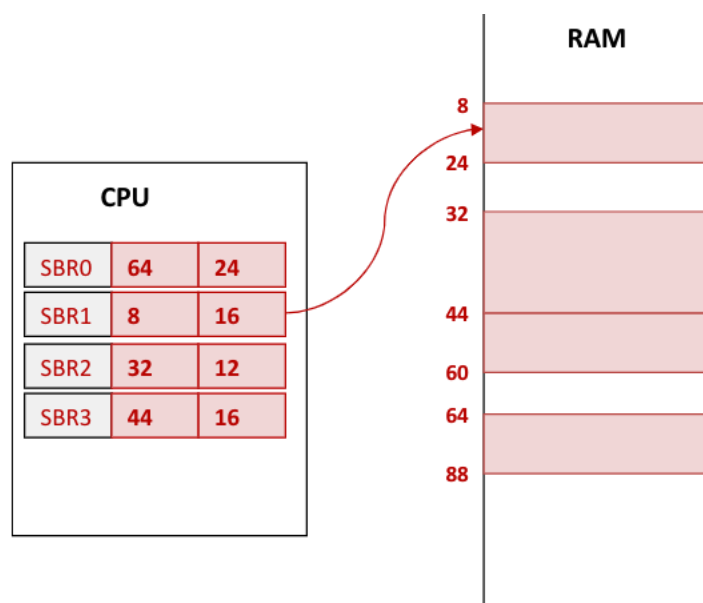
- a. Illustrate (draw) a sample process with the following running information:
 - SBR0 stores the stack region {SAddr = 64, Limit = 24}
 - SBR1 stores the data region {SAddr = 8, Limit = 16}
 - SBR2 stores the heap region {SAddr = 32, Limit = 12}
 - SBR3 stores the text region {SAddr = 44, Limit = 16}

There is no need to worry about the scale of region (how large the region is), however, you need to indicate all addresses / limits correctly as well as the relationship between the SBRs and the regions.

- b. Briefly describe how to handle **malloc()** (dynamic memory allocation) for the sample process in (a).
- c. Do we need caching mechanism like TLB in this machine? Briefly explain.
- d. Briefly explain how **thread switching**, i.e. switching between two threads in the same process, can be supported on this machine **efficiently**. Your answer only need to focus on the memory context (i.e. ignore hardware context etc) and minimally should contain the following 2 points:
 - What should be saved as part of the **thread context**?
 - Use (a) as an example to illustrate the basic ideas.

ANS:

a.



b. As the heap region is already bordering on the text region. We need to find a larger free piece of memory for relocating the heap region. SBR2 will store new base and limit afterwards.

c. We do not need TLB in this machine The SBRs are hardware registers, which give very fast access to the physical address lookup already.

d. Key ideas:

- Only stack region needs to be switched. So, the SBR0 should be saved as part of a text context.

Upon switching, we simply swap the SBR0 to the one saved in the thread context.

4. [Adapted from AY1516 Exam Paper] As discussed in lecture, we can protect a memory page by adding permission bits to the page table entry (PTE). Suppose we add 3 access right bits: {**R**: Readable, **W**: Writable, **X**: Executable} to each PTE. When a processor

instruction violates the access permission of a page, OS will be invoked to handle the problem. We can utilize this behavior to implement the **copy-on-write** mechanism. Also discussed in lecture, copy-on-write can be used to reduce the memory usage during the `fork()` system call, by allowing memory pages to be shared between parent and child process until modified. In this question, we will explore these ideas using a simplified example.

Suppose process P has only 3 valid page table entries:

Page No	Frame No	R	W	X
0	7	1	0	1
1	2	1	1	0
2	5	1	0	0
3 N-1	---	---	---	---

- a. Give the page table entries for the **child process** after P executes **fork()**. If you need to use any new frame numbers, use them in this order {6, 0, 3, 4, 1}

Using the above scenario, explain the following:

- b. What are the steps required to handle copy-on-write? Indicate any additional information that OS need to maintain. Show the affected PTE(s) for the child process afterwards.

ANS:

a.

Page No	Frame No	R	W	X
0	7	1	0	1
1	2	1	0	0
2	5	1	0	0
3 N-1	---	---	---	---

Page #0 and #2 are read-only pages, which need not be duplicated, i.e. simply shared between parent and child processes. Page #1's write

- b. Memory store triggers against page with no writable permission.

OS checks that this is a copy-on-write page (additional information).

- a. allocate new frame
- b. copy current content over to (a)
- c. Re-perform the memory store

Note: Linux maintains a separate data structure (called virtual memory area – VMA) in addition to the page table to keep track of memory mappings and what are the access permissions for every memory area and page. If a particular page is marked as writable in the VMA and the corresponding entry in the page table for that page is marked as non-writable, the OS recognizes that page as a copy-on-write page (COW). When a process tries to modify such a page, a page fault occurs, and the page fault handler handles the COW and performs the memory store.

Page #1's PTE become:

Page No	Frame No	R	W	X
1	6	1	1	0

Where frame #6 contains mostly the original page#1 information and the new memory store value.

For your own exploration

5. (Dynamic Allocation, adapted from [SGG]) It is possible for a program to dynamically allocate (i.e., enlarge the memory usage) during runtime. For example, the system call **malloc()** in C or **new** in Java/C++ can enlarge the **heap region** of process memory. Discuss the OS mechanisms needed to support dynamic allocation in the following schemes:
- a. Contiguous memory allocation (both fixed and dynamic size partitioning)
 - b. Pure Paging
 - c. Pure Segmentation

ANS:

- a. It is simpler to have the heap region to be allocated at the end of the logical memory space. Then, we can enlarge the heap region by enlarging the partition allocated to the process.

Under fixed partitioning:

- No extra work is needed, as heap region can simply use up the free space (the internal fragmentation) between the partition size and the actual memory size

Under dynamic partitioning:

- If the adjacent partition is free:
- Simple: Simply modify the partition information, i.e. change the length of current partition and shorten the free partition.


- If adjacent partitions are occupied:
 - More troublesome: The current partition cannot be enlarged. Relocation is required. OS need to look for a large enough free partition to fit the enlarged partition. Once located, the current partition is moved over.
 - b. Due to internal fragmentation of the paging scheme, it is possible that the allocation can use the remaining free space in the page. Suppose the allocation overshoot the page boundary, then OS needs to look for a free physical frame f . Afterward, update the page table by changing the first invalid page table entry from invalid to valid and fill in frame number f .
 - c. Idea is similar to the dynamic partitioning. If there is free memory at the end of the heap segment, then OS can simply update the limit of the segment and reduce the size of the affected free partition. If there is no free memory, then relocation is required. After relocation, both base and limit of the heap segment needs to be updated.
6. Calculate the minimum and maximum percentage of memory capacity lost to internal fragmentation in a system that uses the Buddy allocator. Can the Buddy allocator suffer from external fragmentation?

ANS:

For an allocation request of N bytes: min fragmentation = 0% (exact fit, $N=2^k$), max fragmentation = ~50% ($N=2^{(k-1)+1}$) Note that internal fragmentation >50% is not possible.

External fragmentation can still happen!

7. (Growing/Shrinking of 2 Regions) This question looks at the problem of maximizing the **logical memory space** for two growing/shrinking regions (e.g., Stack and Heap regions). Suppose we have only a piece of 1,000 bytes of memory space. Which of the following placements of the stack and heap regions is the best choice? The arrows represent the growing direction of the regions. Briefly justify your choice:



ANS:

Designs #1 and #2 place a pre-defined upper limit on the two regions. These upper limits may not make sense for certain execution behaviour. You can easily find execution scenario that utilize heavily one of the regions over the other (e.g., heavy recursive function calls → stack region may be exhausted, heavy use of malloc → heap

region may be exhausted instead). In such cases, one of the regions may reach the upper limit even though there may still be free memory space.

Design three allows free growing of the two regions. The regions reach their maximum size only when the growing boundary of the two regions meet, i.e. the memory space is fully utilized. This is the most commonly seen memory layout (heap and stack grows toward each other).